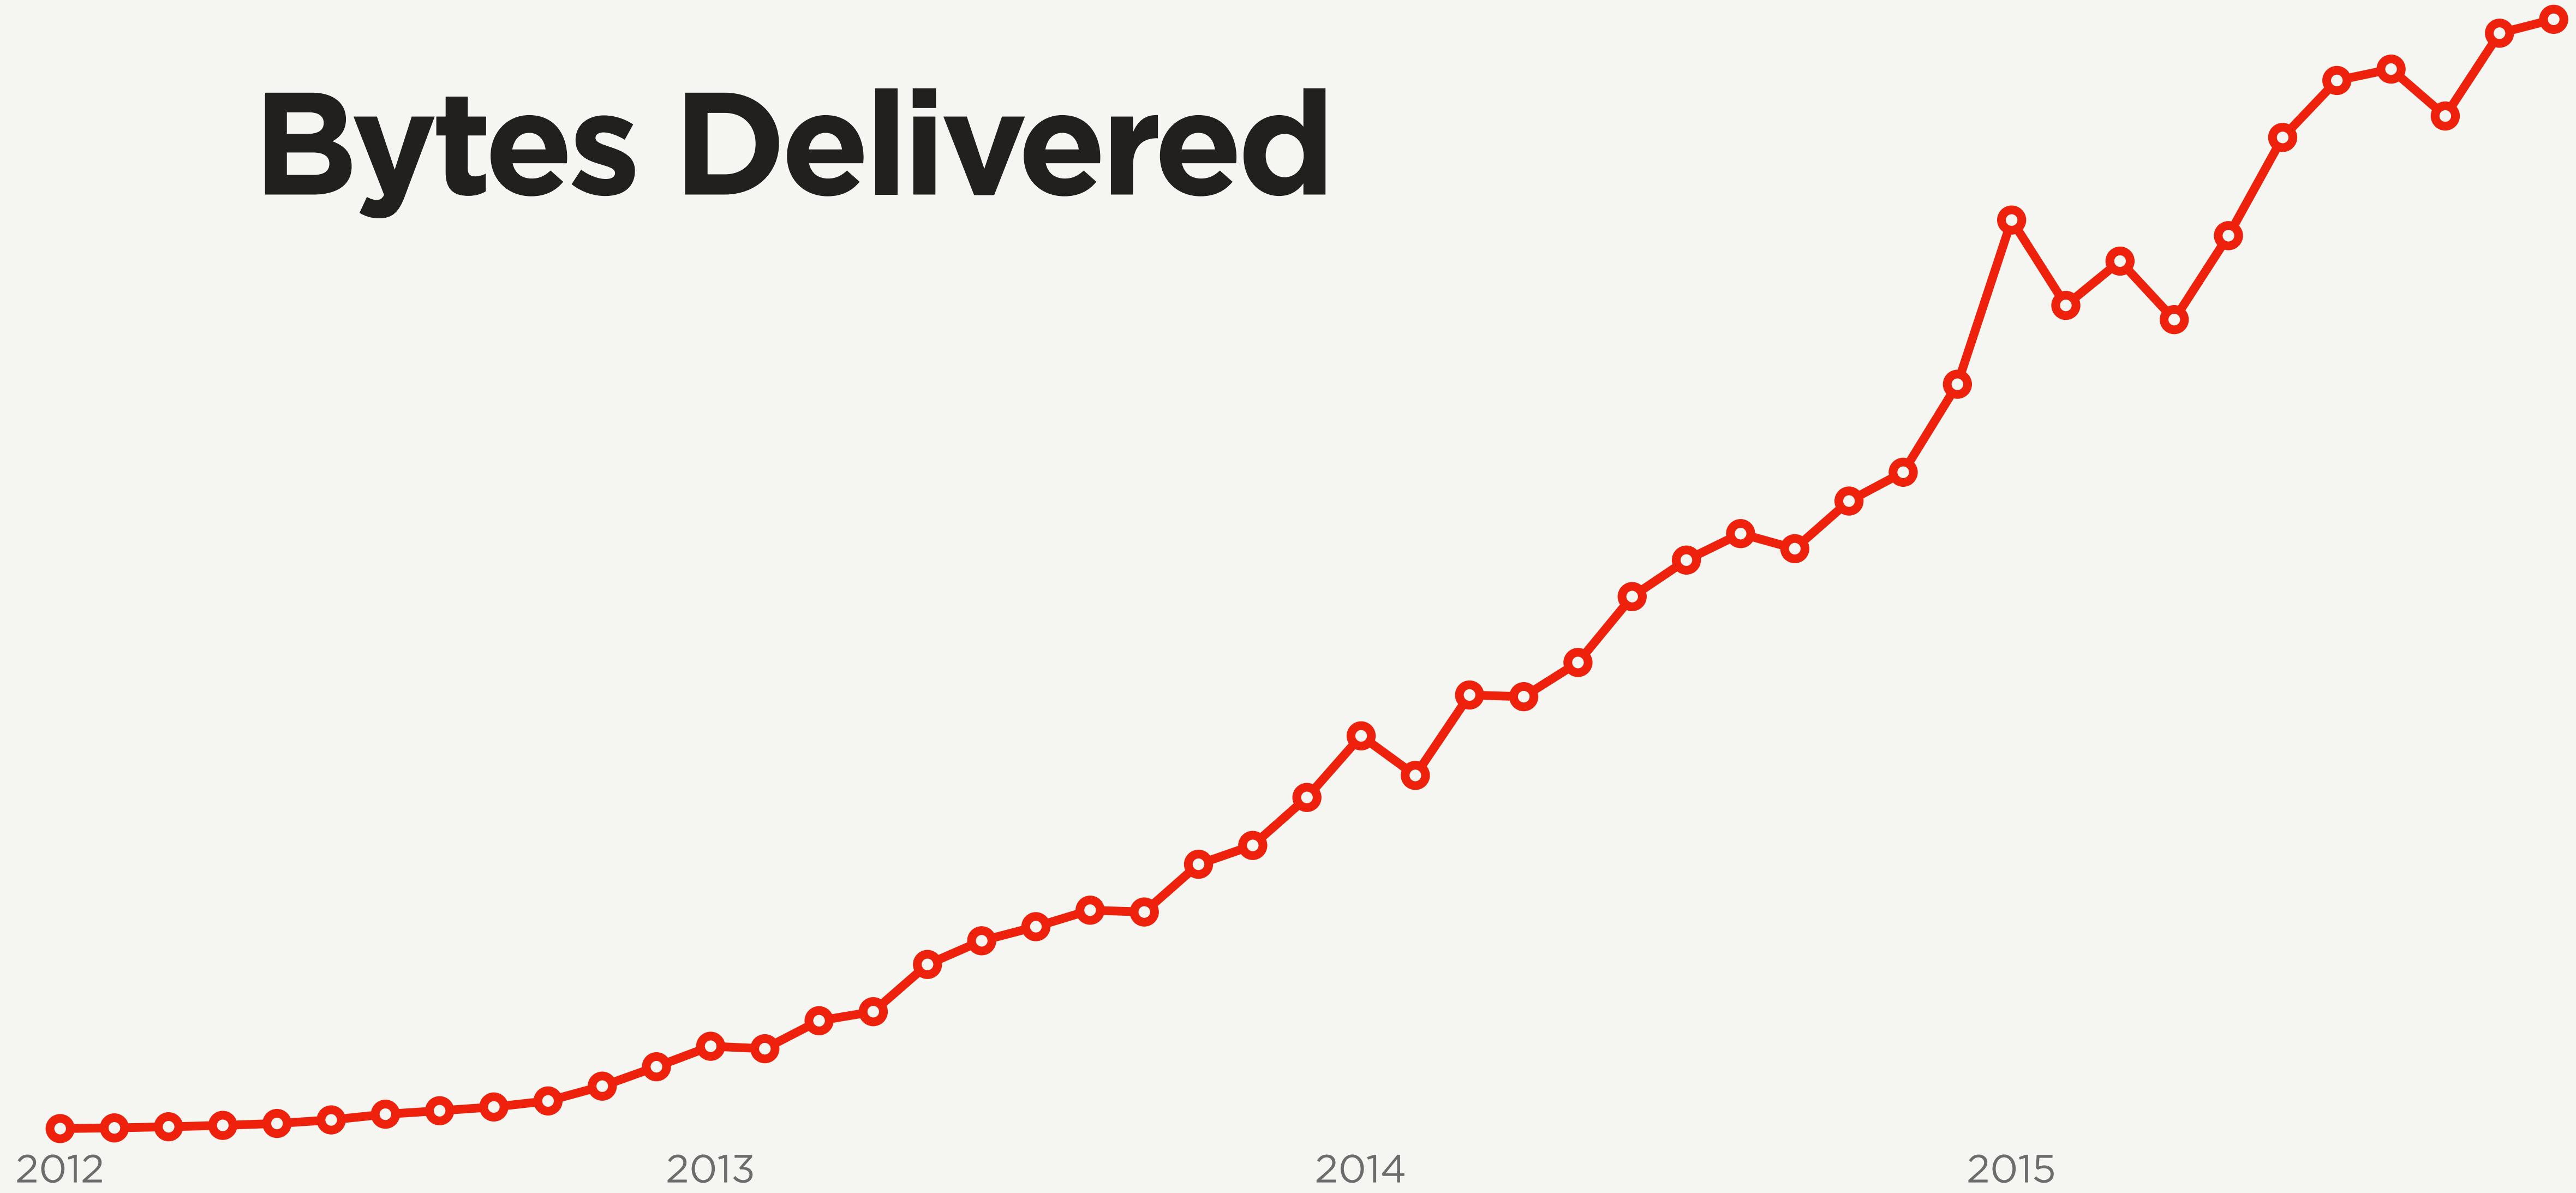


Zuul's Journey to Non-Blocking

Arthur Gonigberg
Cloud Gateway

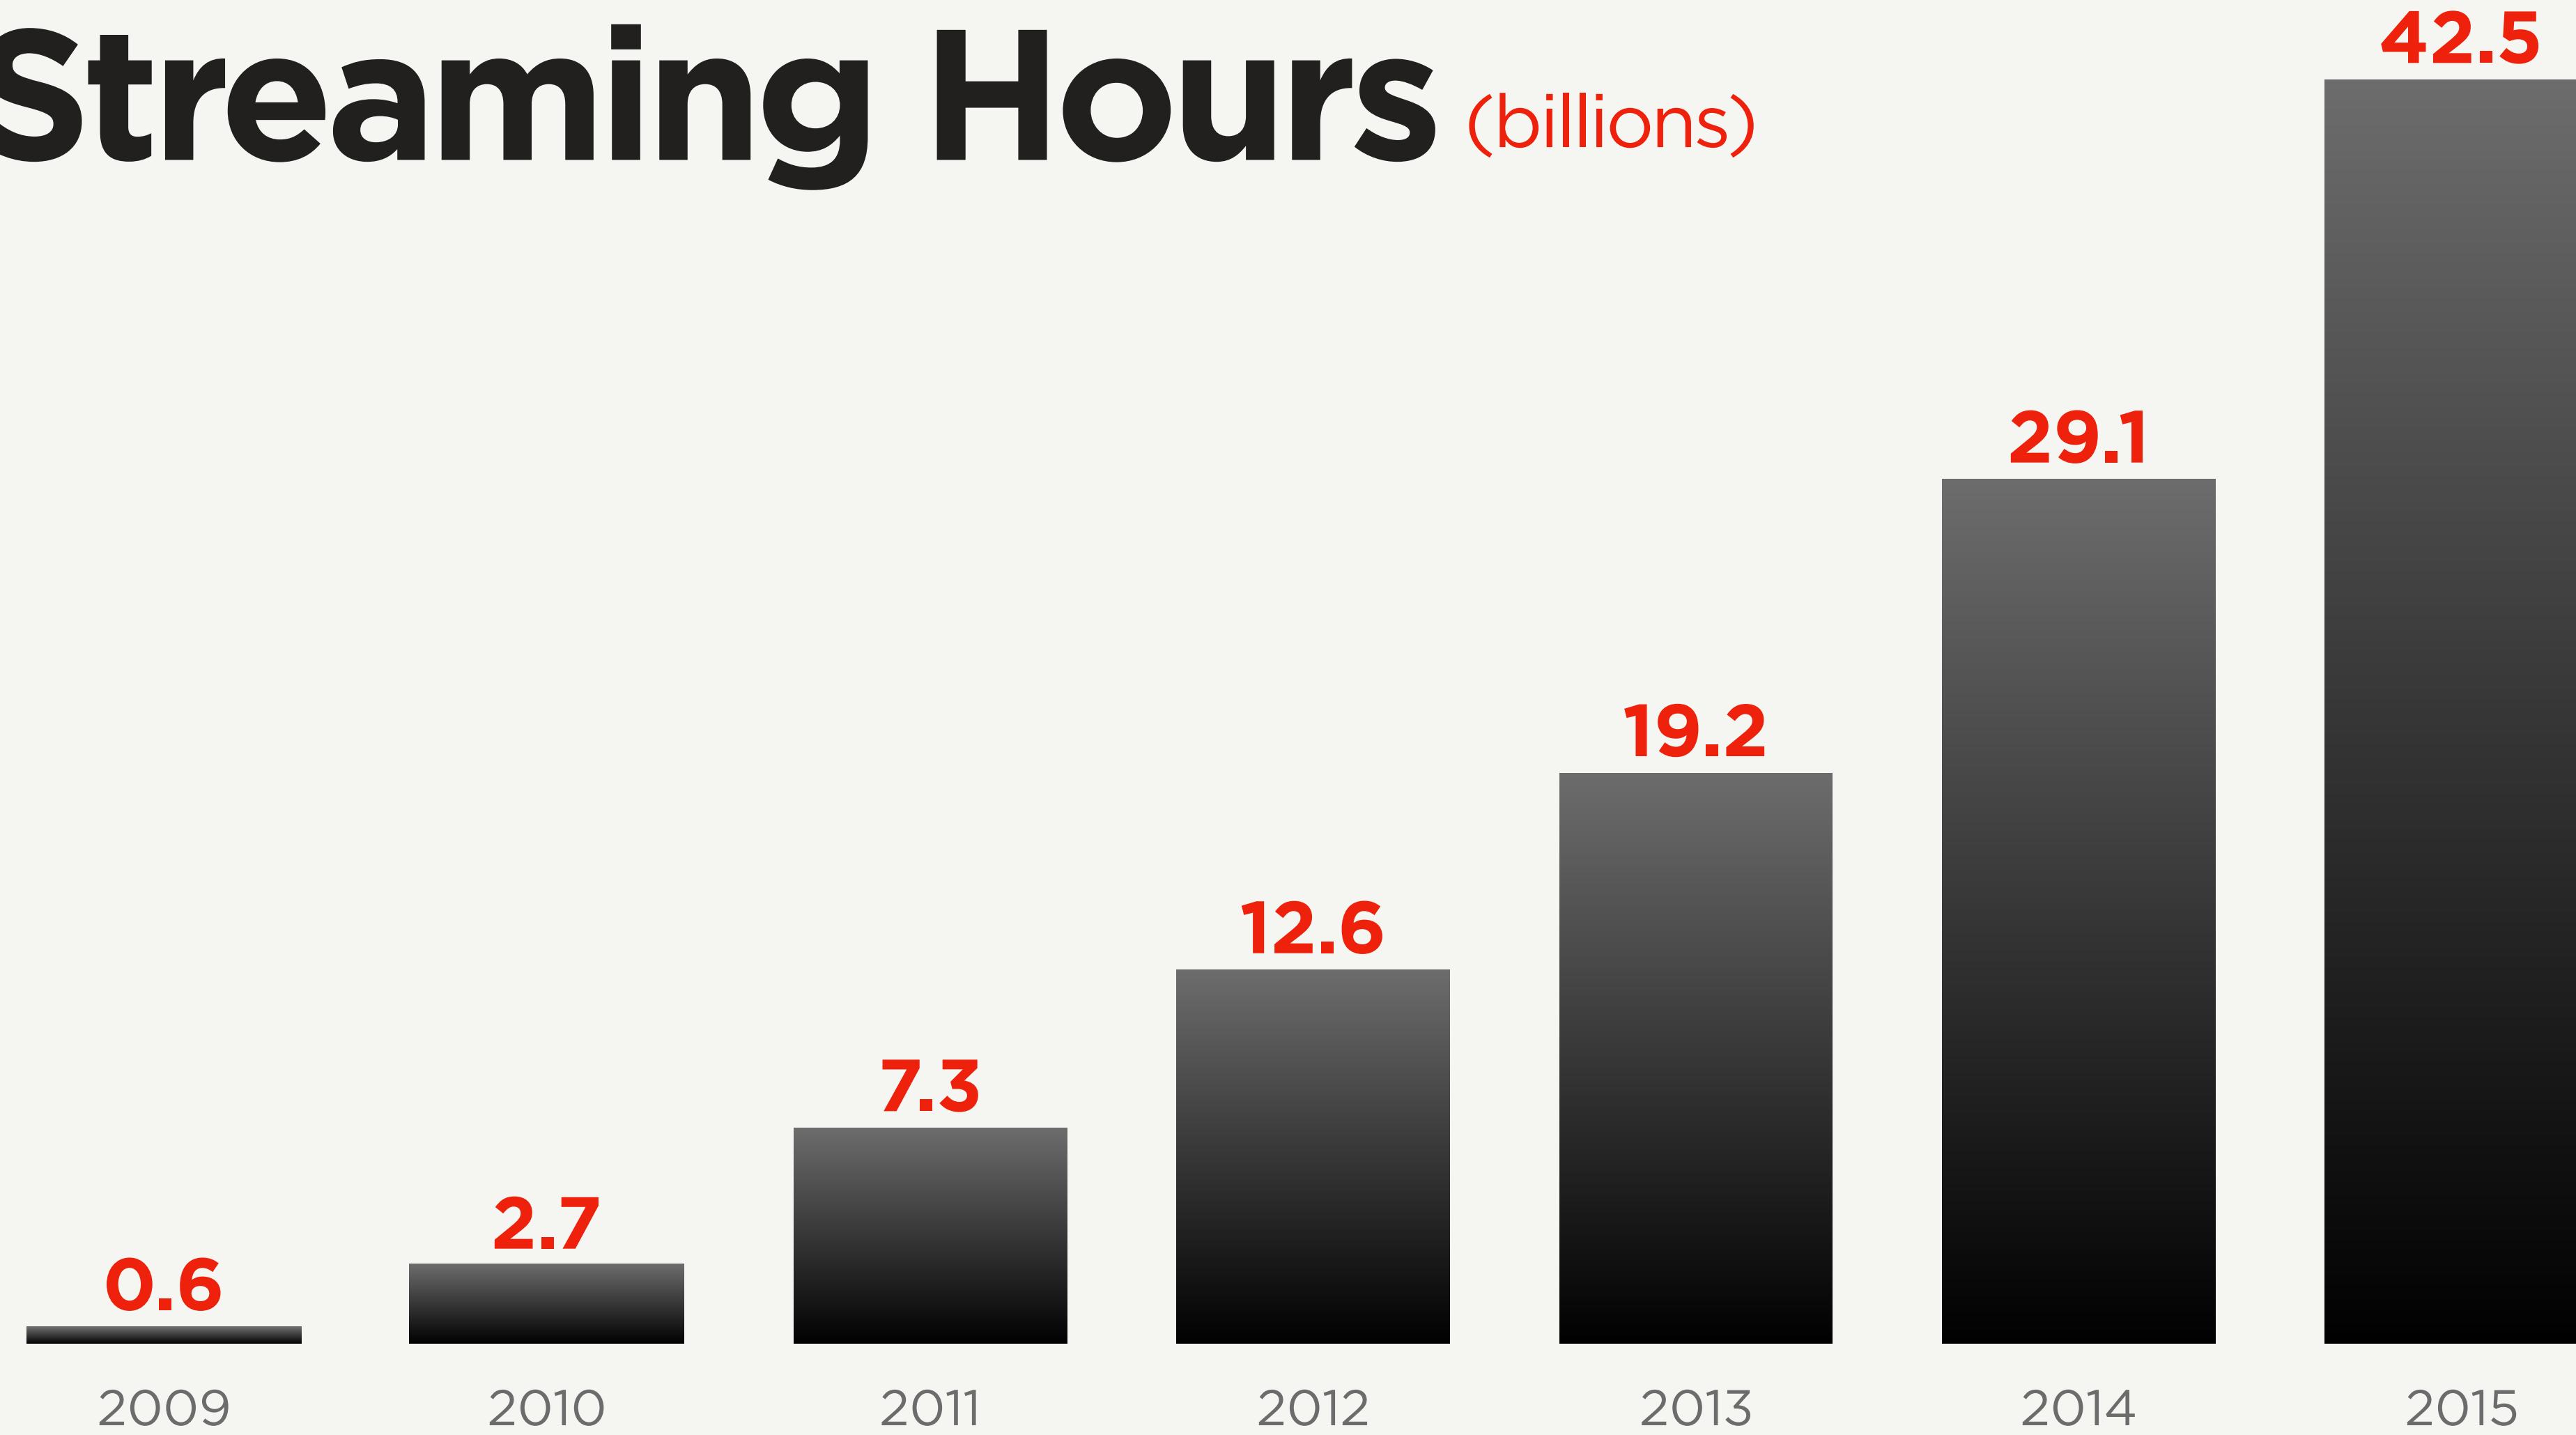


Bytes Delivered



Streaming Hours

(billions)



Zuul.

Routing

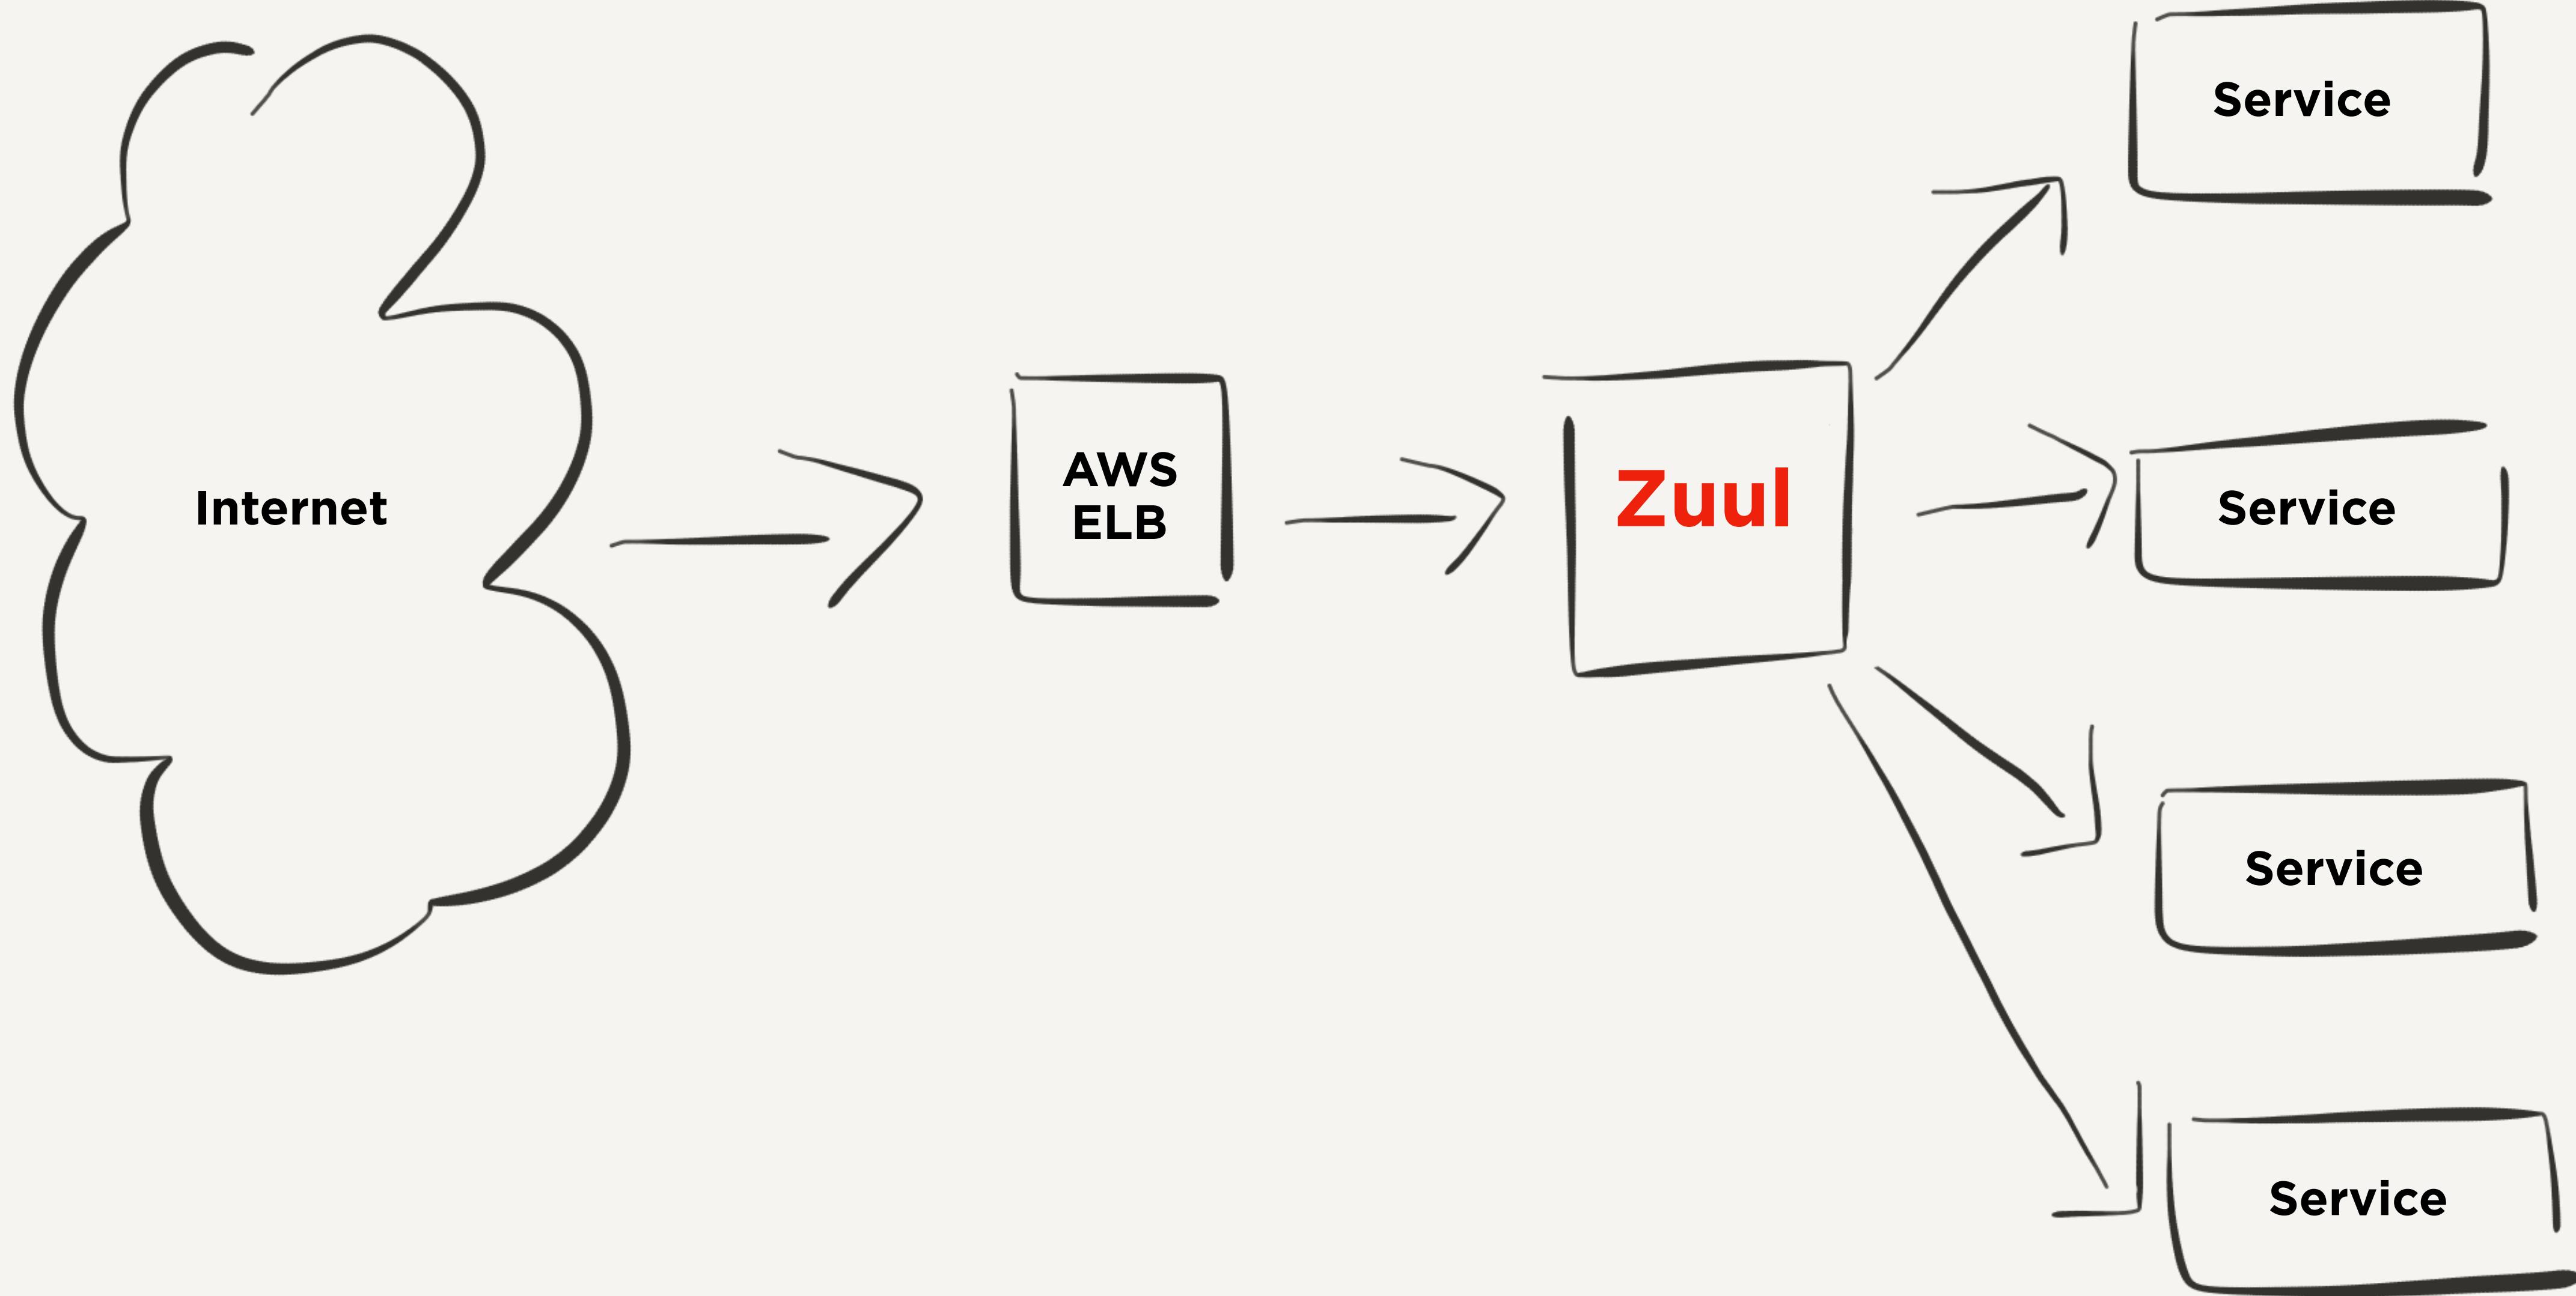
Monitoring

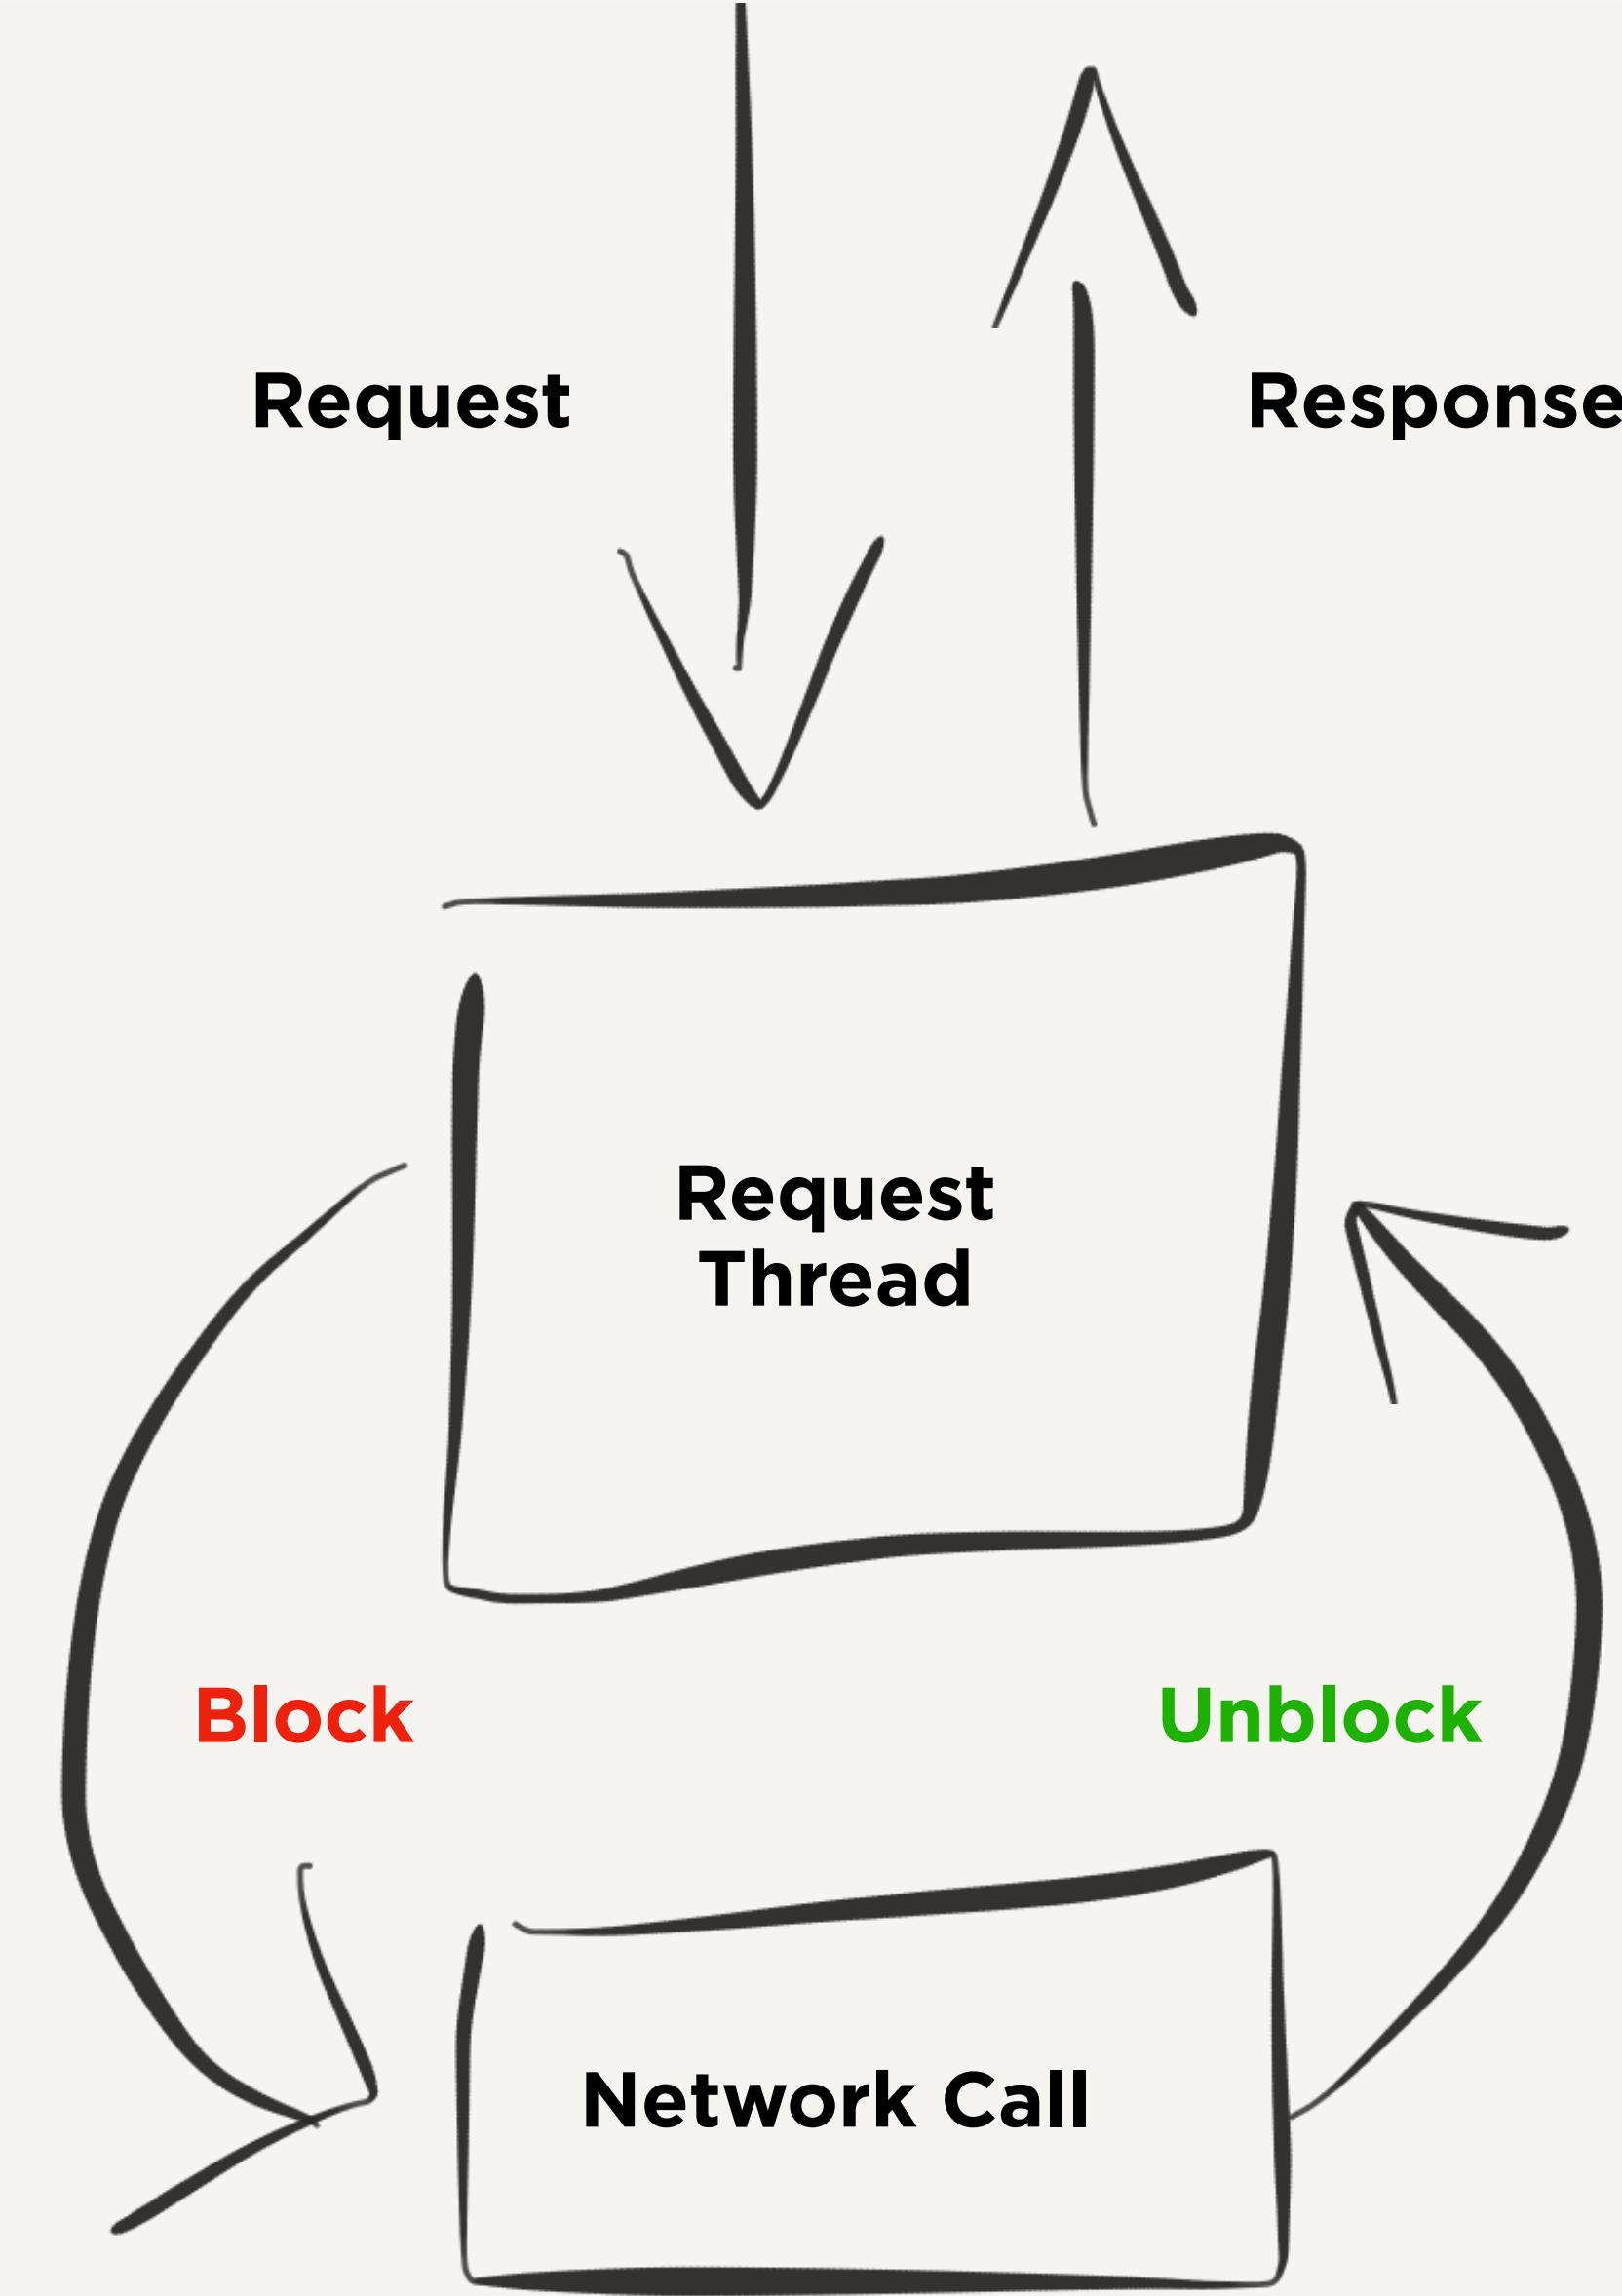
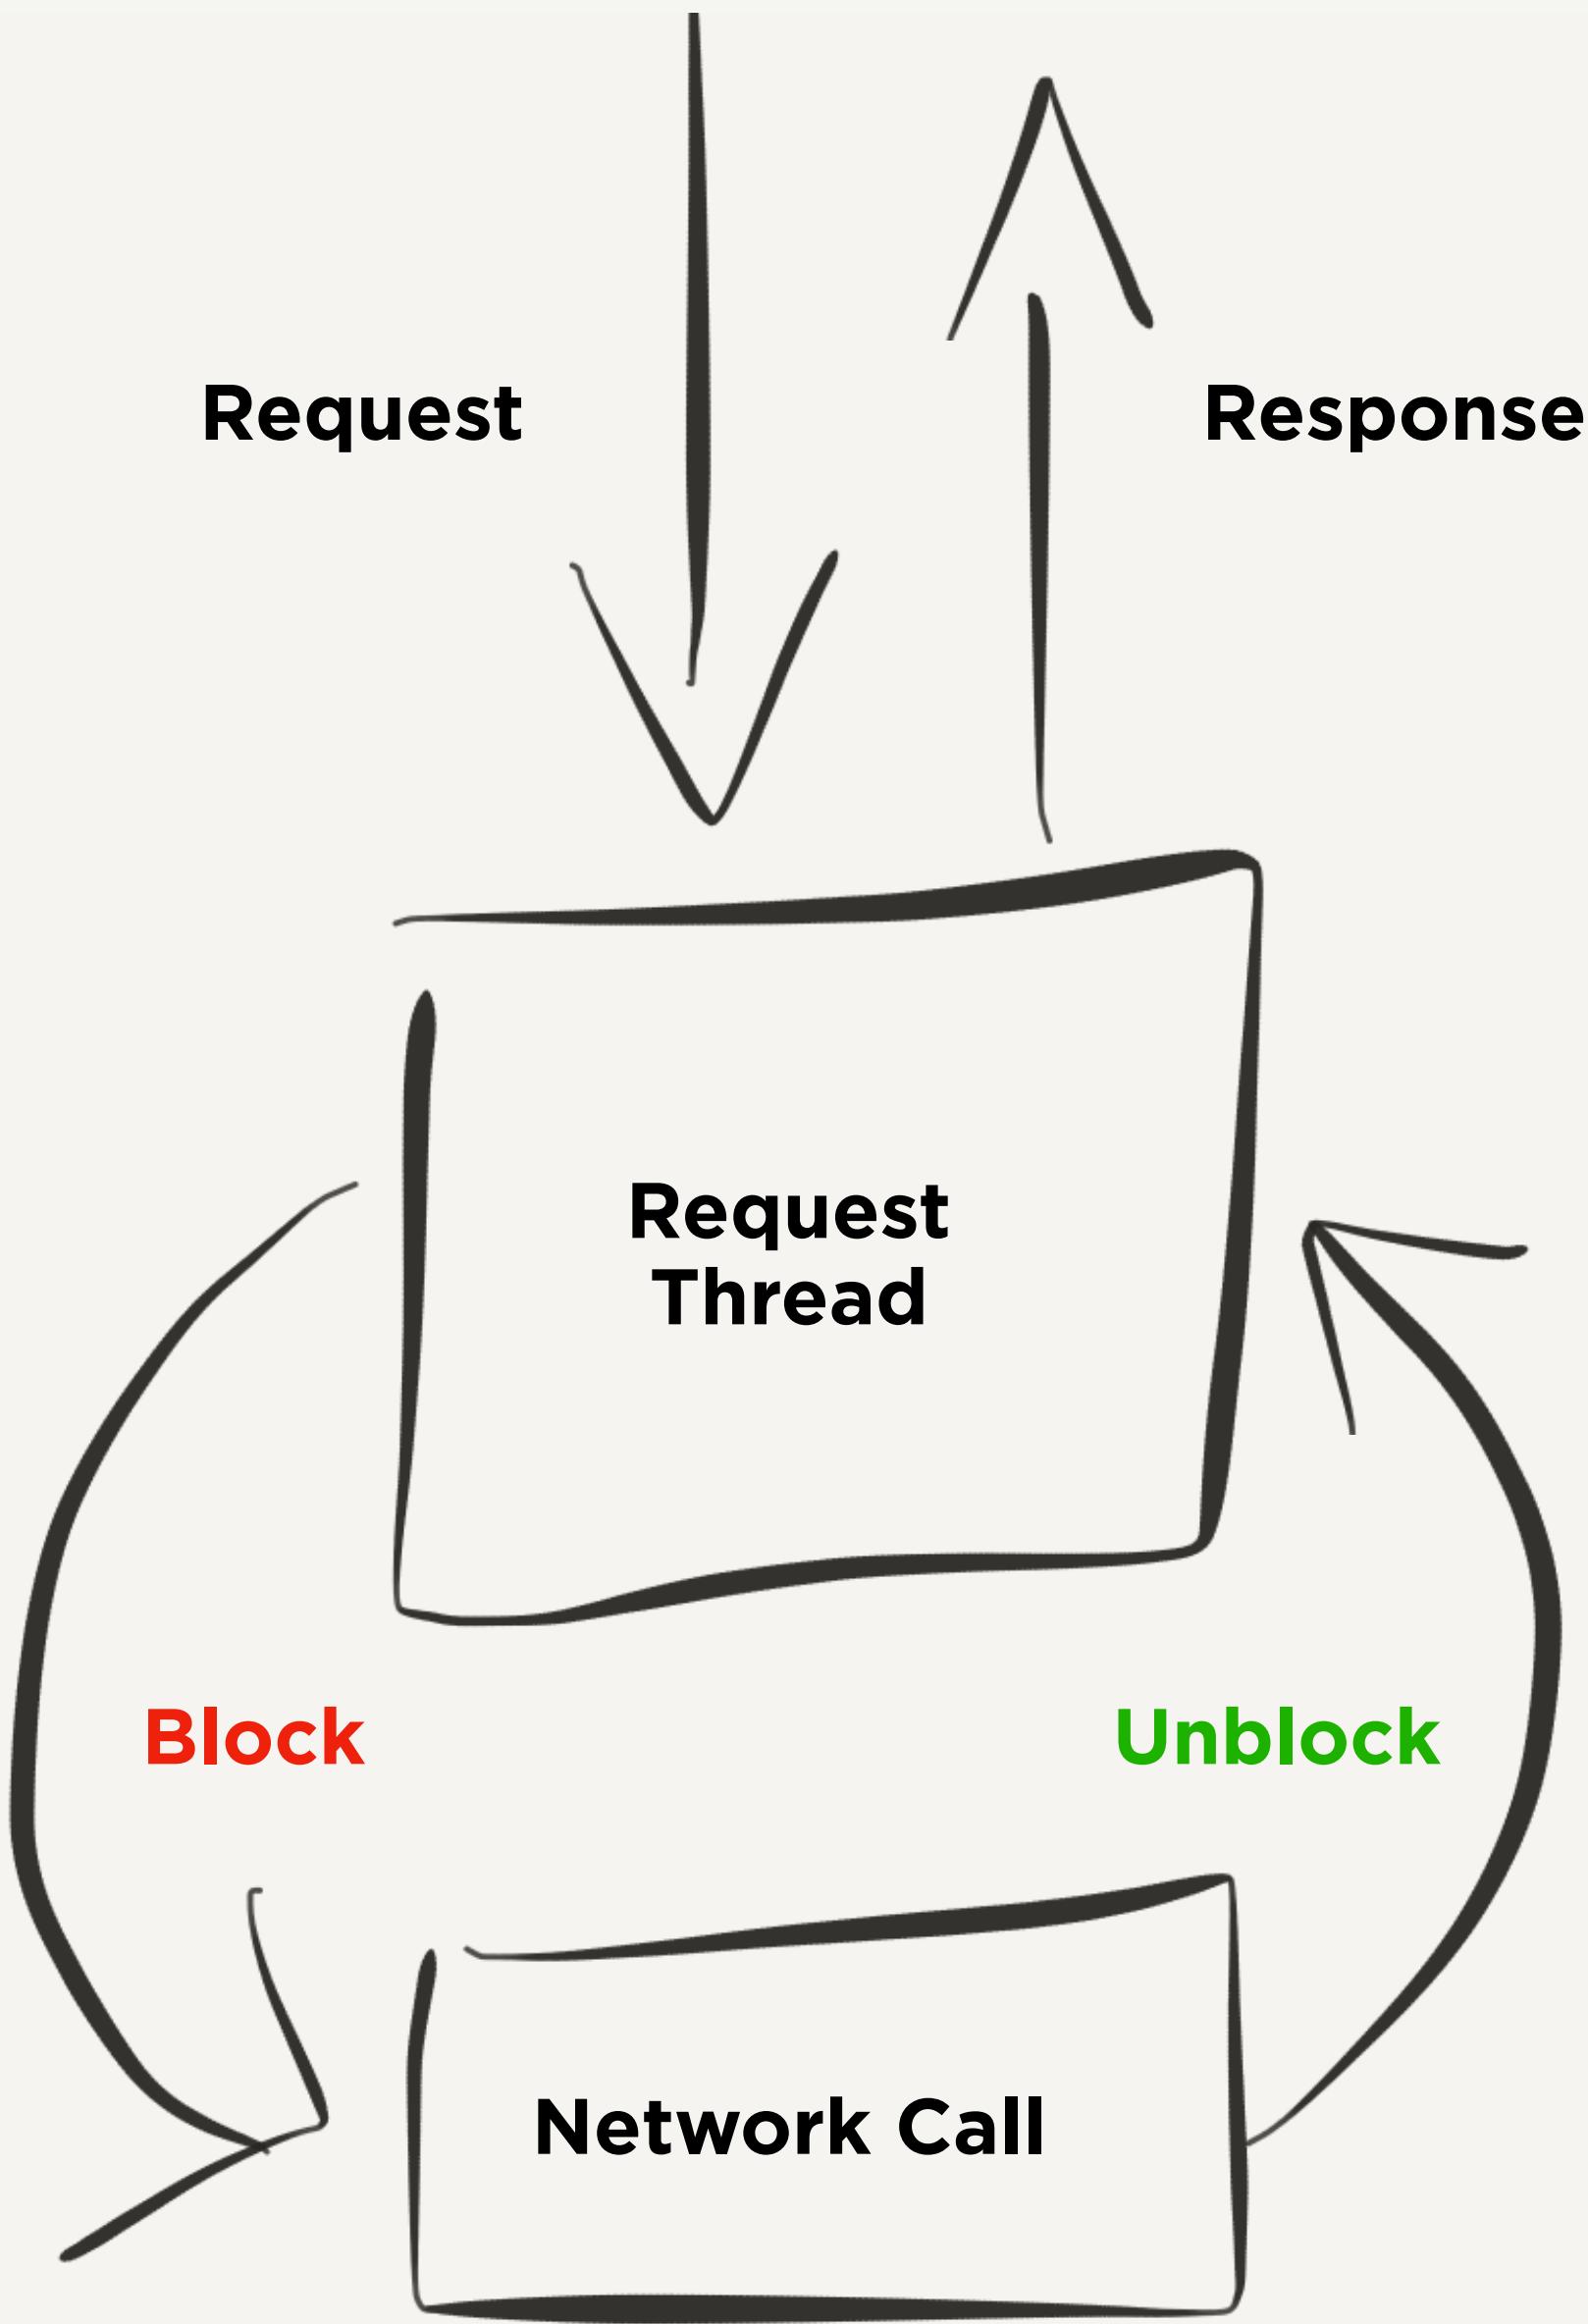
Security

Resiliency

Flexibility









Netflix Technology Blog [Follow](#)

Learn more about how Netflix designs, builds, and operates our systems and engineering organizati...

Nov 26, 2012 · 2 min read

Introducing Hystrix for Resilience Engineering

Netflix Brings Reactive Extensions to Java



Like

| by [Jonathan Allen](#) on Feb 06, 2013. Estimated reading time: 1 minute |

[Discuss](#)



Netflix Technology Blog [Follow](#)

Learn more about how Netflix designs, builds, and operates our systems and engineering organizati...

Feb 4, 2013 · 7 min read

Reactive Programming in the Netflix API with RxJava

The Reactive Manifesto

Published on September 16 2014. (v2.0)

Hello Netty

```
public class HttpServerHandler extends SimpleChannelInboundHandler<HttpObject> {

    @Override
    protected void channelRead0(ChannelHandlerContext ctx, HttpObject msg) throws Exception {
        if (msg instanceof LastHttpContent) {
            ByteBuf content = Unpooled.copiedBuffer("Hello World.", CharsetUtil.UTF_8);
            FullHttpResponse response = new DefaultFullHttpResponse(HttpVersion.HTTP_1_1,
                HttpResponseStatus.OK, content);
            response.headers().set(HttpHeaders.Names.CONTENT_TYPE, "text/plain");
            response.headers().set(HttpHeaders.Names.CONTENT_LENGTH, content.readableBytes());
            ctx.write(response);
        }
    }

    @Override
    public void channelReadComplete(ChannelHandlerContext ctx) throws Exception {
        ctx.flush();
    }
}
```

-> % ab -c 200 -n 1000000 -k http://localhost:8080/

This is ApacheBench, Version 2.3 <\$Revision: 1757674 \$>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/

Benchmarking localhost (be patient)
Completed 100000 requests
Completed 200000 requests
Completed 300000 requests
Completed 400000 requests
Completed 500000 requests
Completed 600000 requests
Completed 700000 requests
Completed 800000 requests
Completed 900000 requests
Completed 1000000 requests
Finished 1000000 requests

Server Software:
Server Hostname: localhost
Server Port: 8080

Document Path: /
Document Length: 11 bytes

Concurrency Level: 200
Time taken for tests: 13.587 seconds
Complete requests: 1000000
Failed requests: 0
Keep-Alive requests: 1000000
Total transferred: 100000000 bytes
HTML transferred: 11000000 bytes

Requests per second: 73602.19 [#/sec] (mean)

Time per request: 2.717 [ms] (mean)
Time per request: 0.014 [ms] (mean, across all concurrent requests)
Transfer rate: 7187.71 [Kbytes/sec] received

Real World Performance



Requests per second per node

WHAT IF I TOLD YOU

IT DEPENDS

Today's Journey

Rebuilding Zuul on Netty and RxJava

Operating Zuul 2.0 in the Real World

When to go Blocking or Non-Blocking

Starting the Journey

Rebuilding Zuul on
Netty and RxJava



Zuul

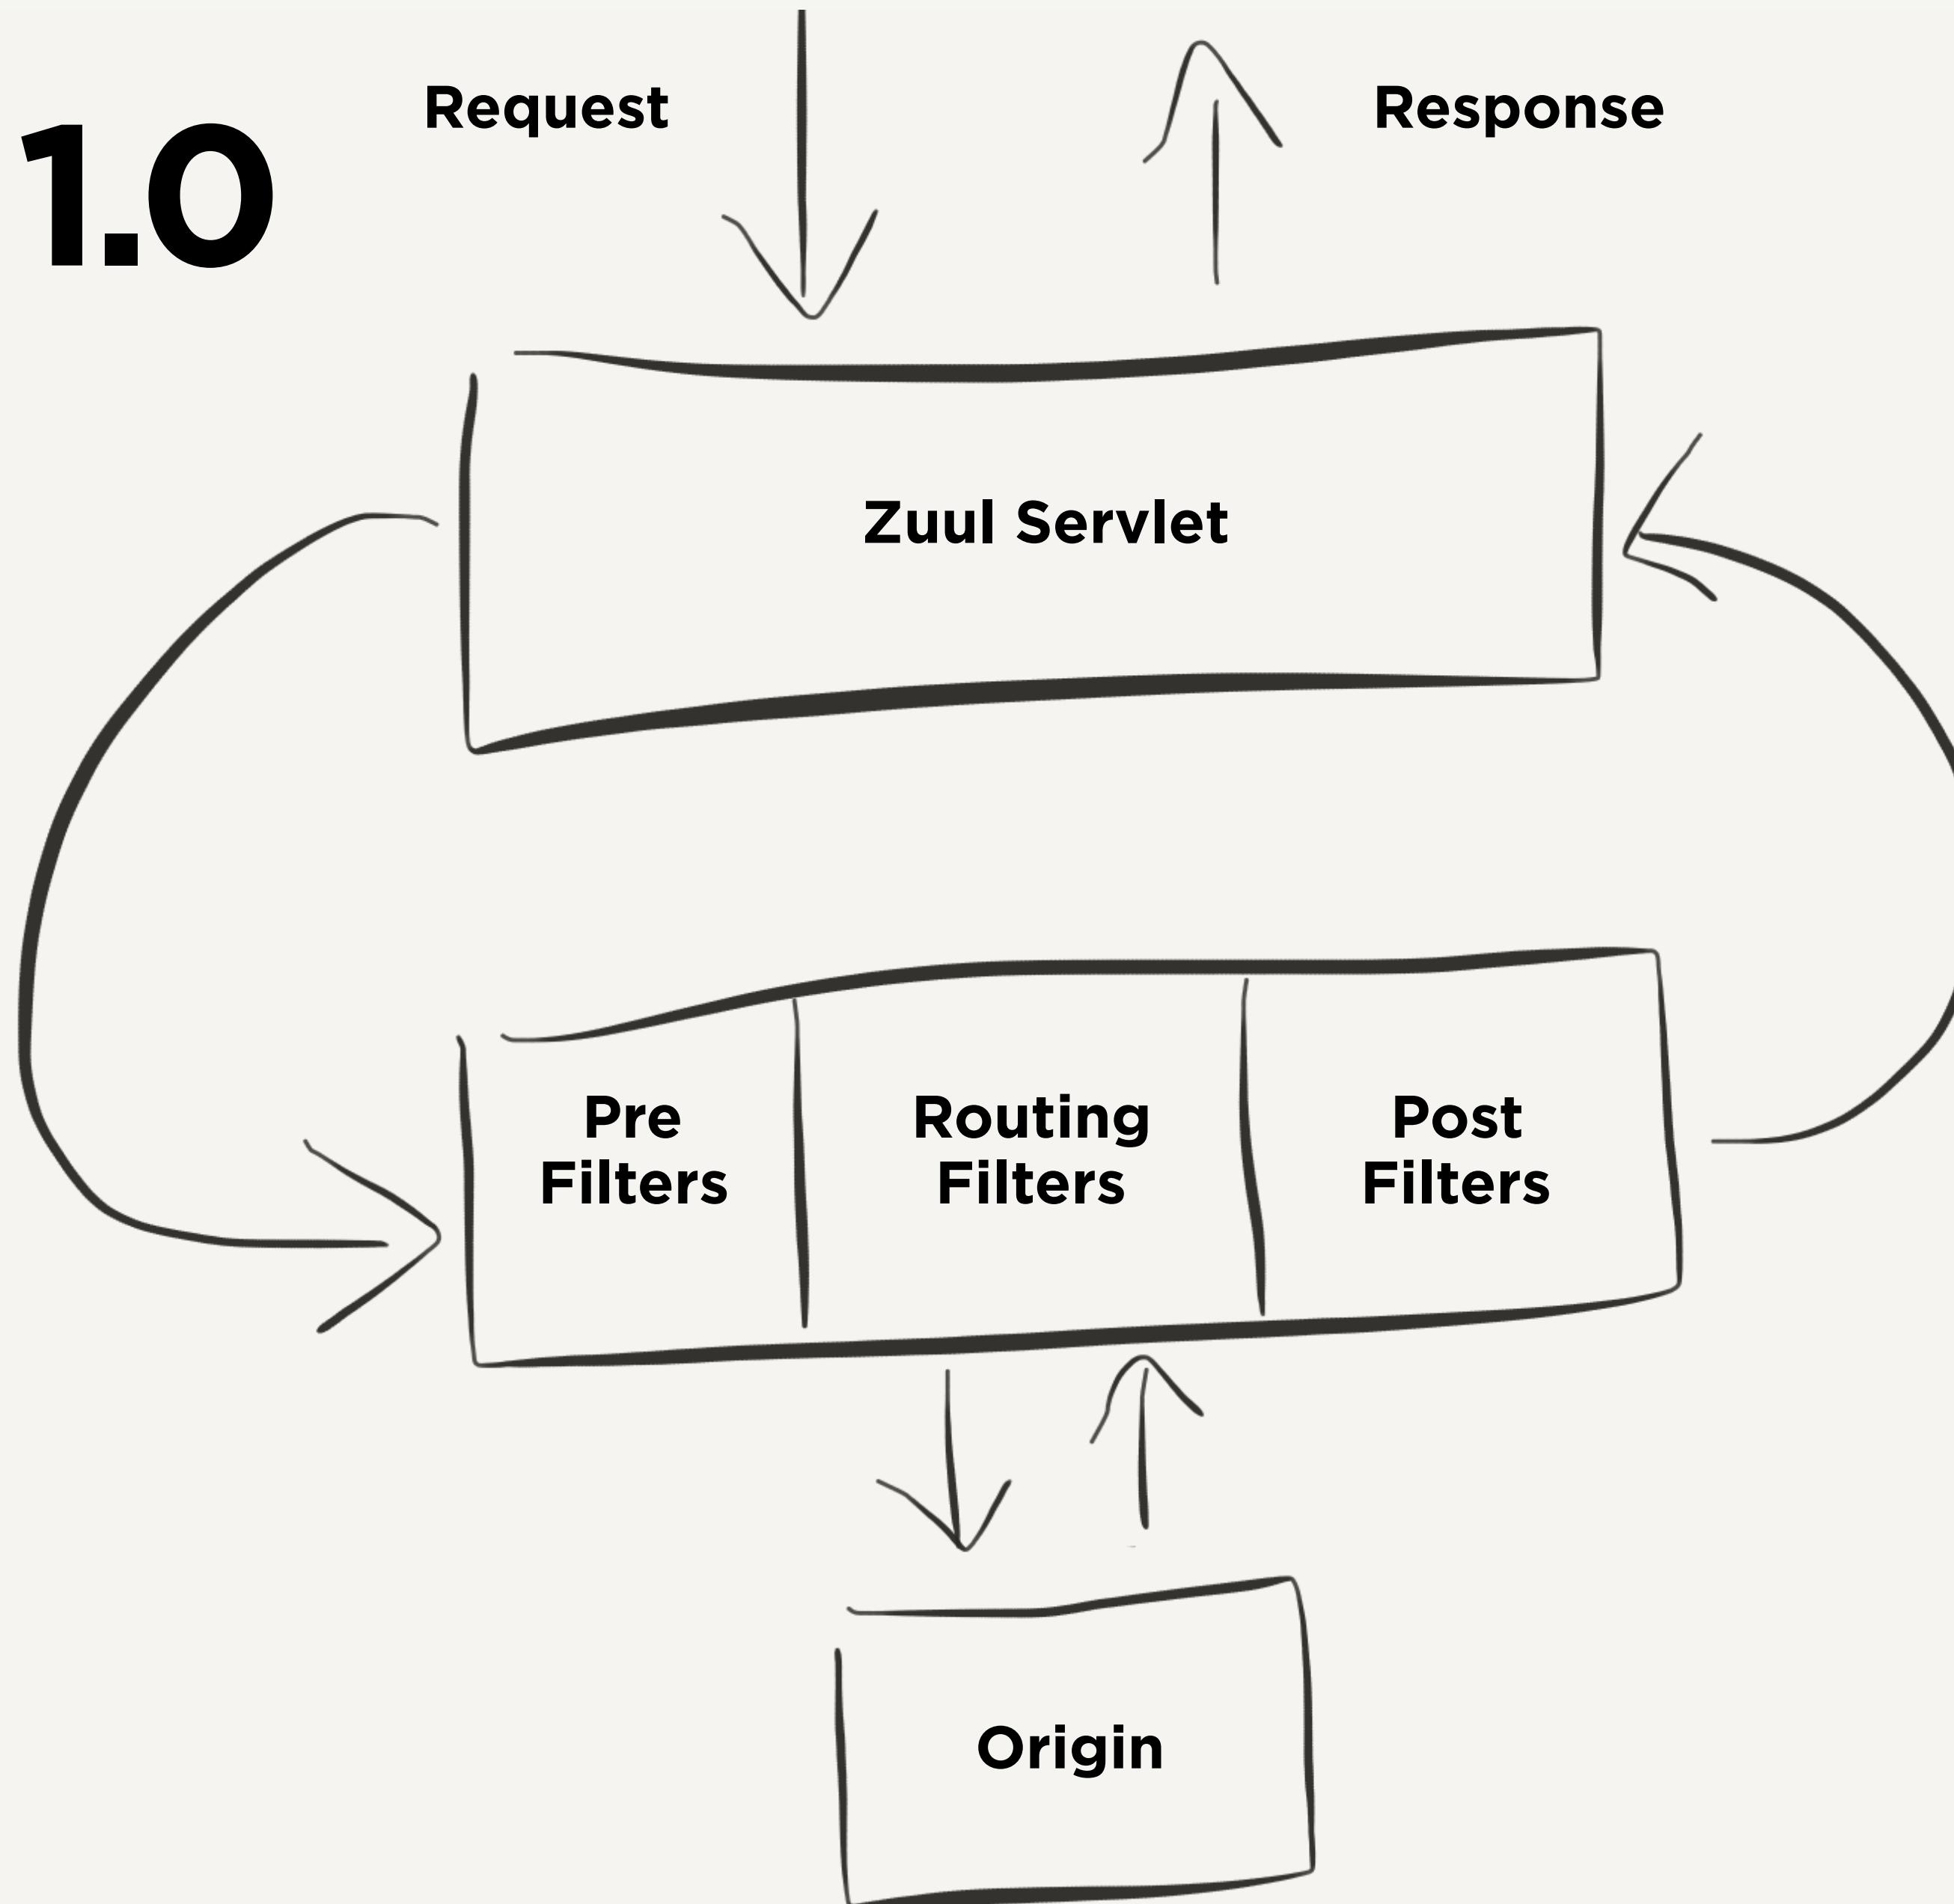
Proxy all traffic coming into AWS

Traffic from (almost) every country

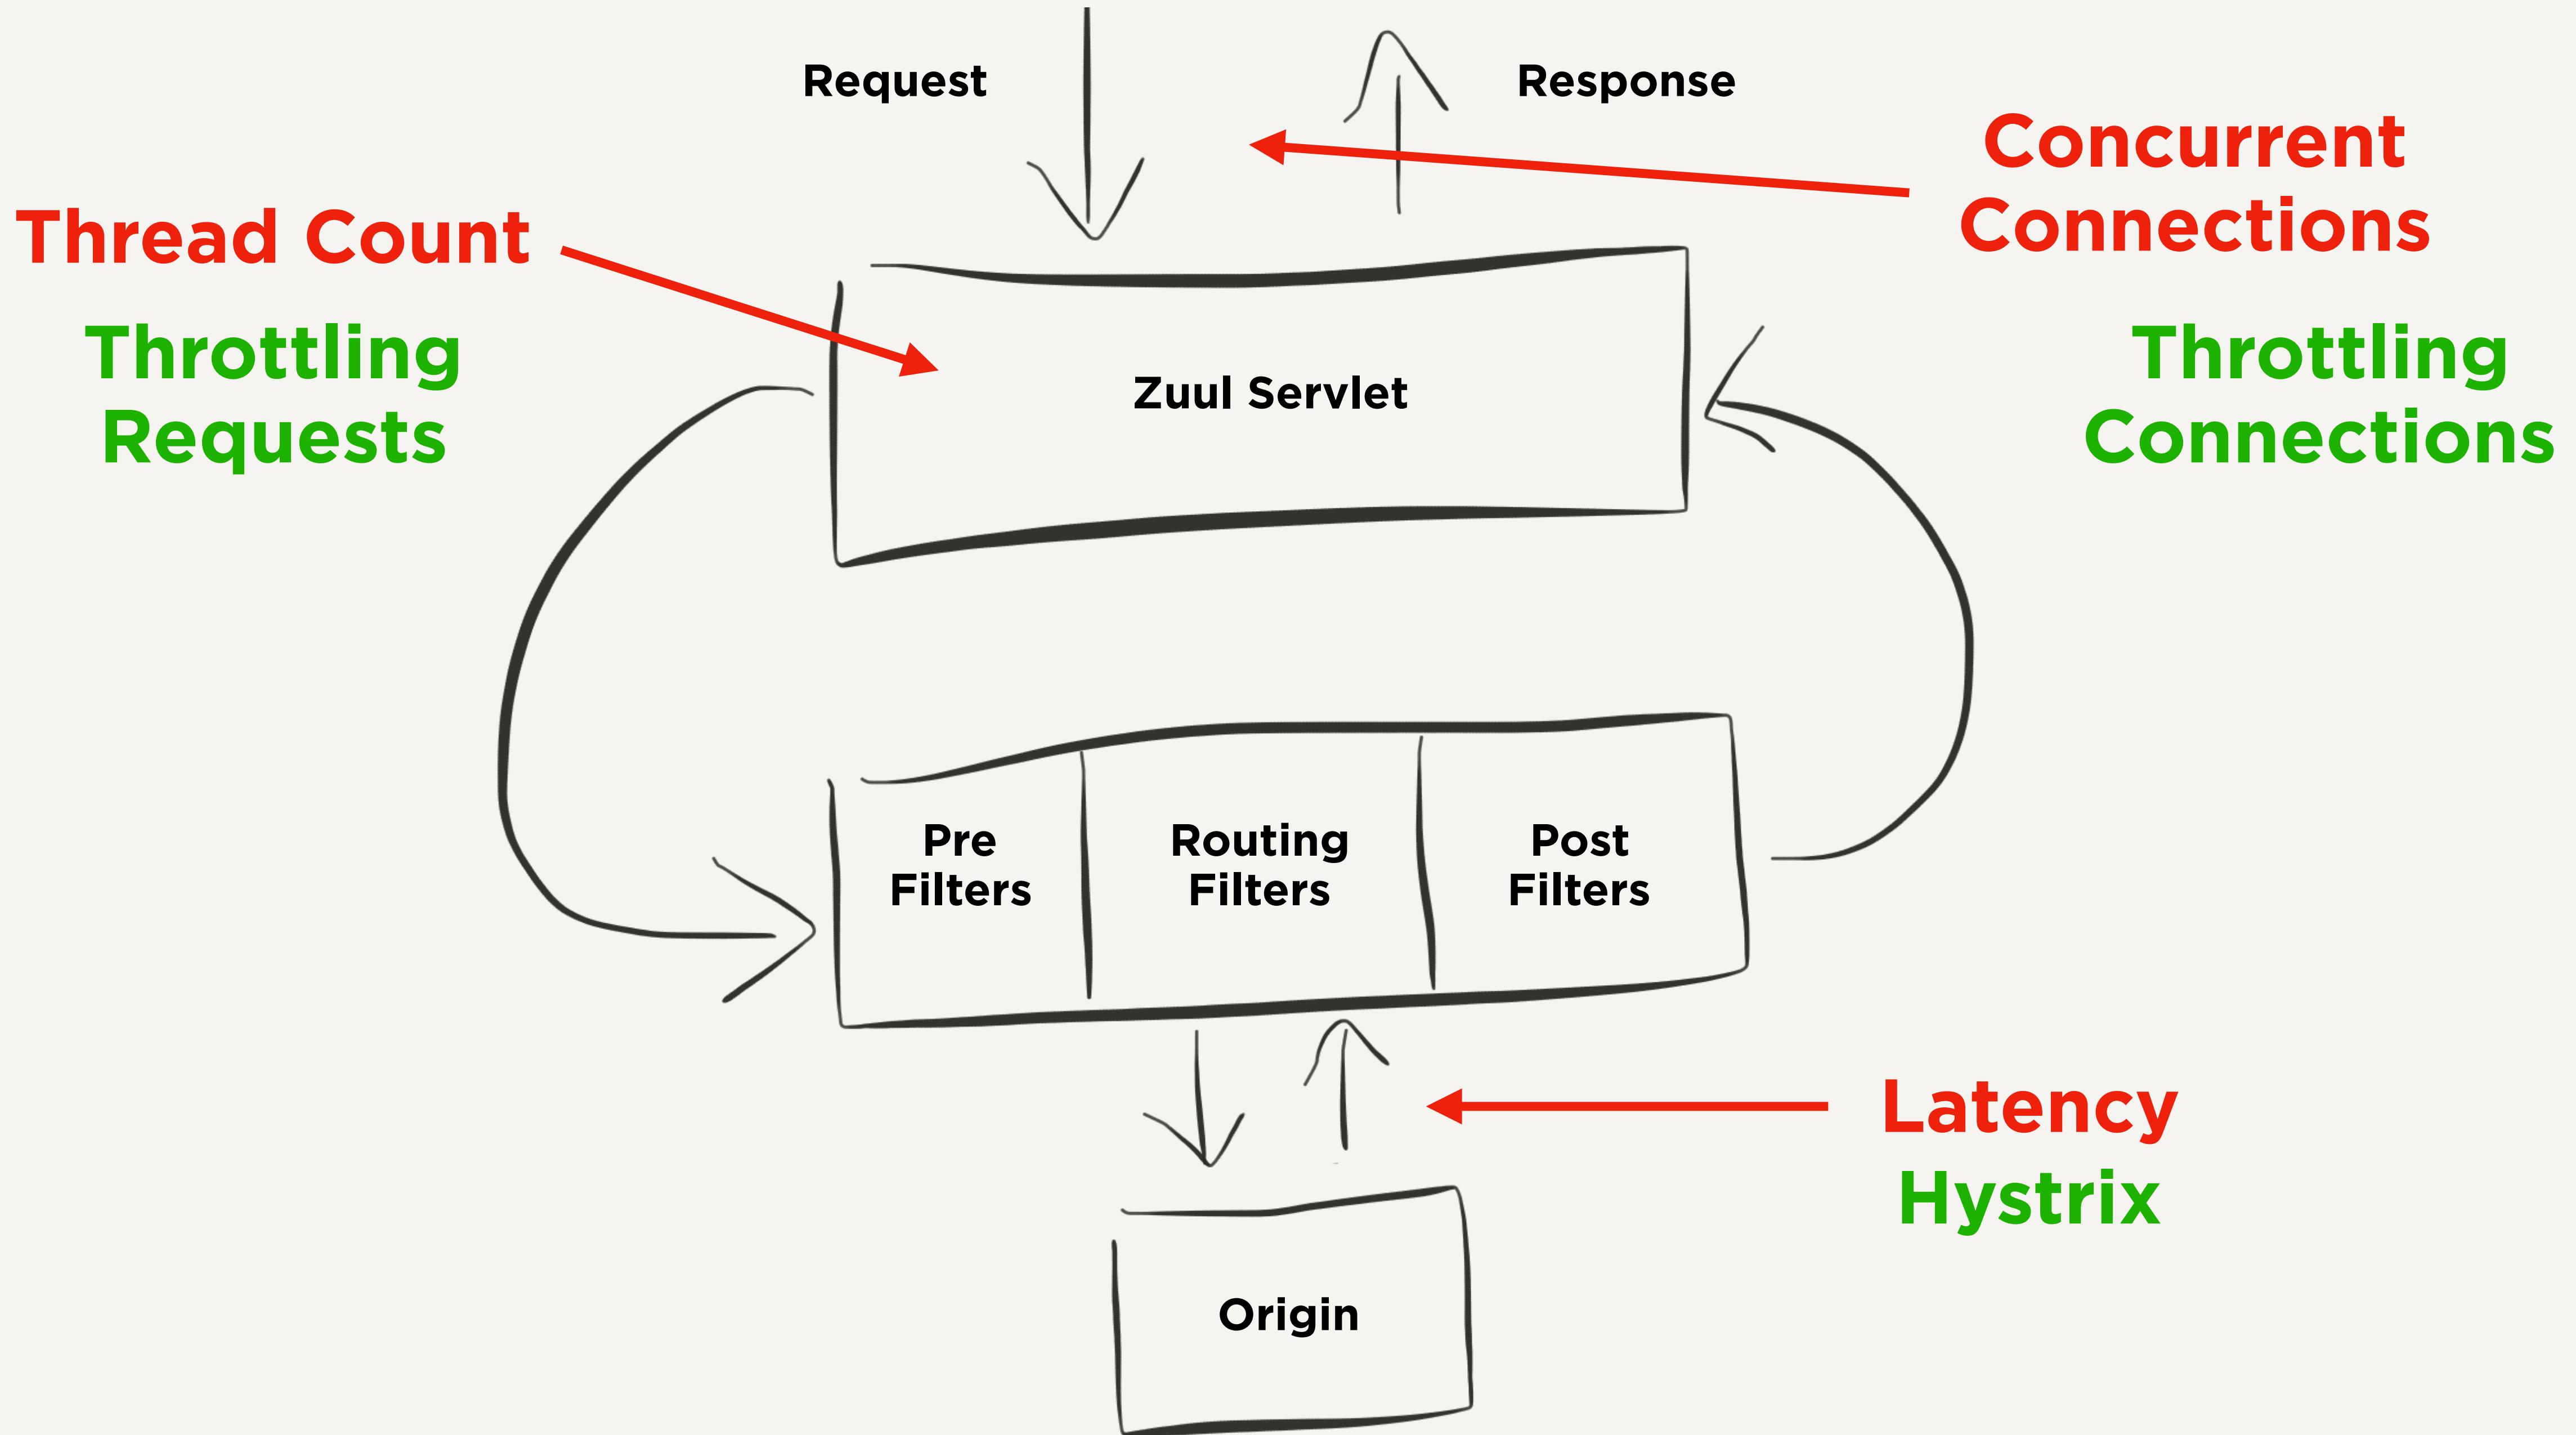
Tens of billions of requests per day

Hundreds of terabytes transferred per day

Zuul 1.0









Noel
@noellinnane

Follow



Netflix is down! Panic!!

1:06 PM - 1 Oct 2016



Bongo
@congobongo93

Follow



#Netflixdown per tre minuti e parte il panico

Translate from Italian

12:57 PM - 15 Jun 2017



P@ - Kevin
@MurmeltierS

Follow



Panik, Panik #netflixdown

1:58 PM - 15 Oct 2015



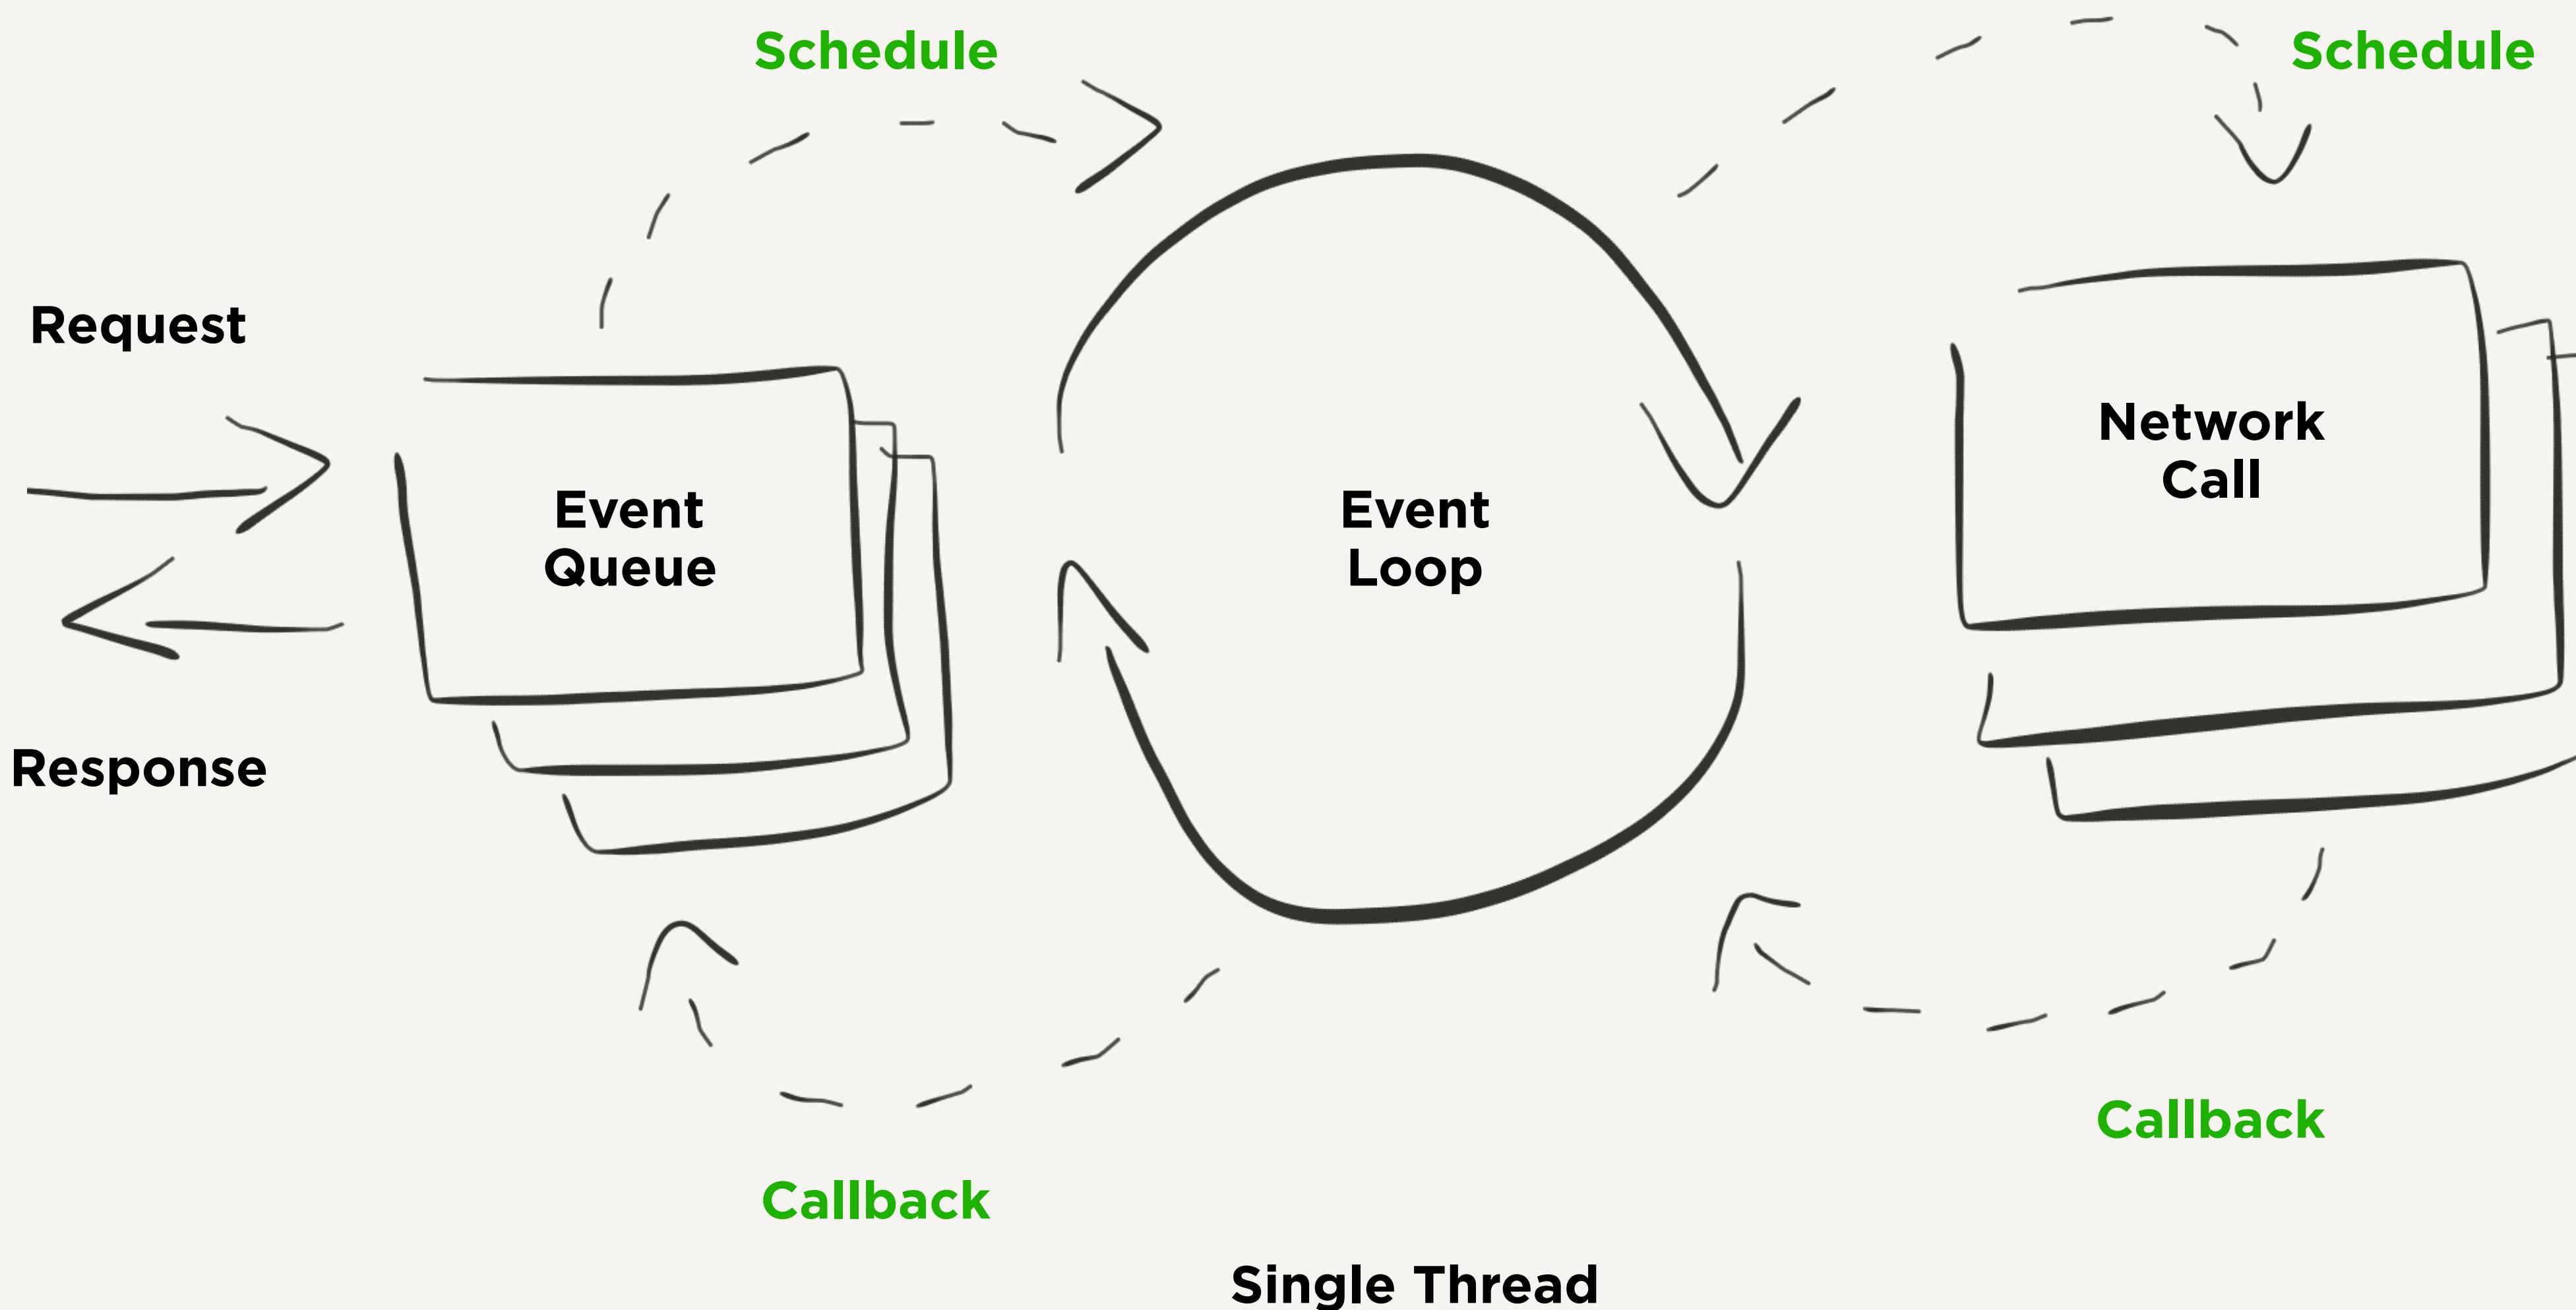
What is **Non-Blocking**?

Non-Blocking I/O Primitives

Asynchronous Programming

Event-loop Pattern

Event Loop



Why Non-Blocking?

Performance → CPU Utilization

Resiliency → Connection Scaling

Zuul 1.5

APR connector

Decouple connections and requests

New filter interfaces

Run the same filter logic on both systems

New **Filter** Interface

Sync:

```
0 apply(I message);
```

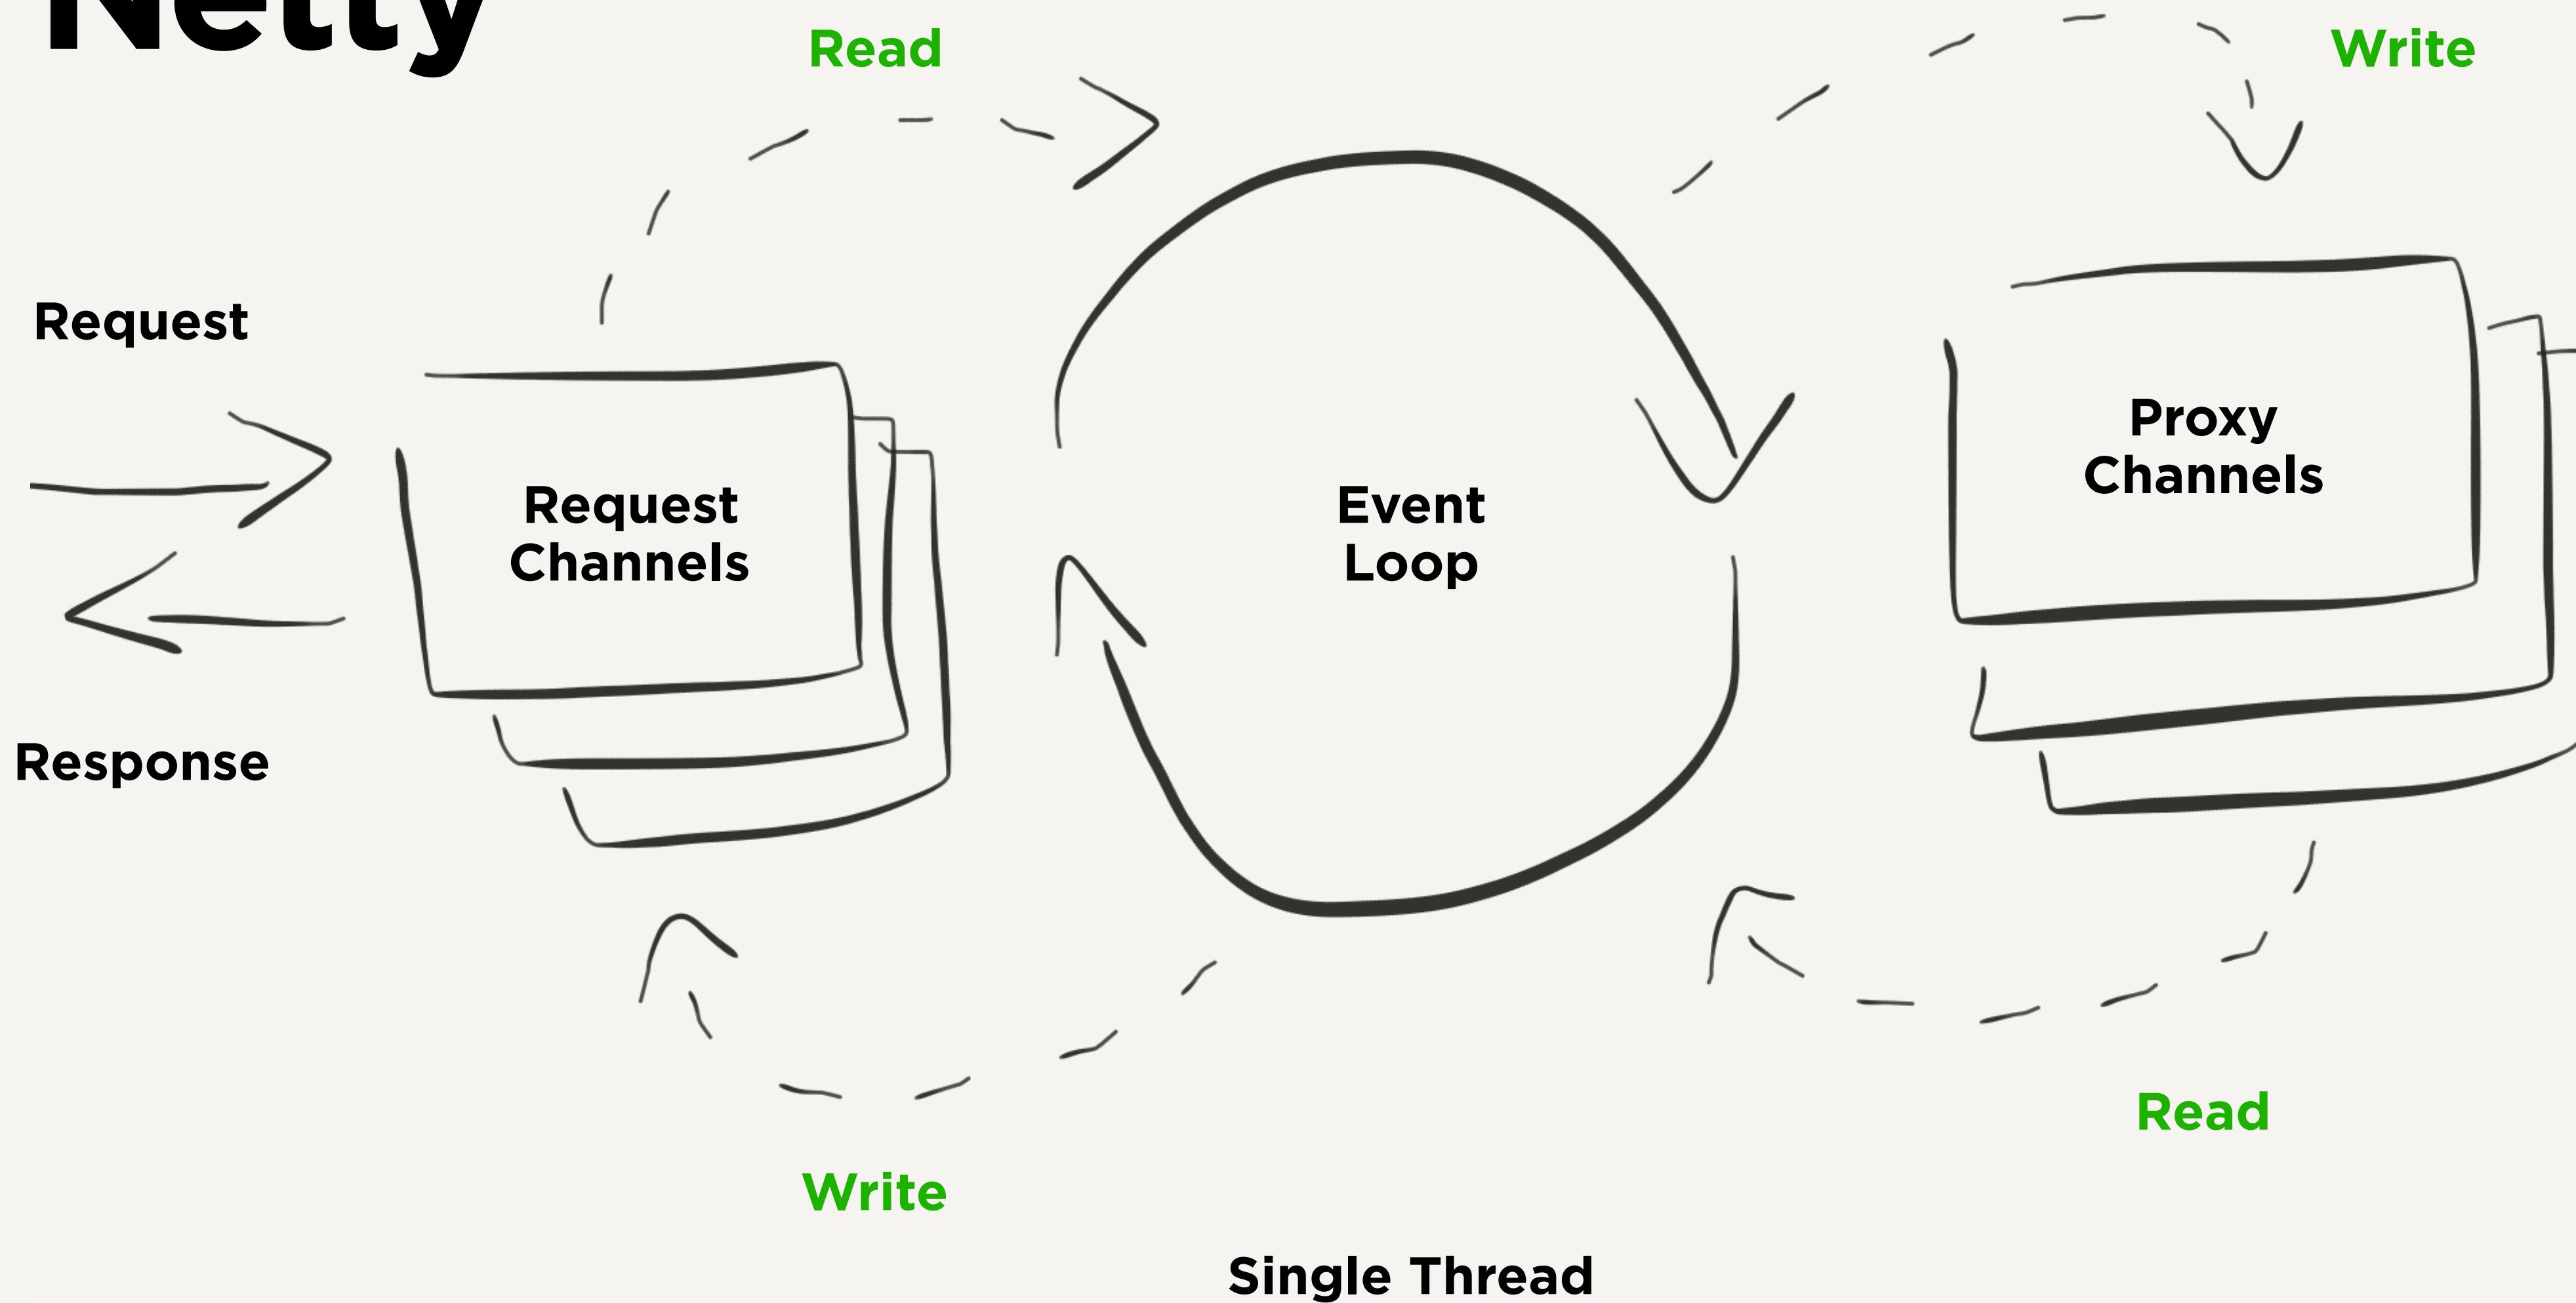
Async:

```
Observable<0> applyAsync(I message);
```

Observable Filter Chain

```
Observable.of(request)
    .flatMap(msg -> applyFilterChain(msg))
    .flatMap(msg -> writeResponse(msg))
    .doOnError(e -> writeErrorResponse(e))
    .doOnNext(msg -> postRequestMetrics(msg));
```

Netty



Building the Core

Adapters between Netty and RxJava

Adapters to wrap ByteBuf to Observable

Maintain the same functionality and metrics

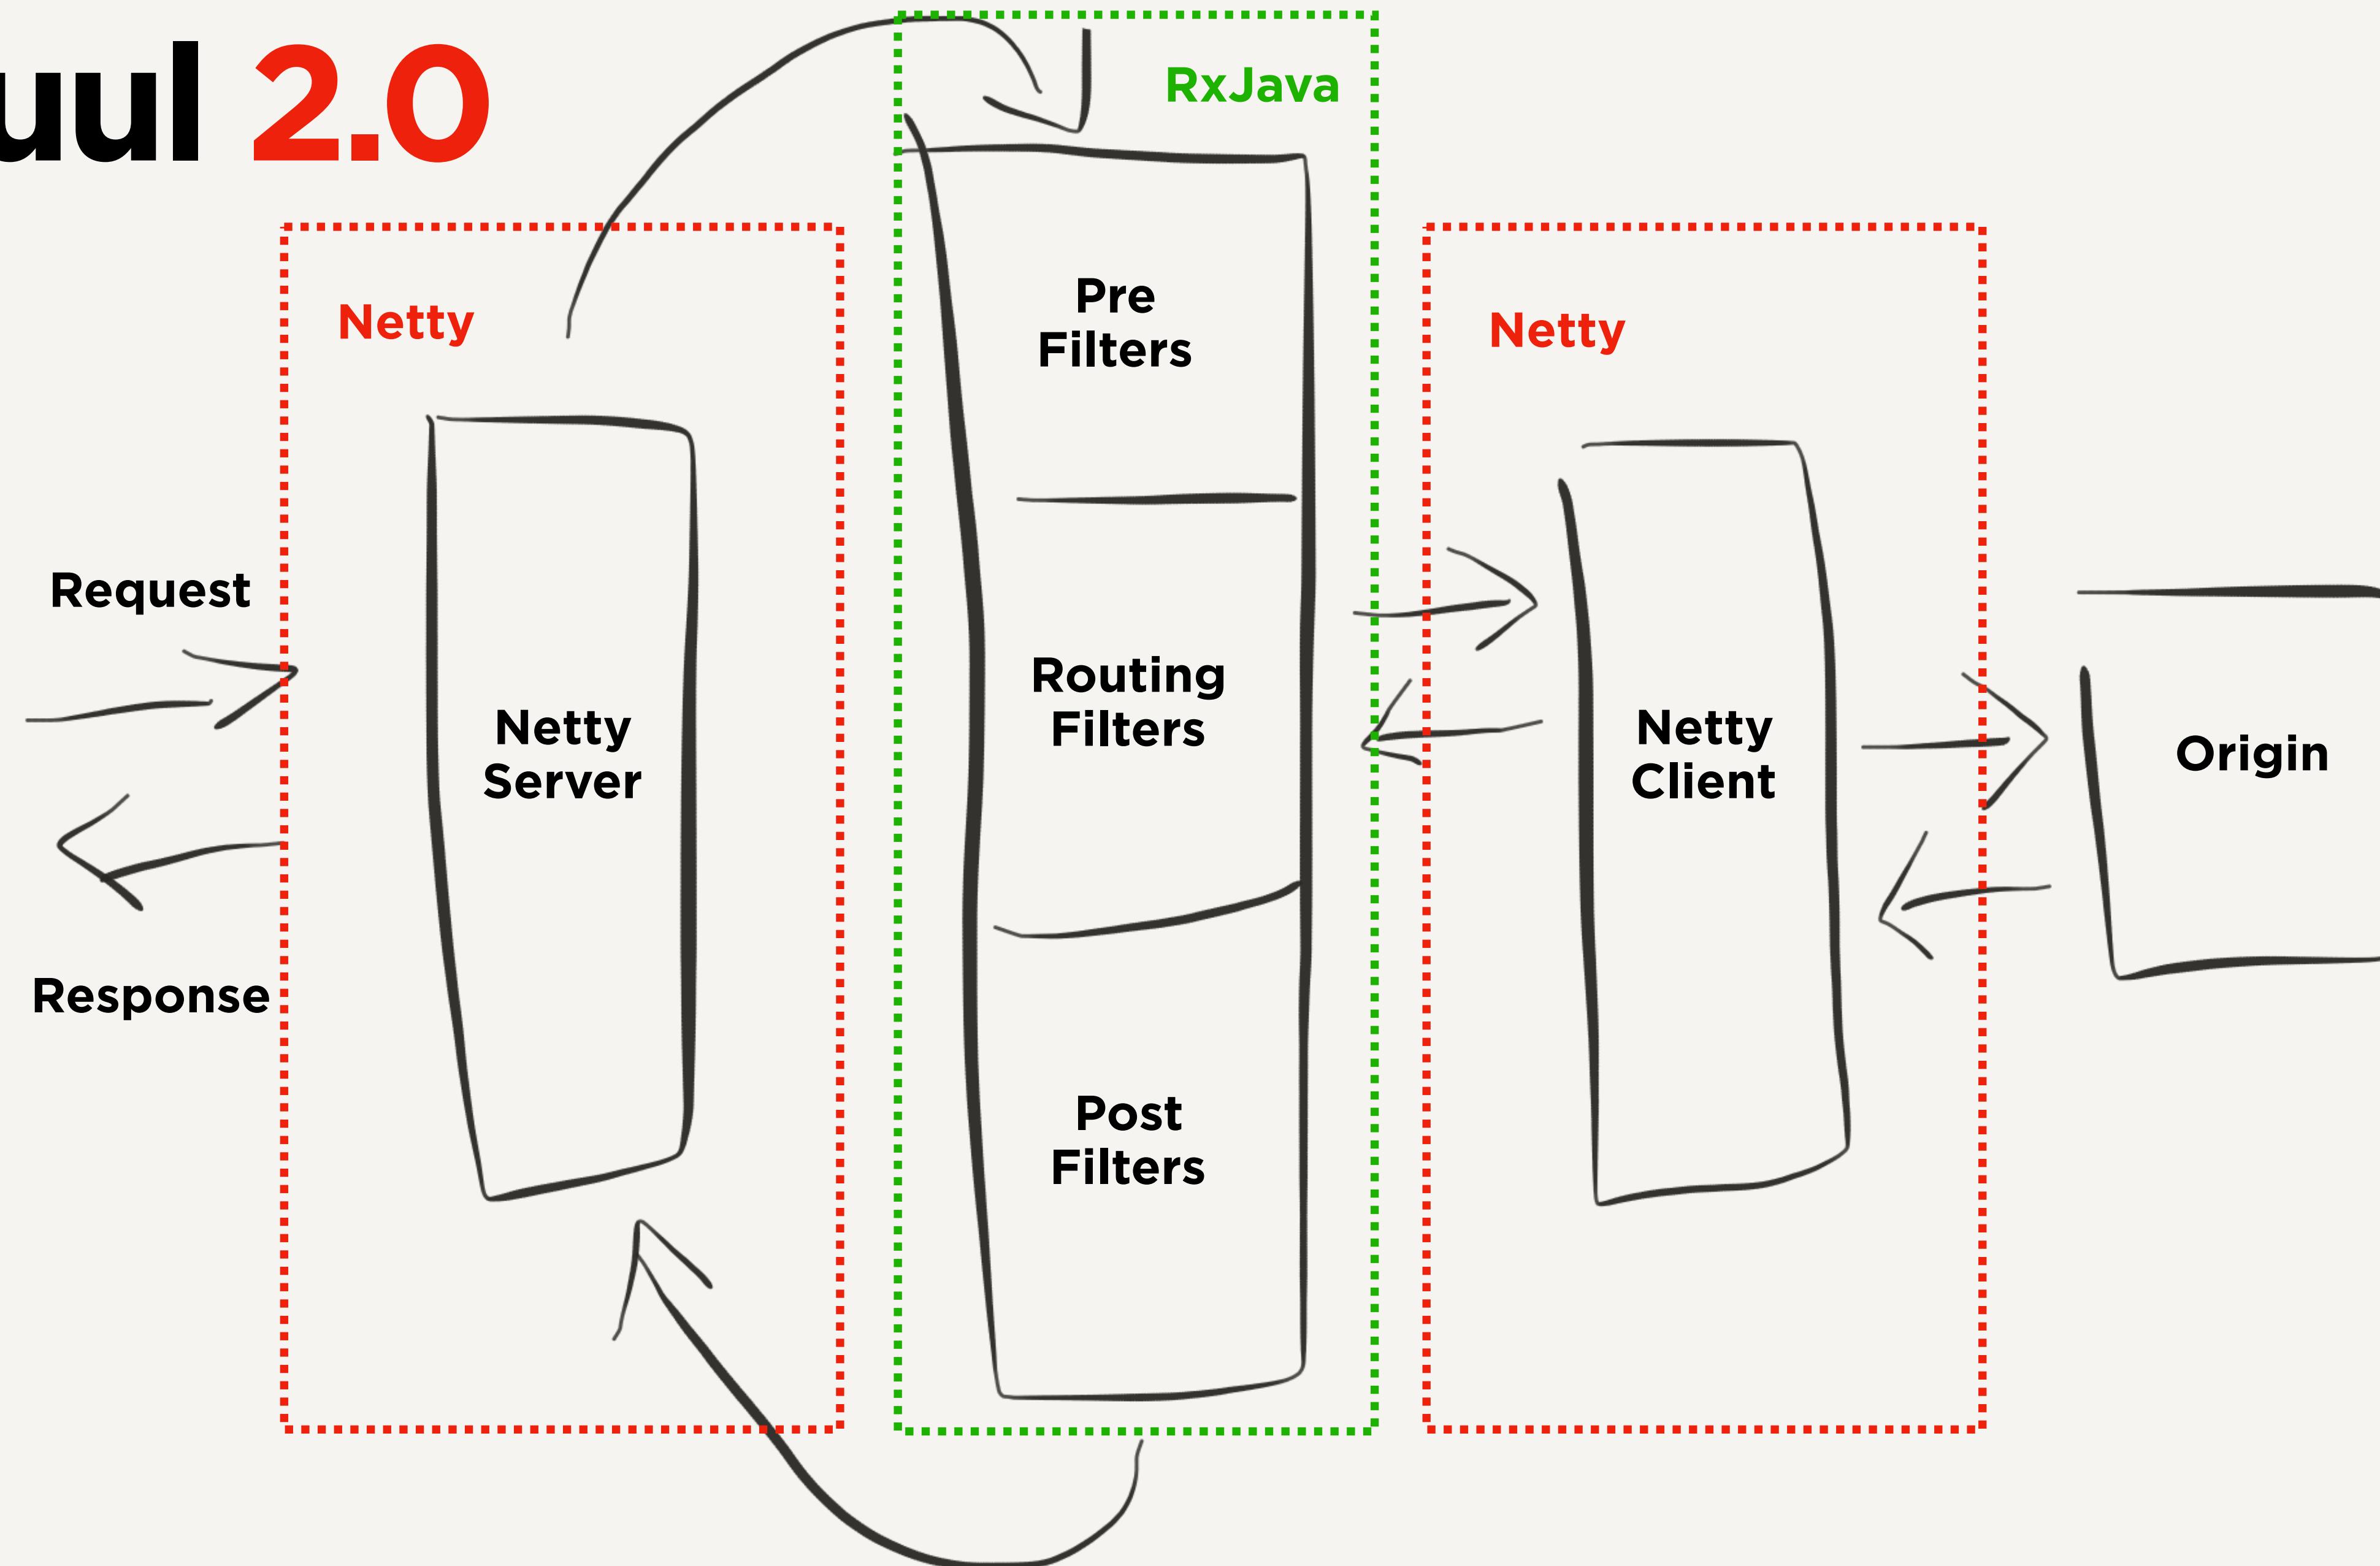
Netty Client

Remove blocking client

Write new Netty-based client

Maintain existing load balancing algorithms

Zuul 2.0



Initial Hurdles

Operating Zuul 2.0
in the Real World



Underlying Assumptions

Unexpected **blocking**

Libraries built with assumption of blocking

Reactive Audit to the rescue!

Reactive Audit Example

HIGH : Call method void java.lang.Thread.sleep(long)
com.netflix.reactive.audit.lib.CPUReactiveAuditException:

Call method void java.lang.Thread.sleep(long)

```
at thread "Salamander-ClientToZuulWorker-3"
at com.netflix.numerus.NumerusRollingNumber.getCurrentBucket(NumerusRollingNumber.java:330)
at com.netflix.numerus.NumerusRollingNumber.getRollingSum(NumerusRollingNumber.java:161)
at com.netflix.zuul.origins.InstrumentedOrigin.getErrorPercentage(InstrumentedOrigin.java:190)
at com.netflix.zuul.origins.OriginDeviceRetryThrottle.shouldThrottle(OriginDeviceRetryThrottle.java:45)
at com.netflix.zuul.origins.InstrumentedOrigin.request(InstrumentedOrigin.java:229)
at com.netflix.zuul.filters.endpoints.NfProxyEndpoint.lambda$applyAsync$1(NfProxyEndpoint.java:46)
at rx.internal.operators.OnSubscribeMap$MapSubscriber.onNext(OnSubscribeMap.java:69)
at rx.internal.operators.OnSubscribeMap$MapSubscriber.onNext(OnSubscribeMap.java:77)
at rx.internal.util.ScalarSynchronousObservable$WeakSingleProducer.request(ScalarSynchronousObservable.java:276)
at rx.Subscriber.setProducer(Subscriber.java:211)
at rx.internal.operators.OnSubscribeMap$MapSubscriber.setProducer(OnSubscribeMap.java:102)
at rx.internal.operators.OnSubscribeMap$MapSubscriber.setProducer(OnSubscribeMap.java:102)
at rx.internal.util.ScalarSynchronousObservable$JustOnSubscribe.call(ScalarSynchronousObservable.java:138)
at rx.internal.util.ScalarSynchronousObservable$JustOnSubscribe.call(ScalarSynchronousObservable.java:129)
at rx.Observable.unsafeSubscribe(Observable.java:10200)
at rx.internal.operators.OnSubscribeMap.call(OnSubscribeMap.java:48)
at rx.internal.operators.OnSubscribeMap.call(OnSubscribeMap.java:33)
at rx.Observable.unsafeSubscribe(Observable.java:10200)
at rx.internal.operators.OnSubscribeMap.call(OnSubscribeMap.java:48)
```

More Underlying Assumptions

Libraries rely on thread locals

Shared state between libs in the thread

Time-consuming, **manual** inspection

Debugging Async Requests

Tracing request failures

Comprehending stack traces

Timing asynchronous requests

Request Passport

Thanks Zach Tellman and Strange Loop!

Maintain "passport" on the channel

Add entries with **state** and **timestamp**

Make passports available on failure

Passport Example

+0=IN_REQ_HEADERS RECEIVED,
+705017=FILTERS_INBOUND_START,
+1949958=MISC_IO_START,
+2027915=IN_REQ_LAST_CONTENT RECEIVED,
+3391350=MISC_IO_STOP,
+4767202=ORIGIN_CONN_ACQUIRE_START,
+4784096=CLIENT_CH_CONNECTING,
+4891510=FILTERS_INBOUND_END,
+5748769=CLIENT_CH_CONNECTED,
+5773298=ORIGIN_CONN_ACQUIRE_END,
+6336479=OUT_REQ_HEADERS_SENDING,
+6388730=OUT_REQ_HEADERS_SENT,
+6402264=OUT_REQ_LAST_CONTENT_SENDING,
+6408016=OUT_REQ_LAST_CONTENT_SENT,
+6416189=CLIENT_CH_ACTIVE,
+45006120586=CLIENT_CH_CLOSE,
+45006221594=CLIENT_CH_READ_TIMEOUT,

+45006359072=ORIGIN_CONN_ACQUIRE_START,
+45006385636=CLIENT_CH_CONNECTING,
+45006492694=CLIENT_CH_CLOSE,
+45006907204=CLIENT_CH_CONNECTED,
+45006928564=ORIGIN_CONN_ACQUIRE_END,
+45007638281=OUT_REQ_HEADERS_SENDING,
+45007690499=OUT_REQ_HEADERS_SENT,
+45007708575=OUT_REQ_LAST_CONTENT_SENDING,
+45007716697=OUT_REQ_LAST_CONTENT_SENT,
+45007725840=CLIENT_CH_ACTIVE,
+59920094115=SERVER_CH_INACTIVE,
+59920182824=IN_REQ_CANCELLED,
+90007646473=CLIENT_CH_CLOSE,
+90007777157=CLIENT_CH_READ_TIMEOUT,
+90007965586=FILTERS_OUTBOUND_START,
+90008909963=FILTERS_OUTBOUND_END,
+90009003762=NOW

Long Tail Problems

Operating Zuul 2.0
in the Real World



Direct Memory Leaks

Servers slowly use more memory

Eventually RPS drops to zero

Extremely difficult to debug

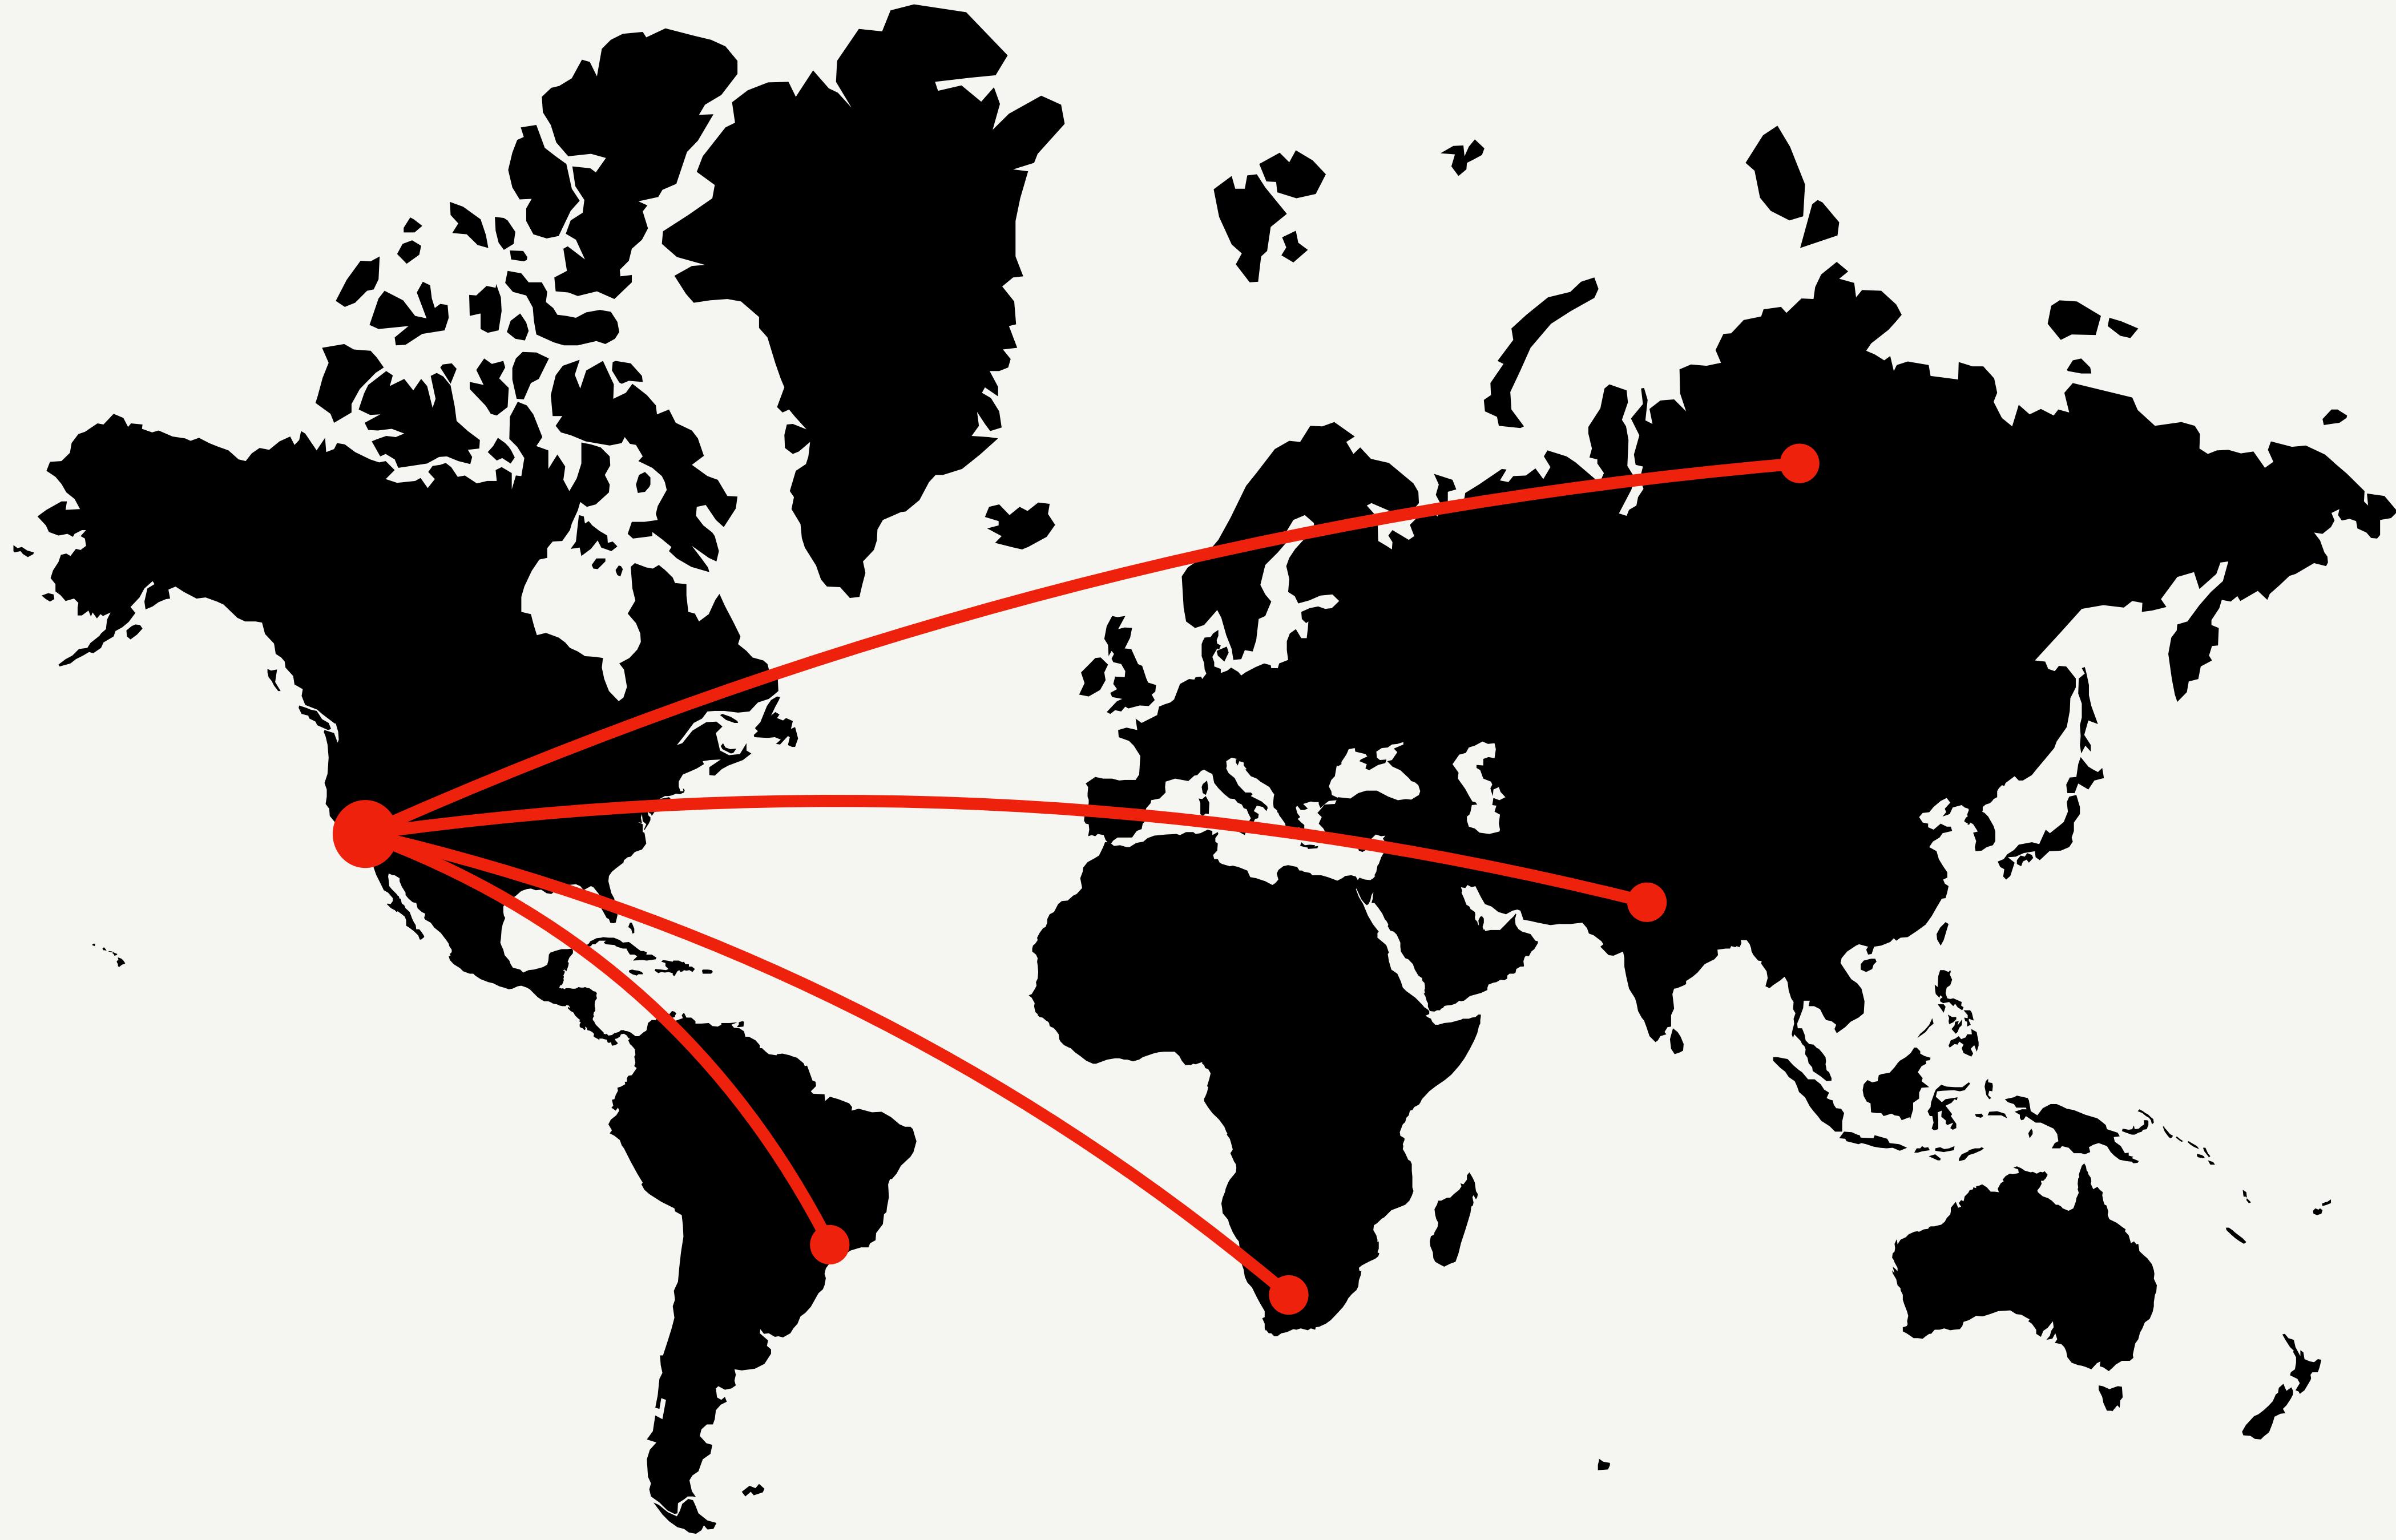
What's going on here?

Analyzing async code is **hard**

Throw in RxJava for even **more** fun

Stack traces become useless

```
java.lang.IllegalStateException: Demo exception.
at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
at sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:62)
at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
at org.codehaus.groovy.reflection.CachedConstructor.invoke(CachedConstructor.java:80)
at org.codehaus.groovy.runtime.callsite.ConstructorSite$ConstructorSiteNoUnwrapNoCoerce.callConstructor(ConstructorSite.java:105)
at org.codehaus.groovy.runtime.callsite.CallSiteArray.defaultCallConstructor(CallSiteArray.java:60)
at org.codehaus.groovy.runtime.callsite.AbstractCallSite.callConstructor(AbstractCallSite.java:235)
at org.codehaus.groovy.runtime.callsite.AbstractCallSite.callConstructor(AbstractCallSite.java:247)
at global.ThreadChannelMismatchTrackerOutbound.apply(ThreadChannelMismatchTrackerOutbound.groovy:23)
at global.ThreadChannelMismatchTrackerOutbound.apply(ThreadChannelMismatchTrackerOutbound.groovy)
at com.netflix.zuul.FilterProcessorImpl.processSyncFilter(FilterProcessorImpl.java:388)
at com.netflix.zuul.FilterProcessorImpl.chainFilters(FilterProcessorImpl.java:122)
at com.netflix.zuul.FilterProcessorImpl.lambda$chainFilters$1(FilterProcessorImpl.java:133)
at rx.internal.operators.OnSubscribeMap$MapSubscriber.onNext(OnSubscribeMap.java:69)
at rx.internal.operators.OnSubscribeDoOnEach$DoOnEachSubscriber.onNext(OnSubscribeDoOnEach.java:101)
at rx.internal.operators.OnSubscribeMap$MapSubscriber.onNext(OnSubscribeMap.java:77)
at rx.internal.operators.OperatorOnErrorResumeNextViaFunction$4.onNext(OperatorOnErrorResumeNextViaFunction.java:154)
at rx.internal.util.ScalarSynchronousObservable$WeakSingleProducer.request(ScalarSynchronousObservable.java:276)
at rx.internal.producers.ProducerArbiter.setProducer(ProducerArbiter.java:126)
at rx.internal.operators.OperatorOnErrorResumeNextViaFunction$4.setProducer(OperatorOnErrorResumeNextViaFunction.java:159)
at rx.internal.util.ScalarSynchronousObservable$JustOnSubscribe.call(ScalarSynchronousObservable.java:138)
at rx.internal.util.ScalarSynchronousObservable$JustOnSubscribe.call(ScalarSynchronousObservable.java:129)
at rx.internal.operators.OnSubscribeLift.call(OnSubscribeLift.java:48)
at rx.internal.operators.OnSubscribeLift.call(OnSubscribeLift.java:30)
at rx.Observable.unsafeSubscribe/Observable.java:10256)
at rx.internal.operators.OnSubscribeMap.call(OnSubscribeMap.java:48)
at rx.internal.operators.OnSubscribeMap.call(OnSubscribeMap.java:33)
at rx.Observable.unsafeSubscribe/Observable.java:10256)
at rx.internal.operators.OnSubscribeDoOnEach.call(OnSubscribeDoOnEach.java:41)
at rx.internal.operators.OnSubscribeDoOnEach.call(OnSubscribeDoOnEach.java:30)
at rx.Observable.unsafeSubscribe/Observable.java:10256)
at rx.internal.operators.OnSubscribeMap.call(OnSubscribeMap.java:48)
at rx.internal.operators.OnSubscribeMap.call(OnSubscribeMap.java:33)
at rx.internal.operators.OnSubscribeLift.call(OnSubscribeLift.java:48)
at rx.internal.operators.OnSubscribeLift.call(OnSubscribeLift.java:30)
```



Geo Shuffle

CDN team gets alerted to location issues
We're sending people across the globe

Thread Locals! Argh!

Reset thread locals when jumping threads

Long Term Benefits

Operating Zuul 2.0
in the Real World

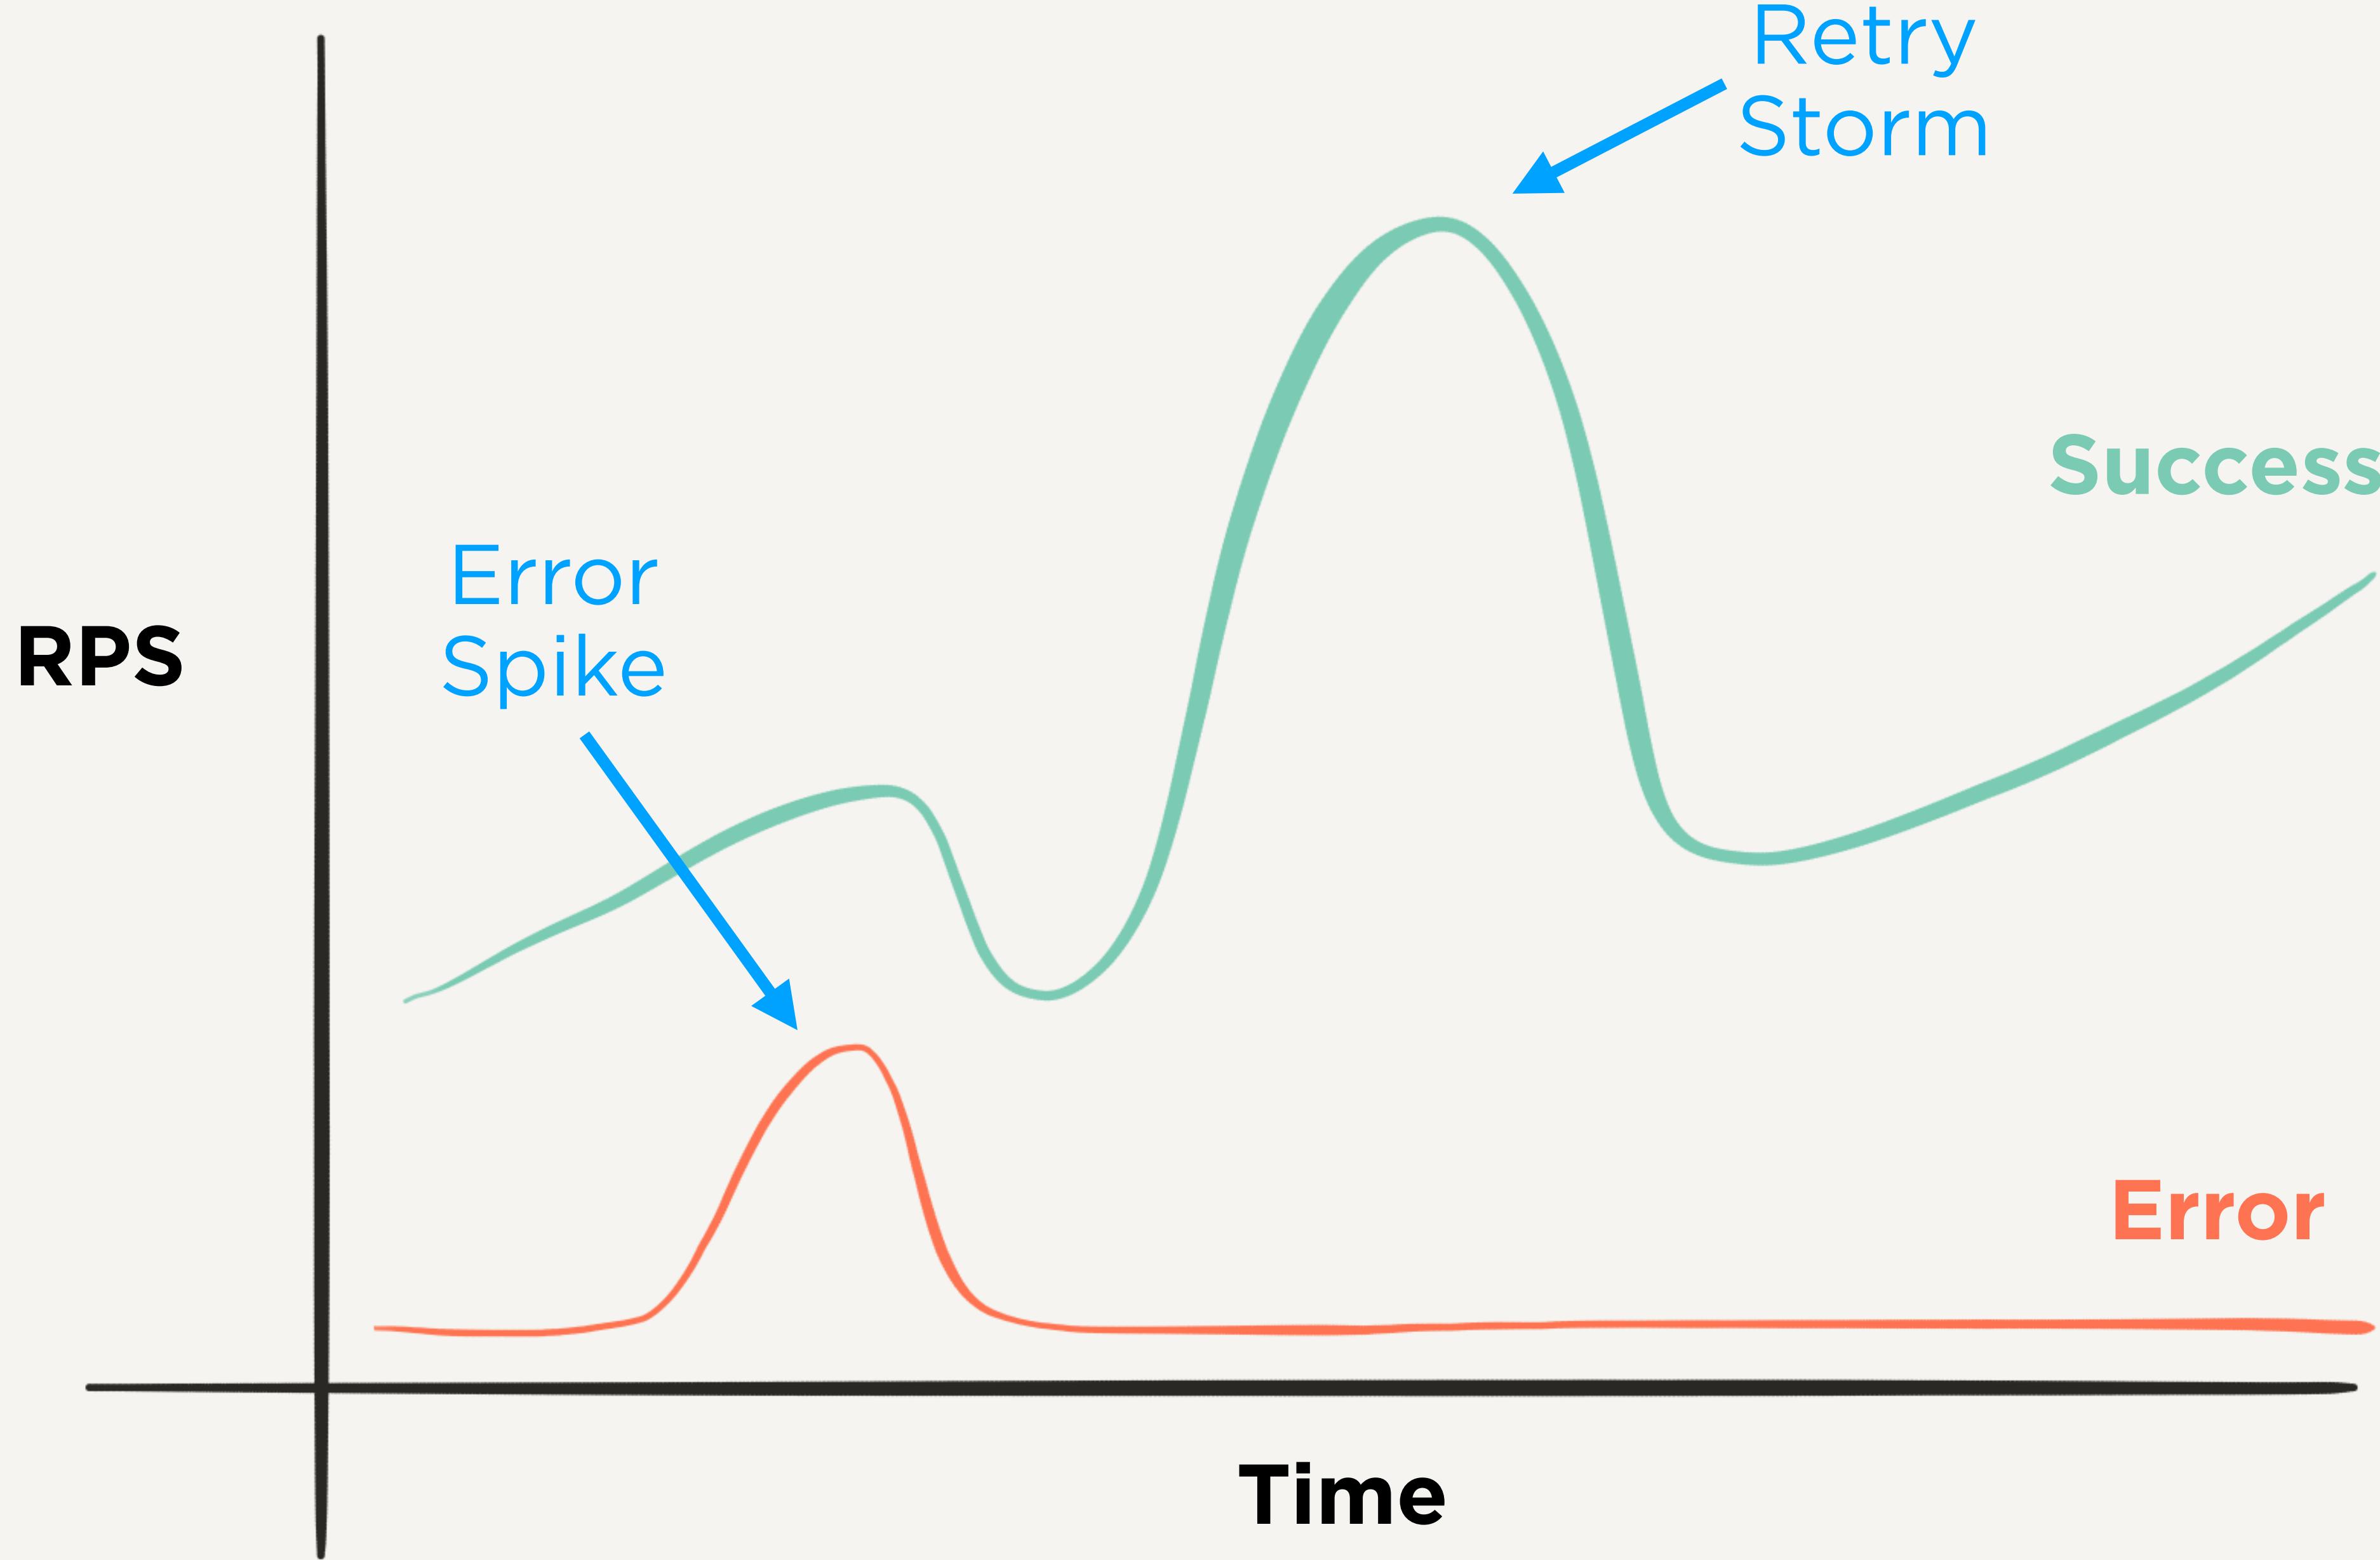


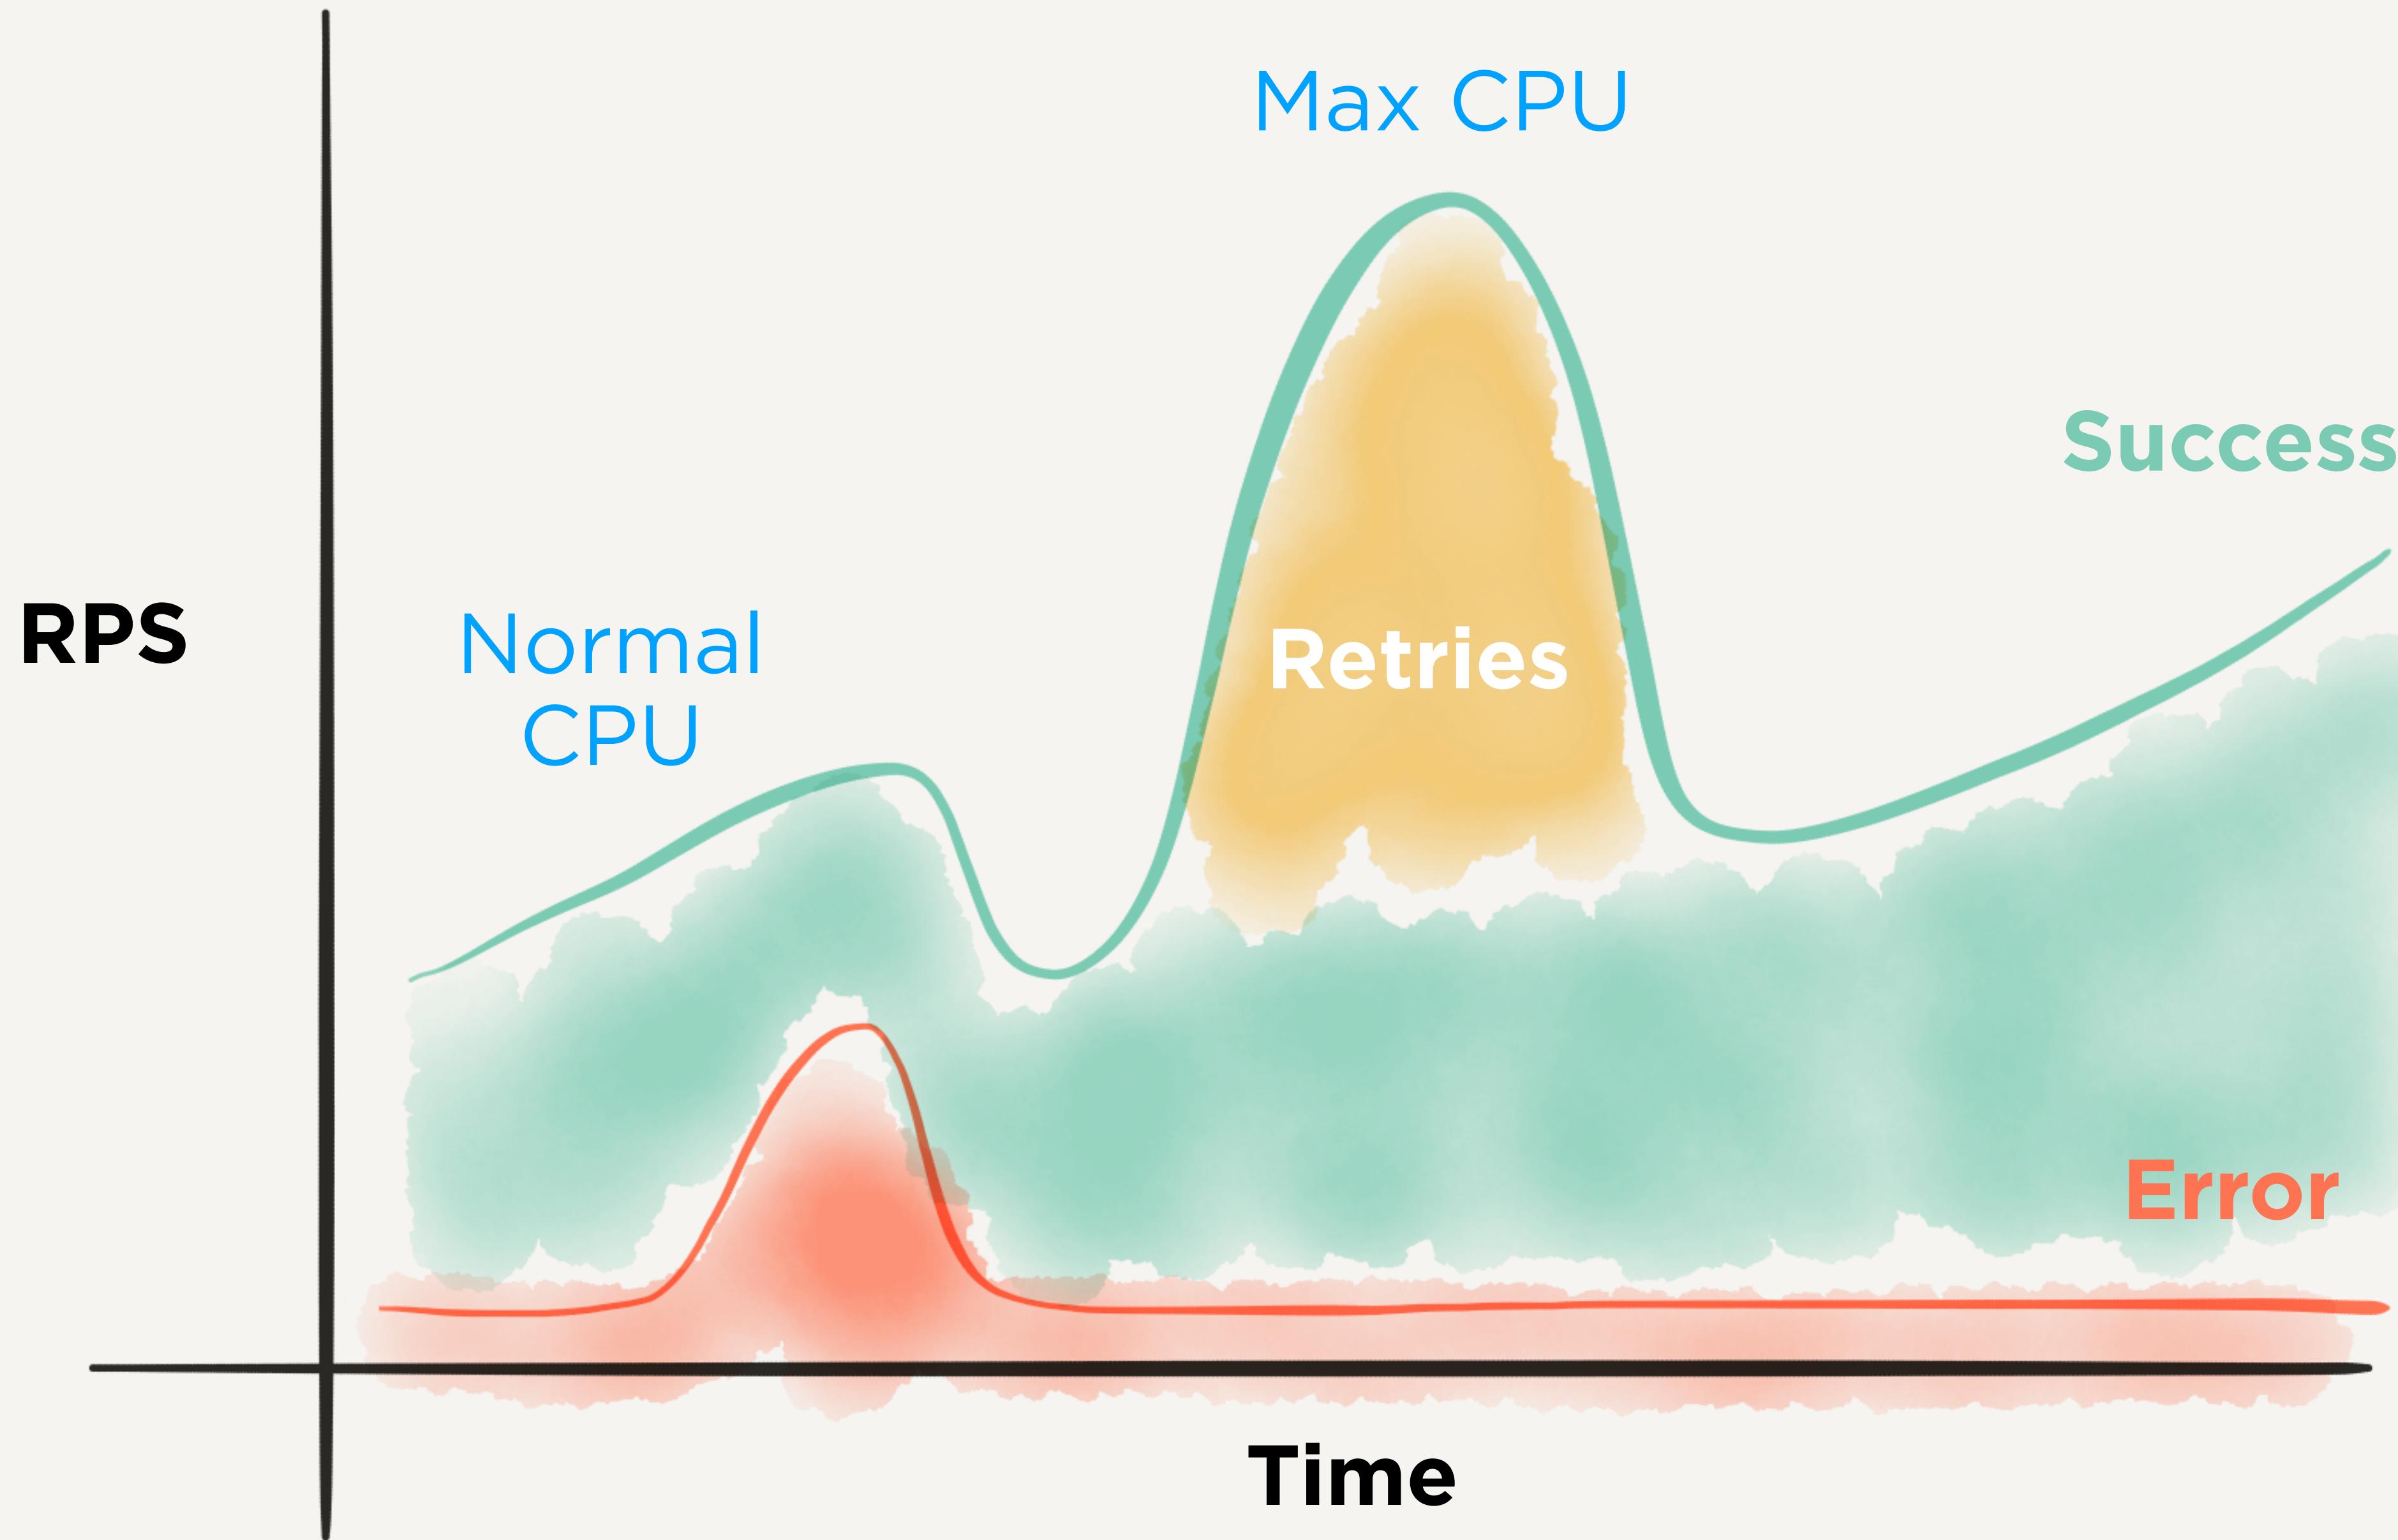
Weathering the Storm

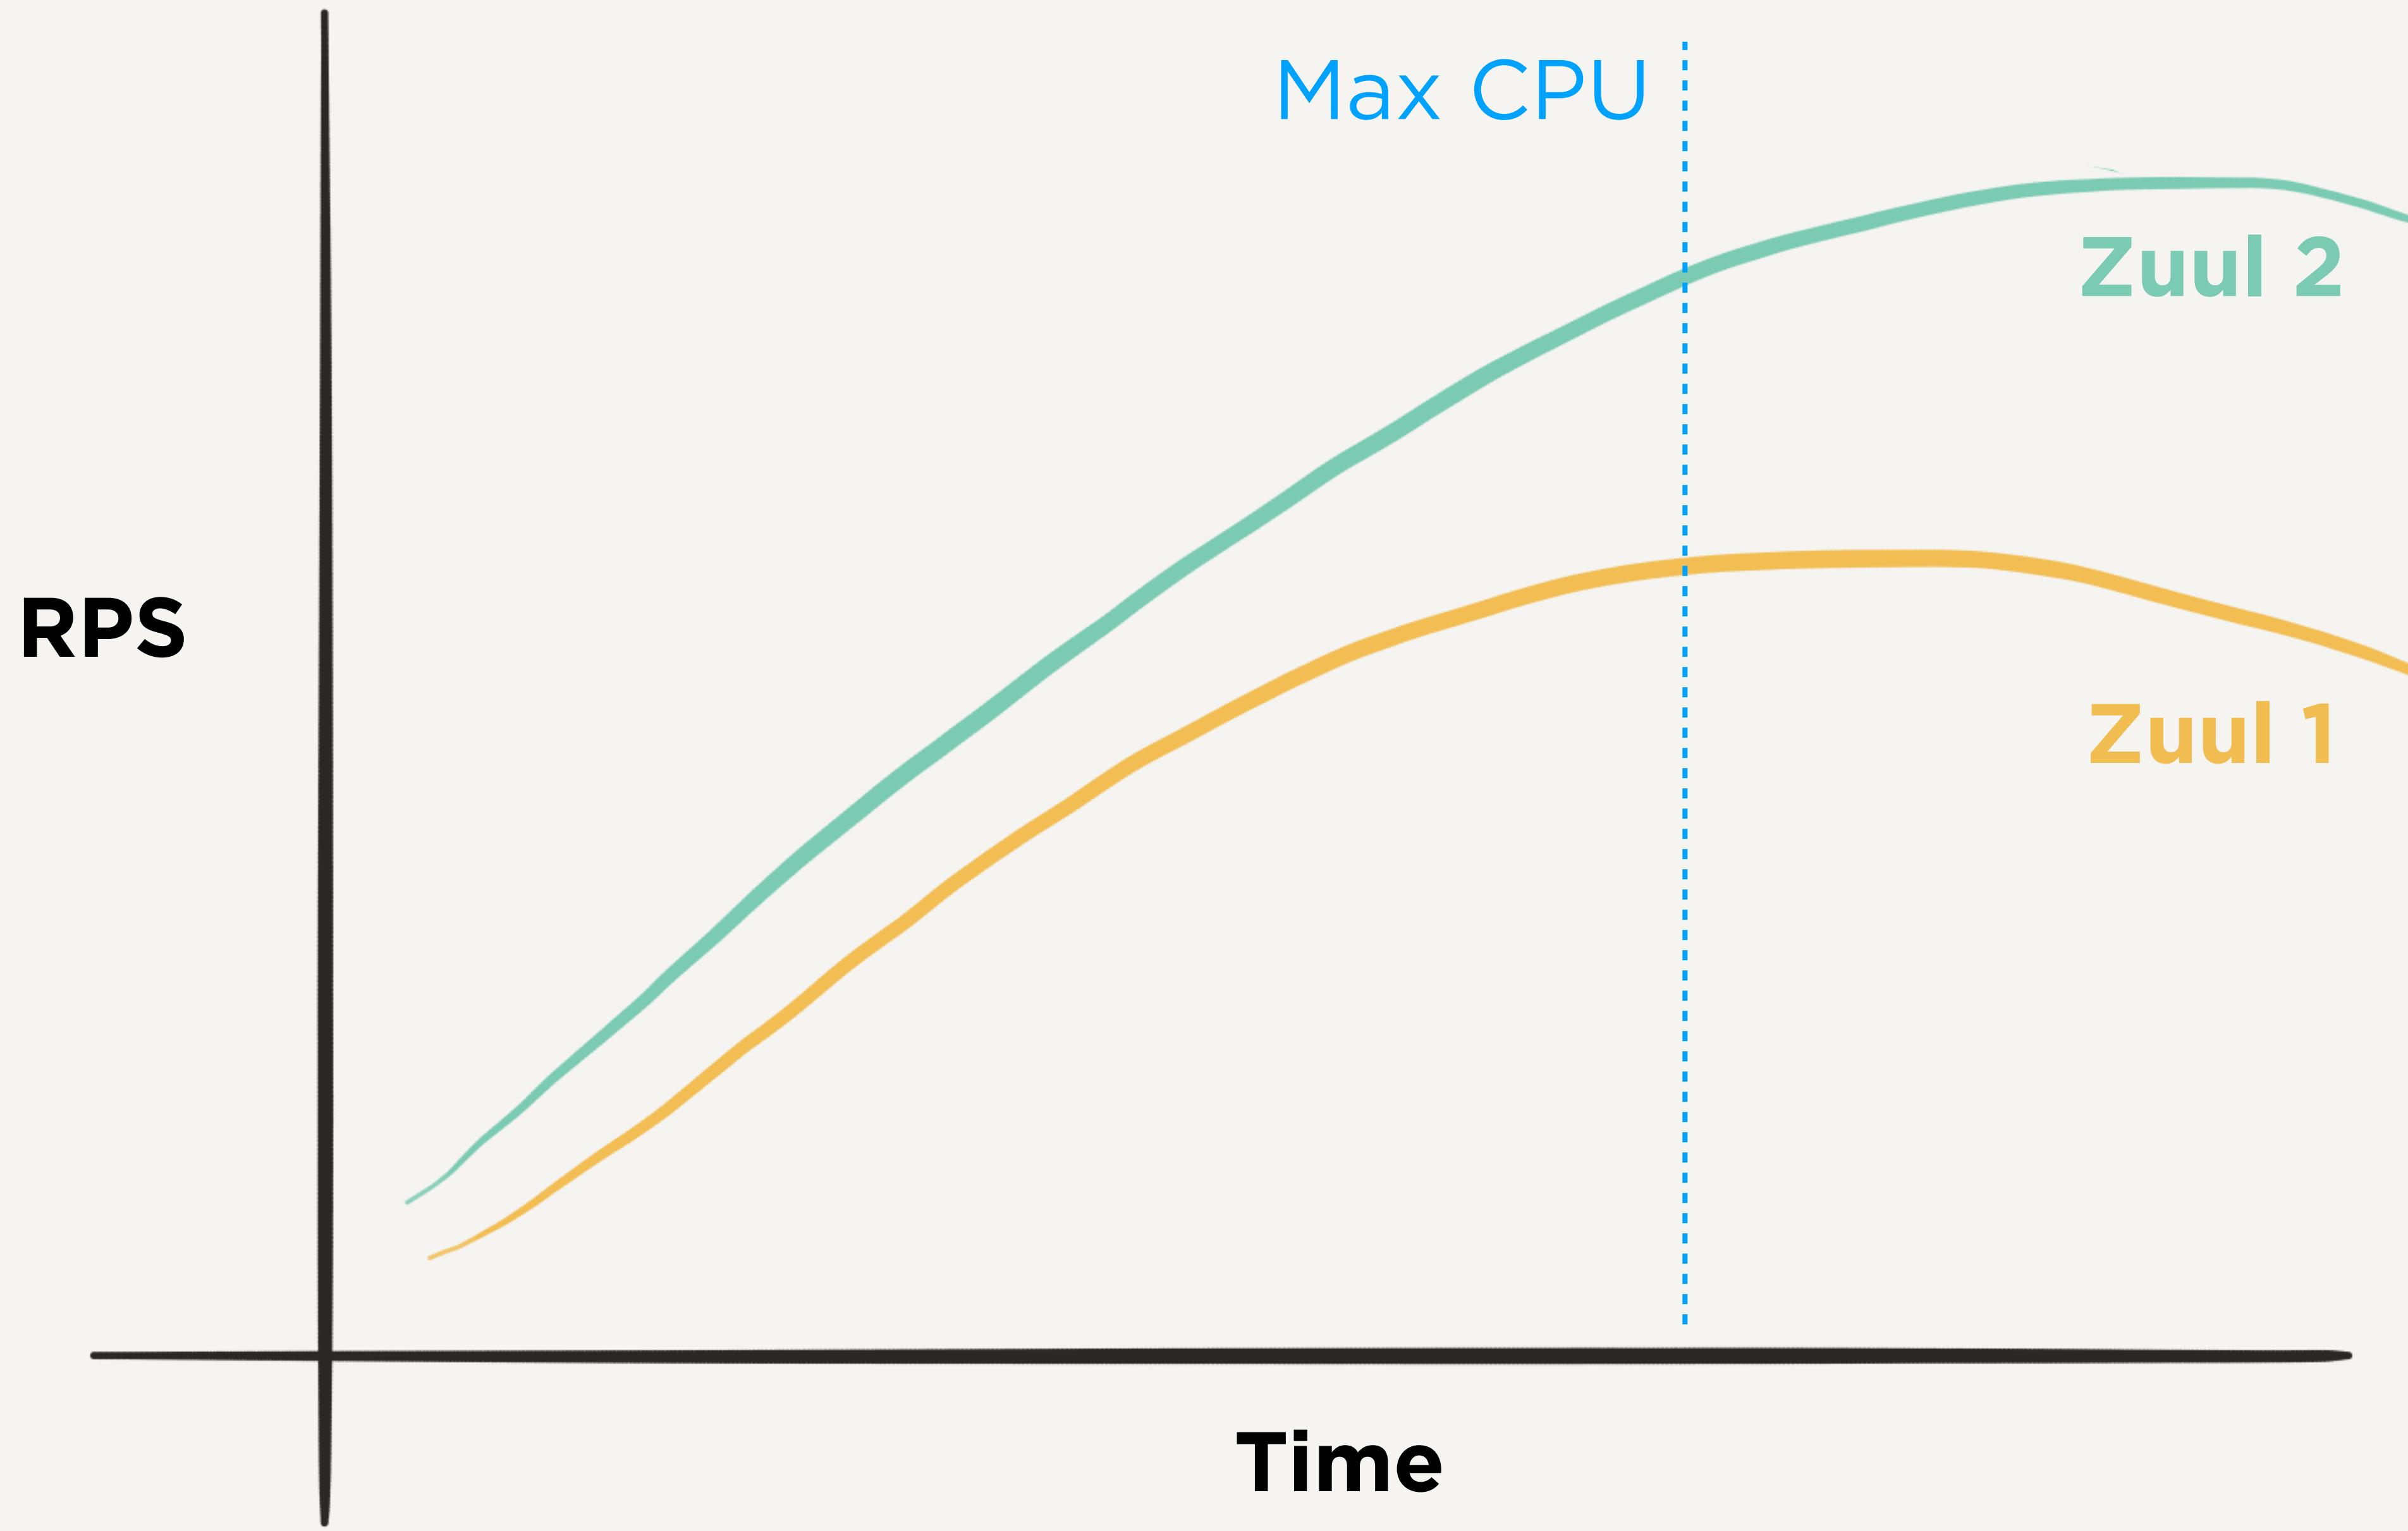
Fully up during massive **retry storm**

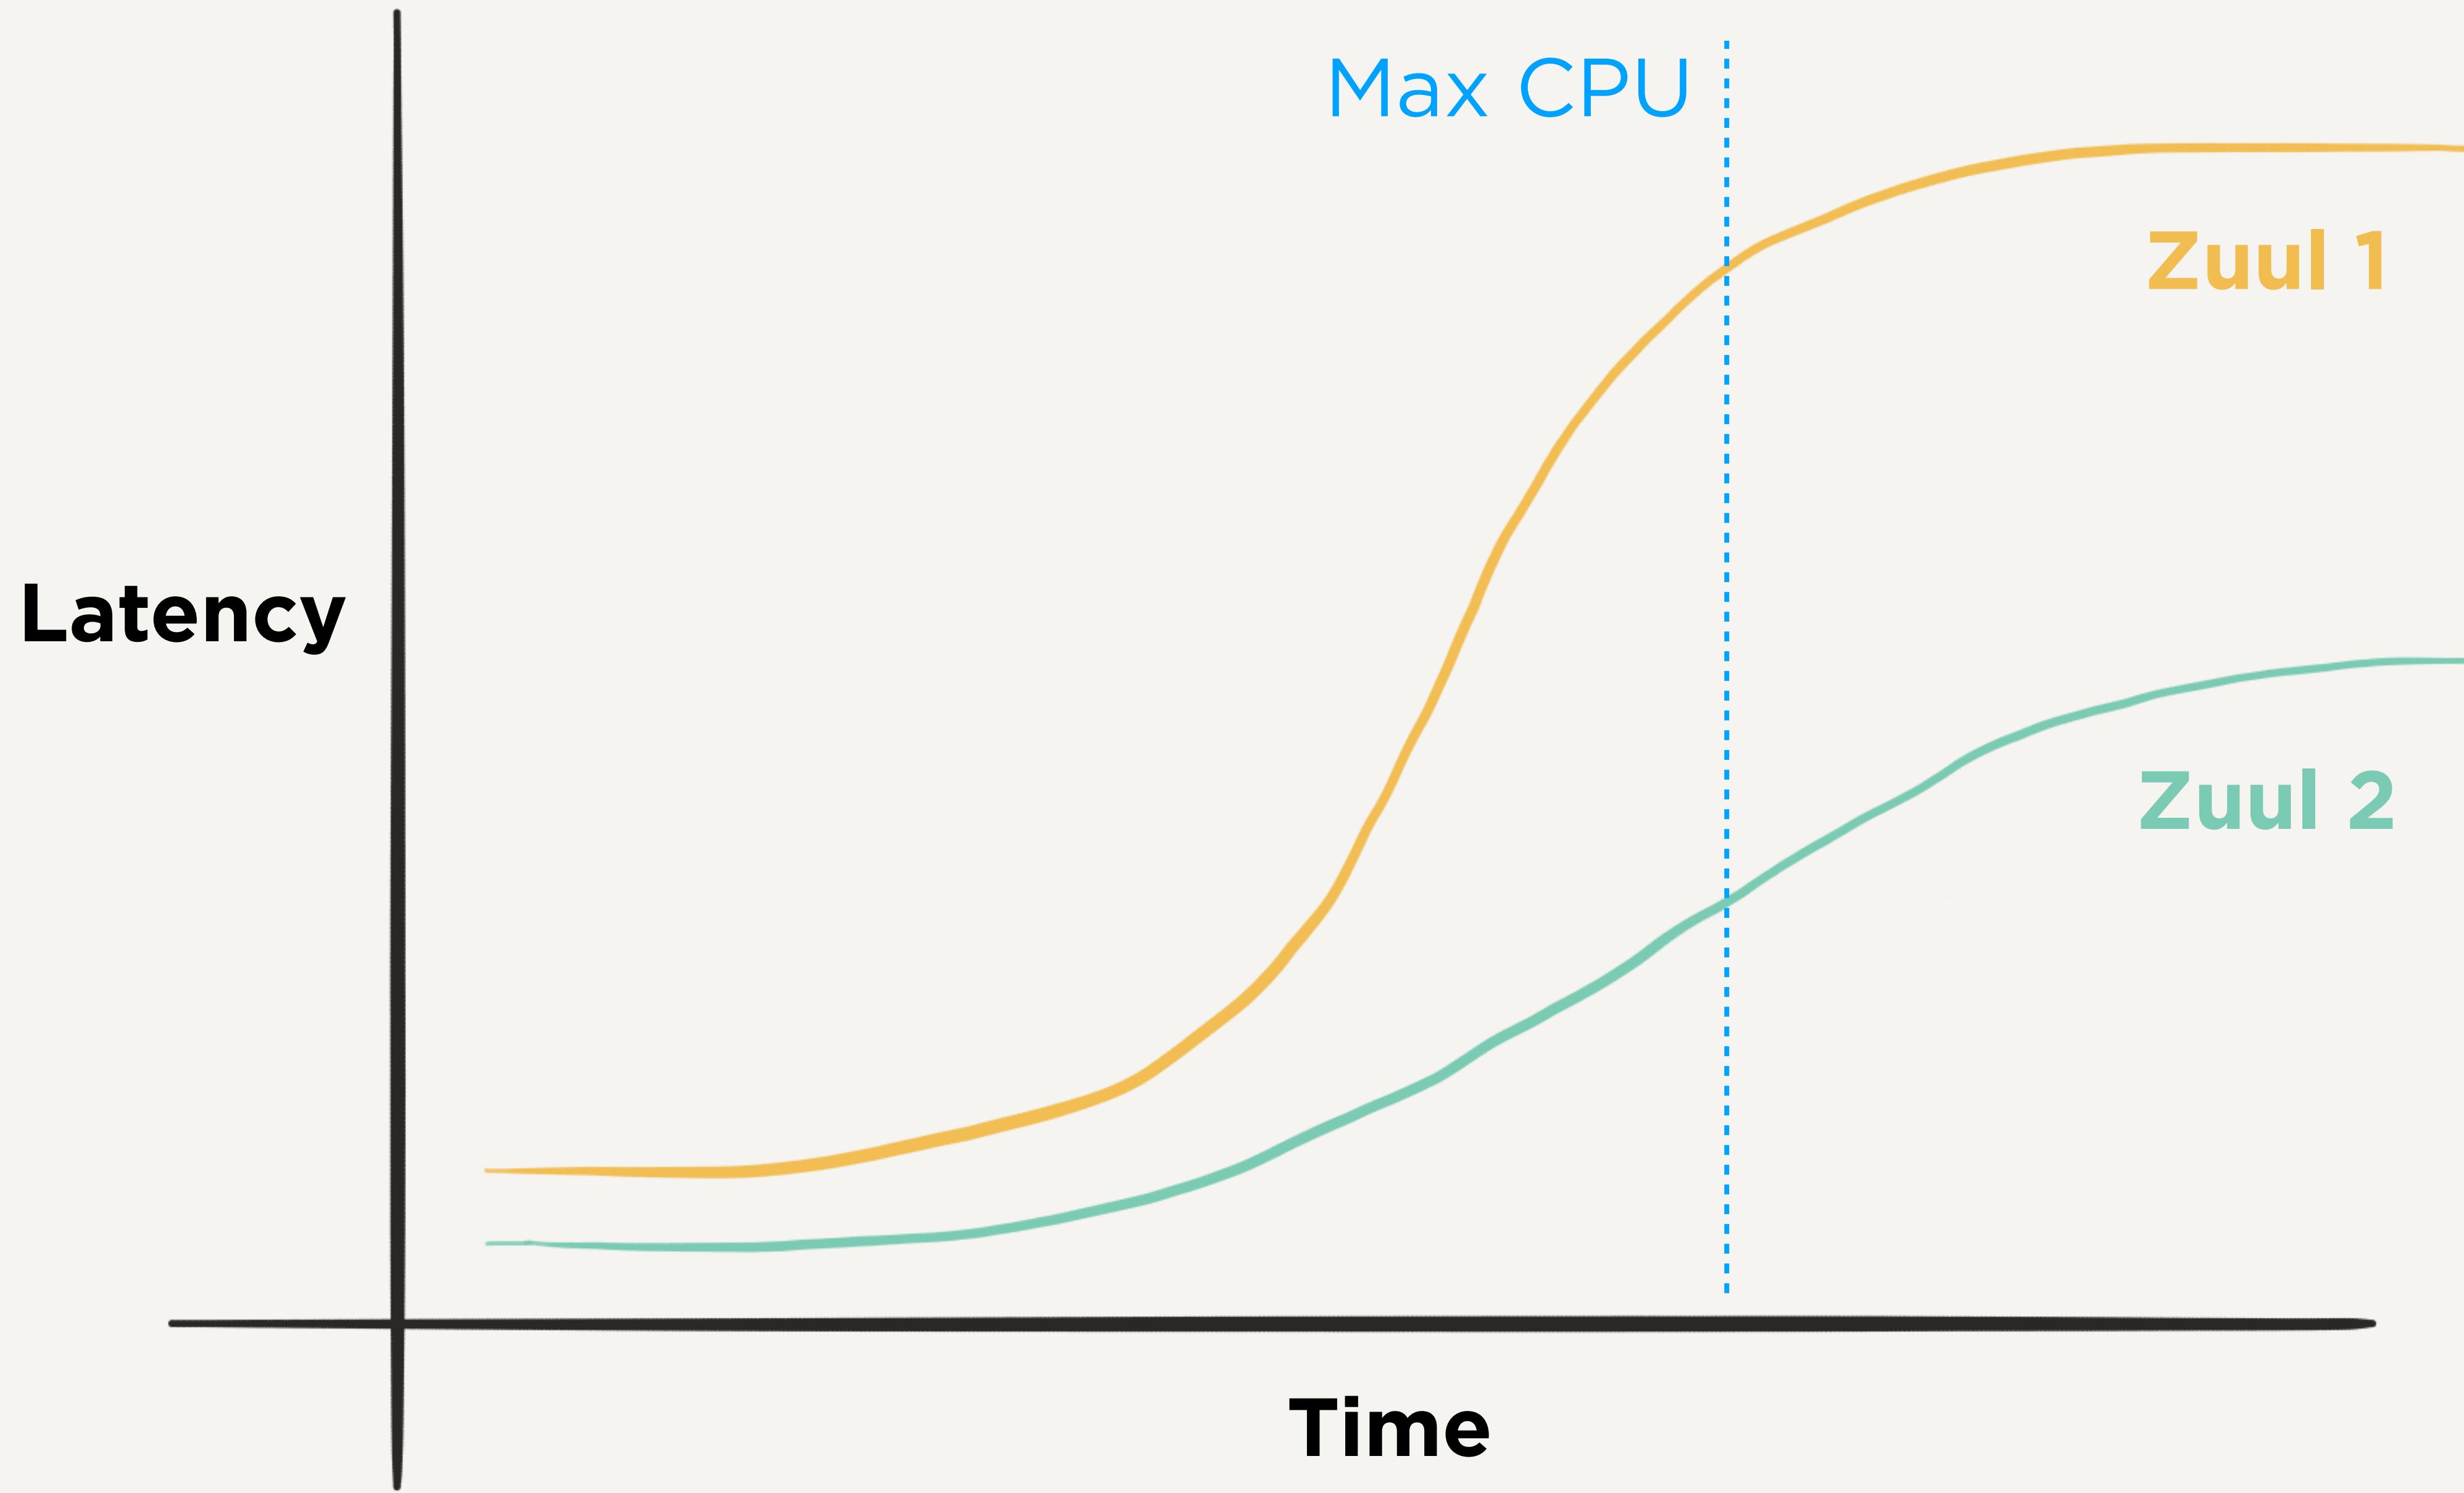
Continued serving 200s at max CPU

What is a retry storm?









Resiliency Benefits

Allows Zuul to **withstand** huge spikes
Shields origins from retry storms

Prevents throttling origins from going
underwater

Long Lived Connections

Tens of thousands of connections per node

Serving tens of millions of concurrent clients

Websocket support for push notifications

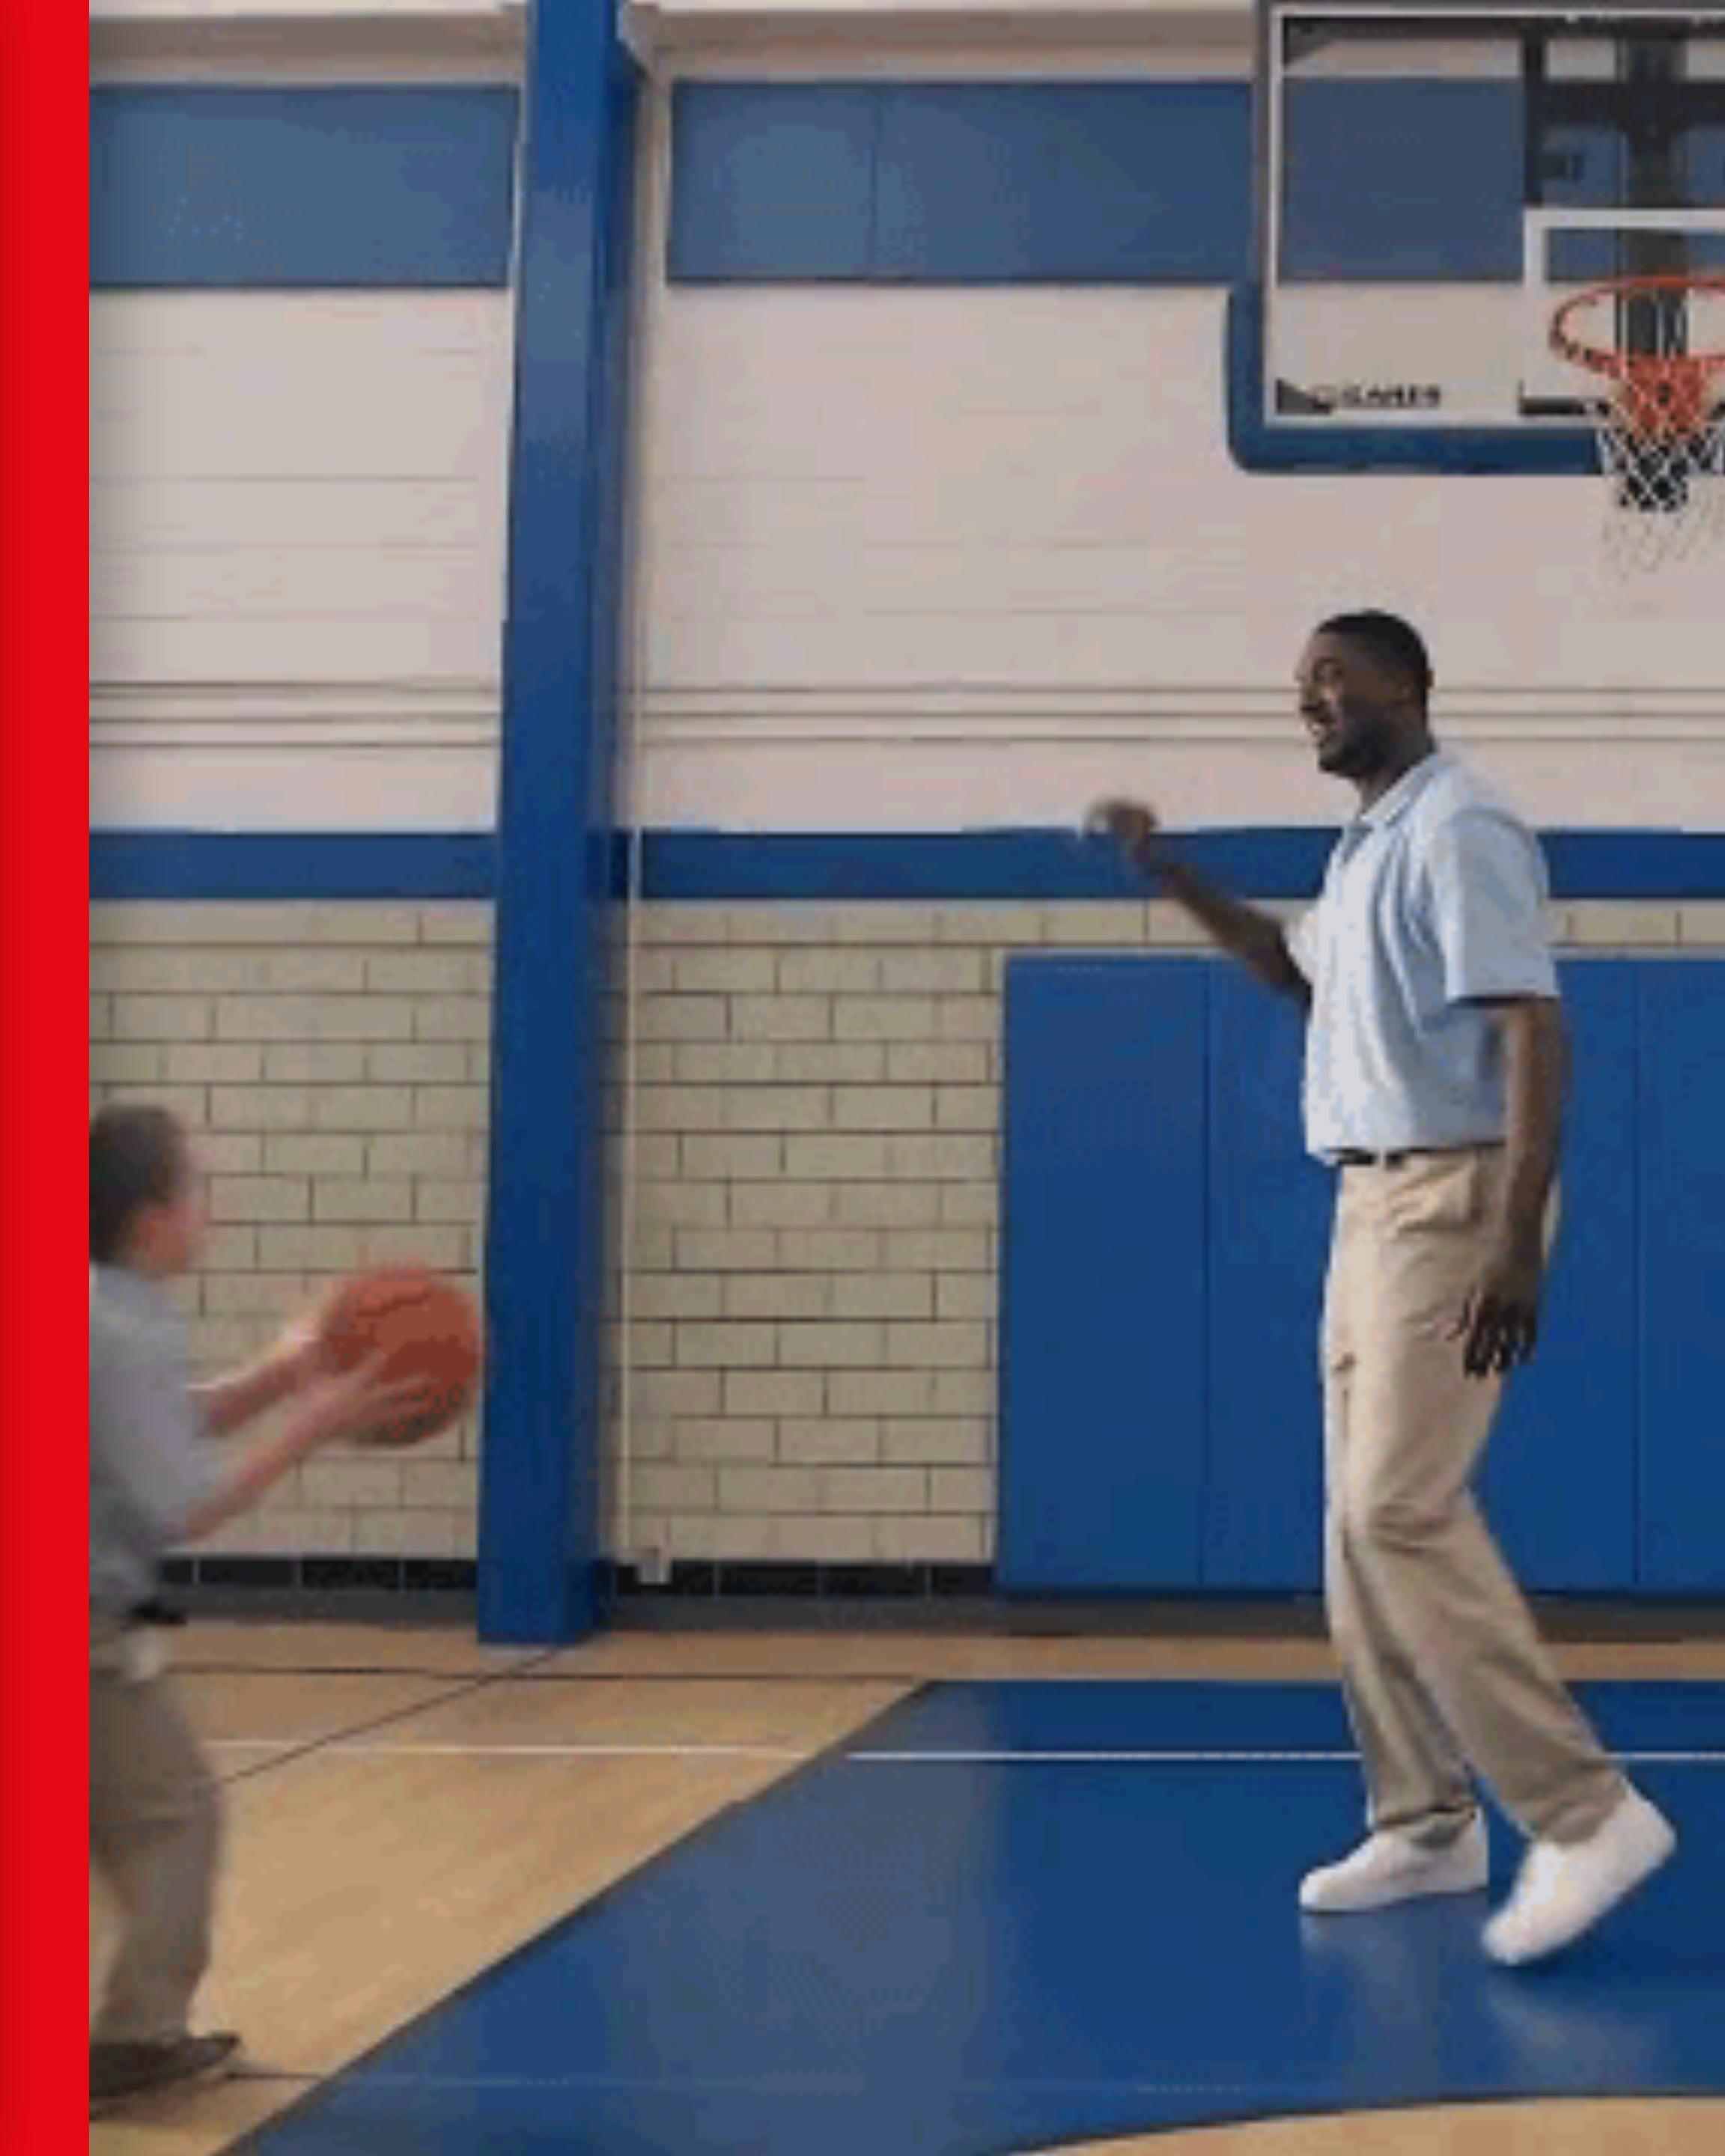
Allows for new features for partner teams

Relaxed **Latency** Constraints

Timeouts are no longer strictly guarded
Can be configured by origins or devices
Hystrix circuit breaking not necessary
Zuul will operate consistently regardless

Blocking or Non-Blocking

Which one is right for you?

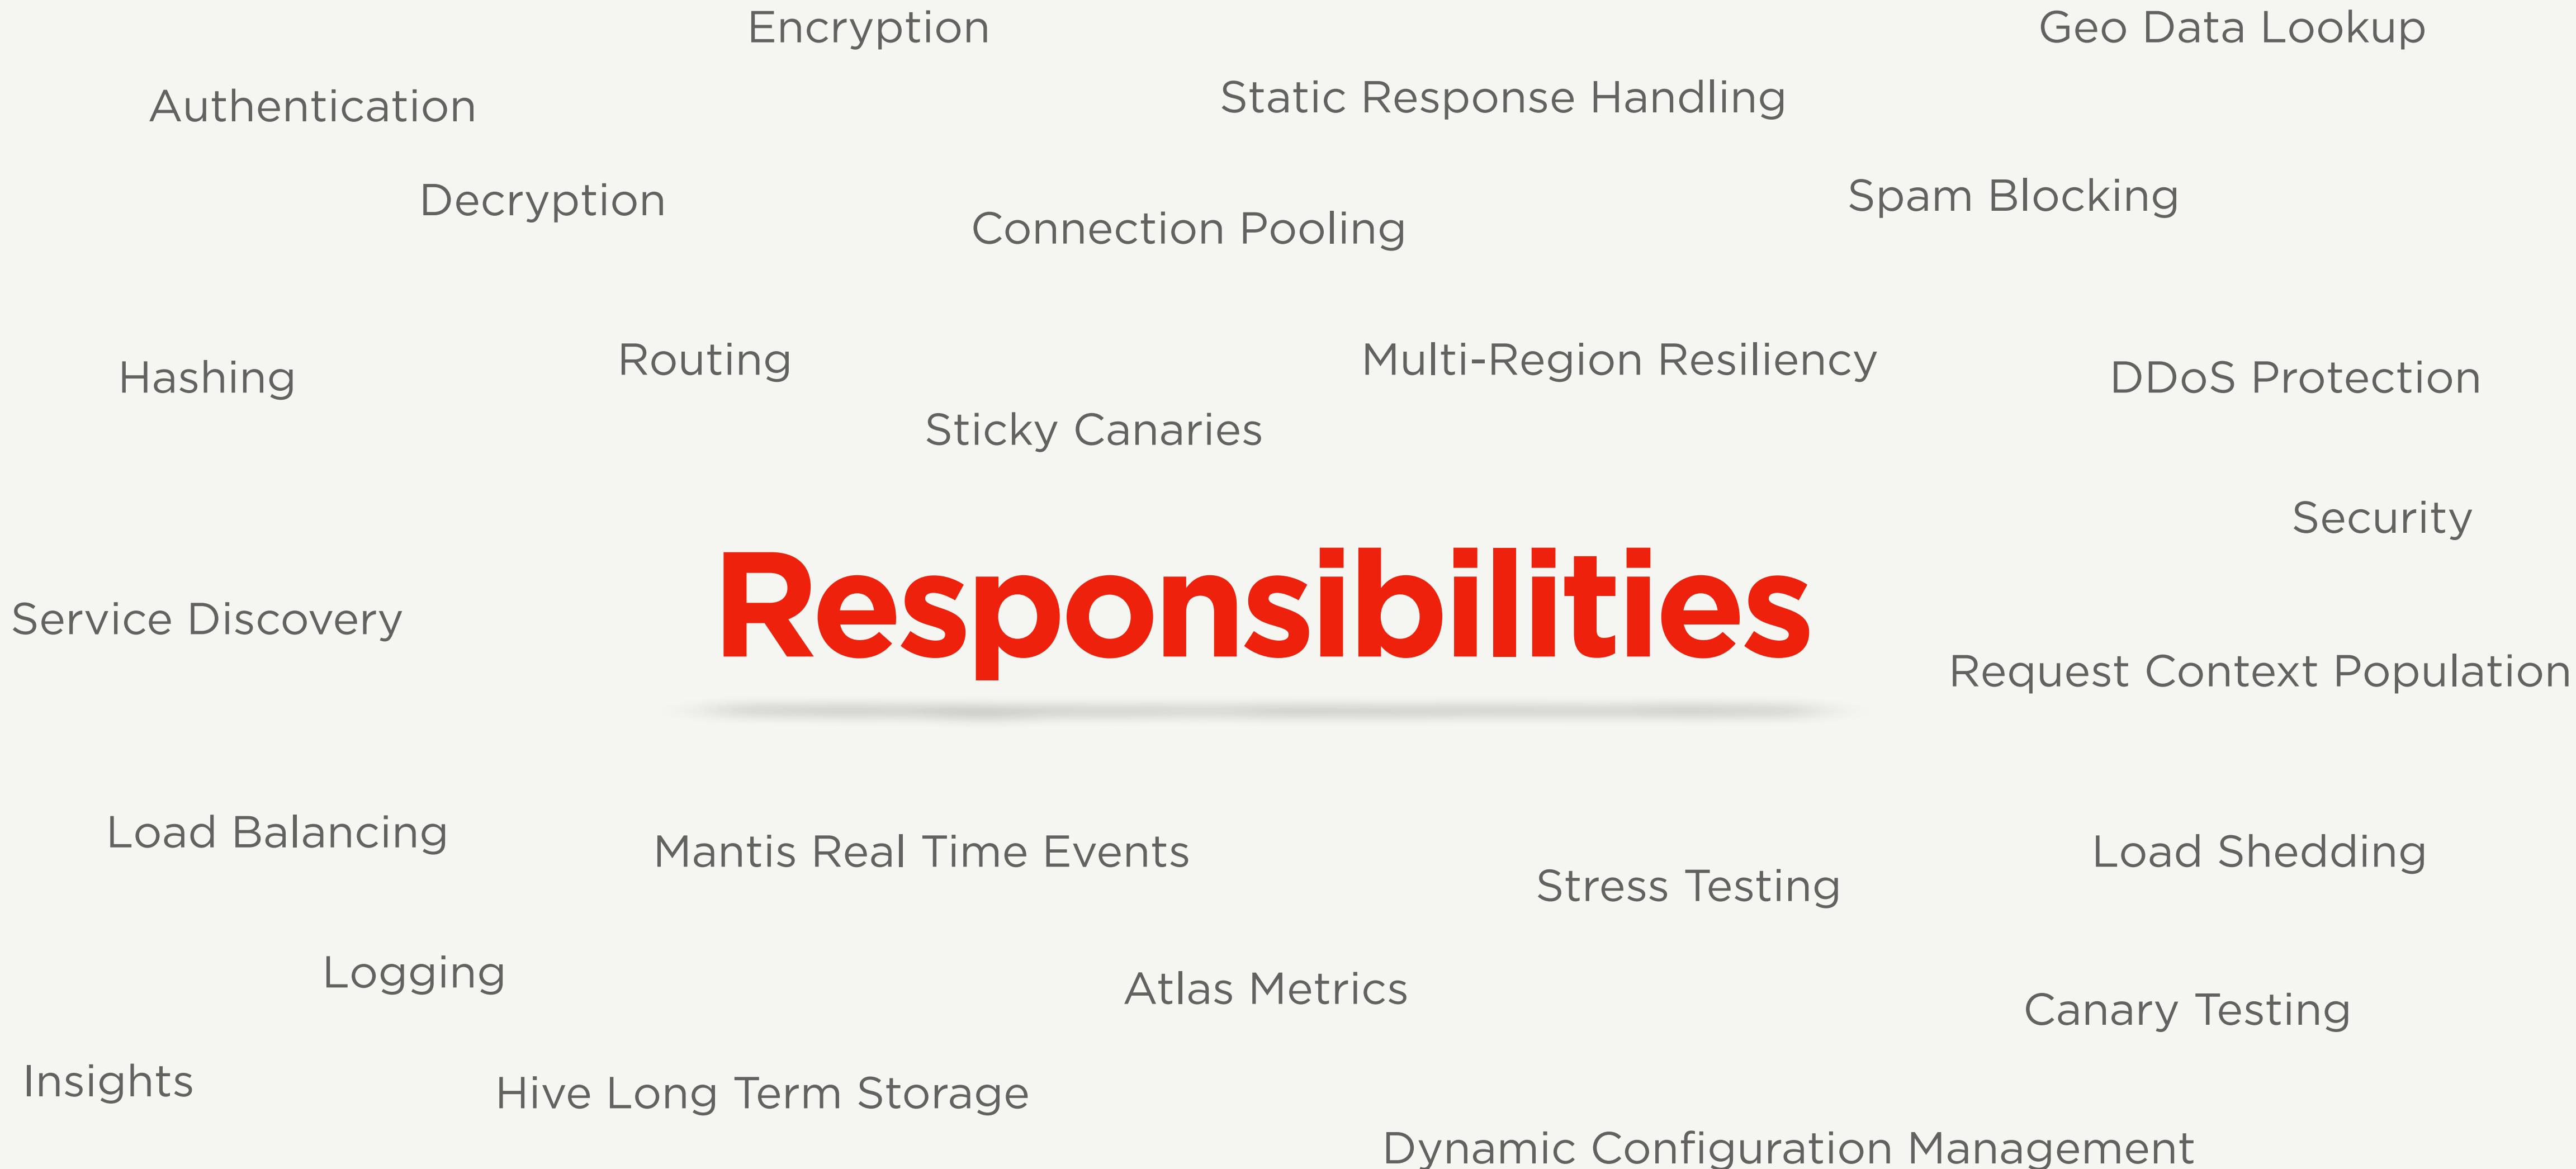


Real World Performance



Request per second per node

Responsibilities



When to Block

Highly **CPU-bound** work loads

Desire operational simplicity

Desire development simplicity

Run legacy systems that are blocking

When **Not** to Block

Highly **I/O-bound** workloads

Long requests and/or large files

Streaming data from queues

Massive amounts of connections

Benchmark Accuracy

Important to set up **realistic** benchmarks
Simulate real-world workloads

Verify assumptions and trendy dogma

Long Term Costs

Account for lack of debuggability

Ongoing maintenance cost of async code

Support and operations burden

Critical Features

What features do you really need?
Are those features worth the costs?
What's your answer to "**it depends**"?

Was it worth it?

Thank you.

Office Hours: Netflix Booth 10am-12pm tomorrow

NETFLIX