

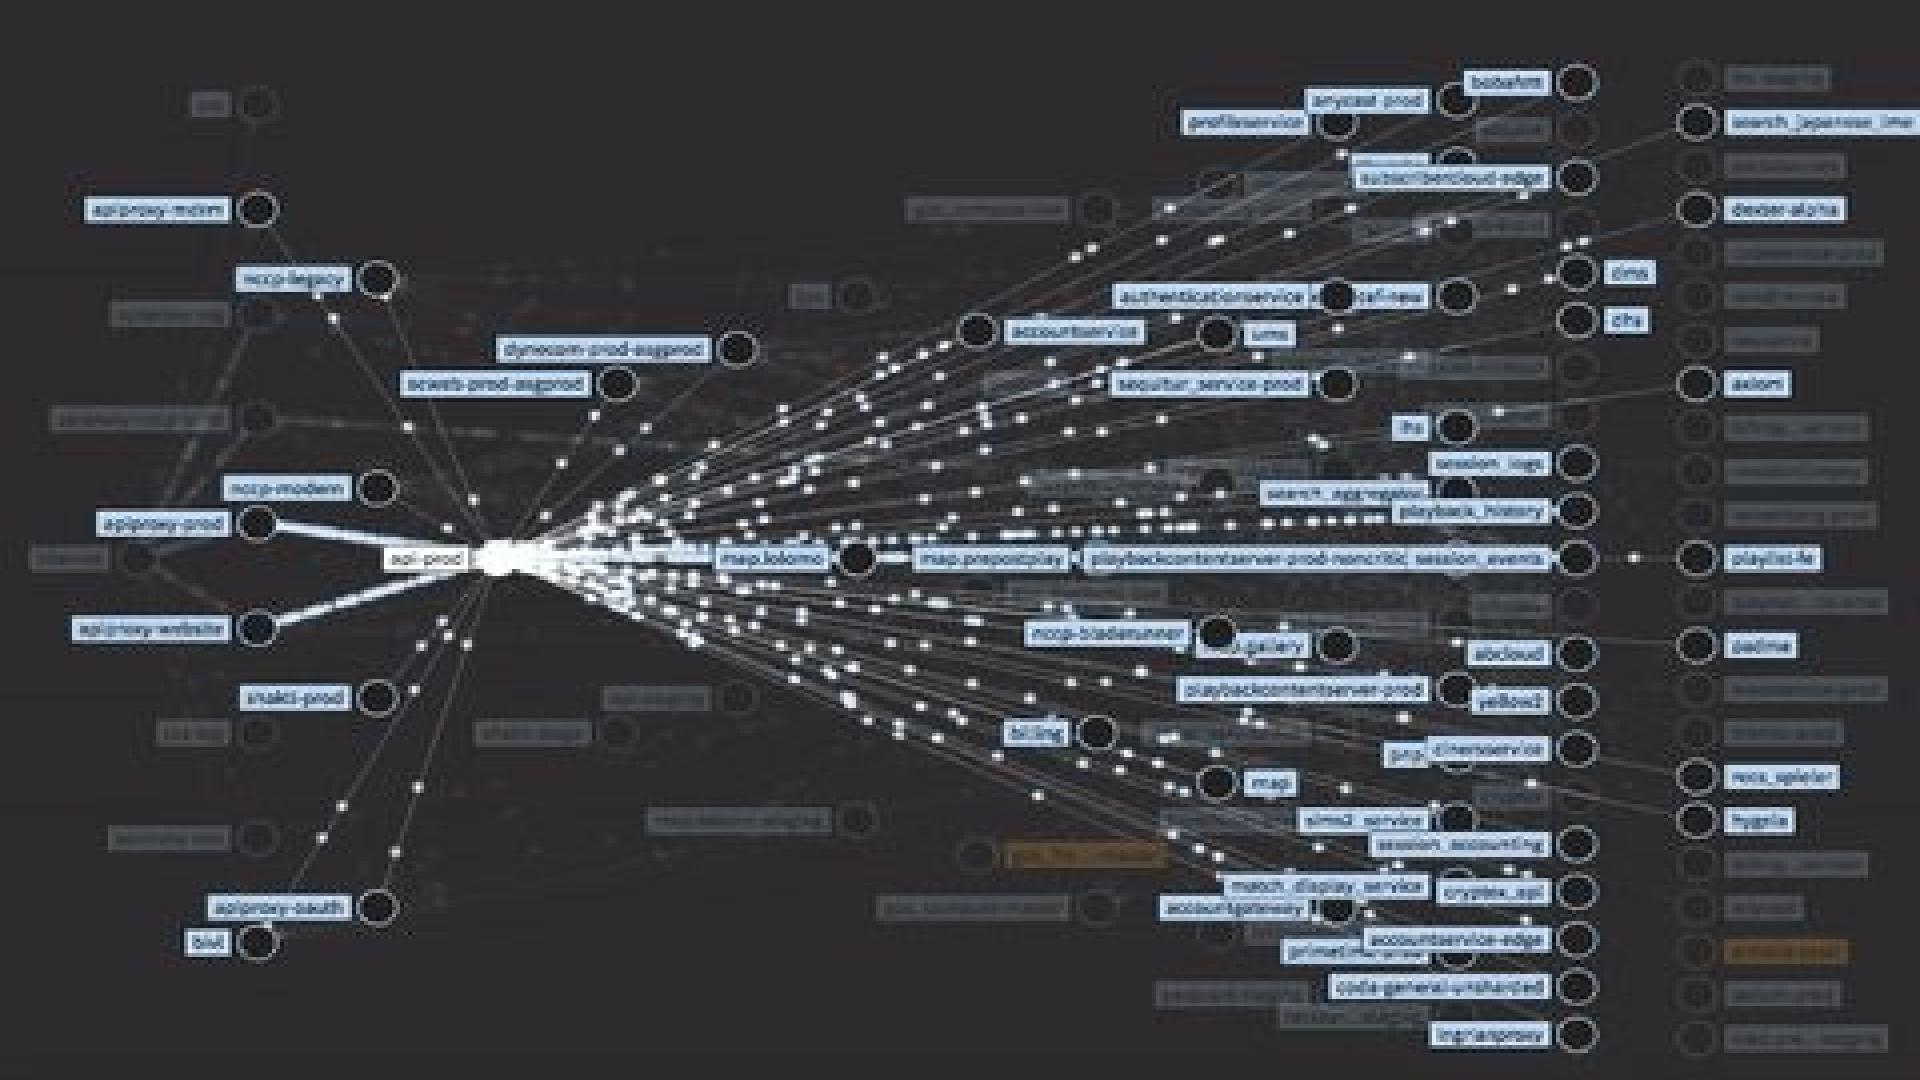
# **Running a Massively Parallel Self-serve Distributed Data System At Scale**

Zhenzhong Xu |  @ZhenzhongXu  
Real-time Data Infrastructure









# Running a Massively Parallel Self-serve Distributed Data System at Scale

- What is Netflix Real-time Data Infrastructure
- Challenges
- Solutions and Principles

# What is ...

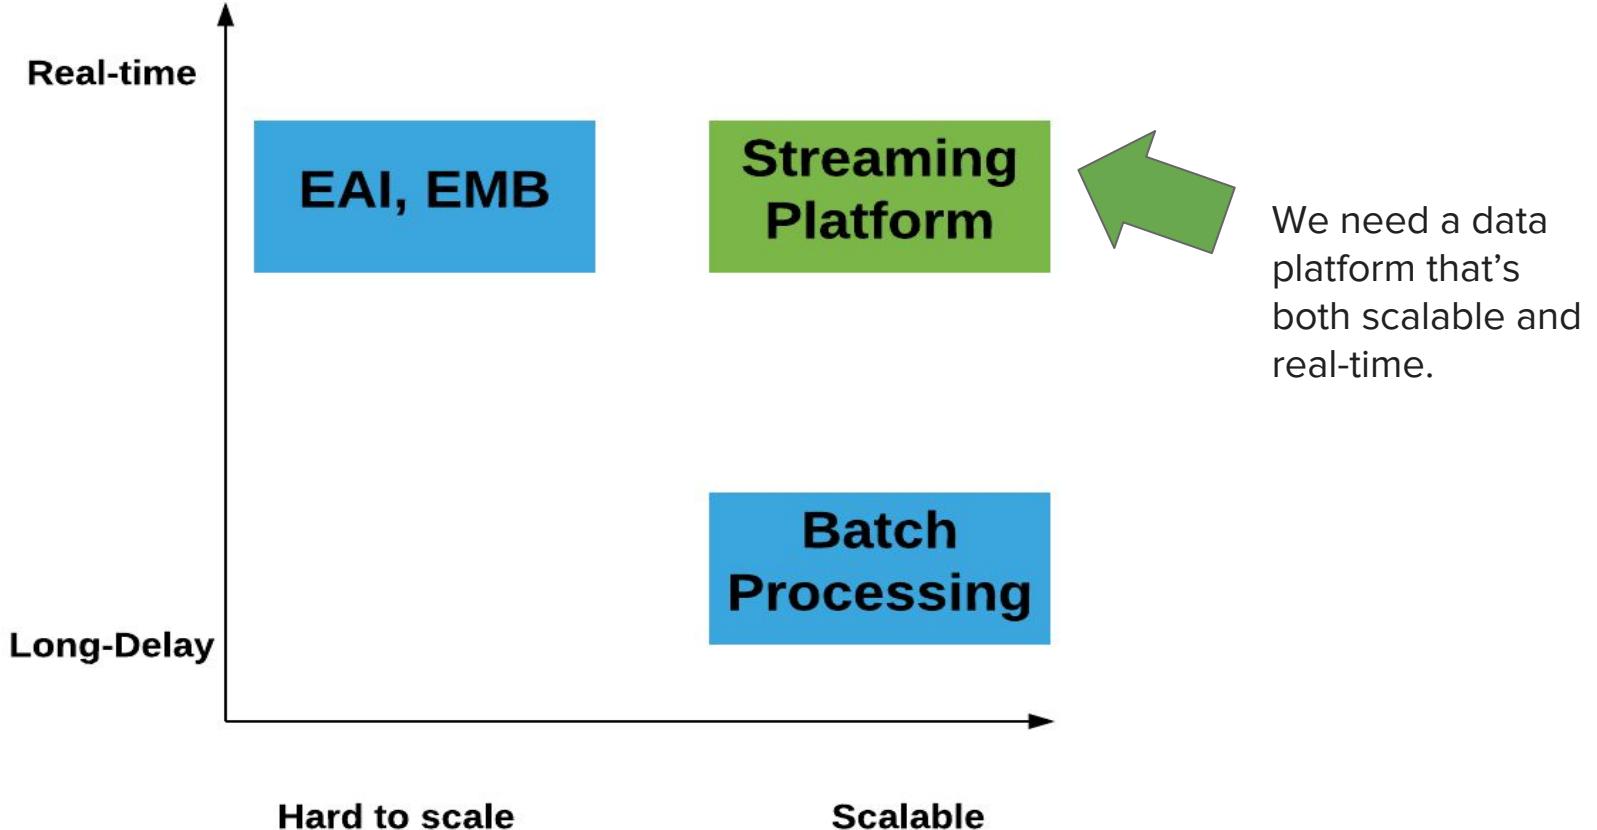
| ... Real-time Data Infrastructure at Netflix

# **Business/Product Driven Analytics**

- Recommendations / Personalization Algorithms
- Customer Experience
- Content Operation
- A/B Testing
- Marketing
- etc

# Rise of Event Driven Architecture

- Notification
- Event Sourcing
- CQRS
- etc



# Data-driven Culture



**So what exactly is Keystone Streaming Platform?**

**Publish, Collect, Move & Compute**

**event data in near real time @ Cloud Scale**

# Keystone is ...

... a collection of microservices & components

Self Service UI

Control Plane



Producer API



Pub/Sub Queuing Service



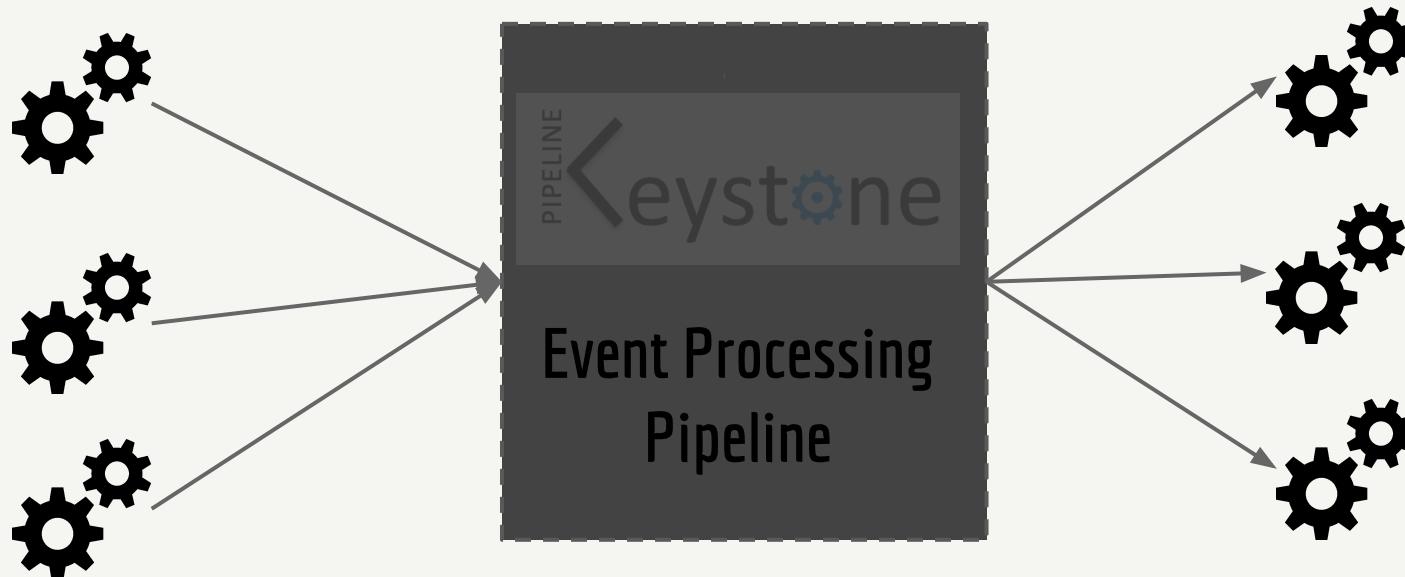
Stream Processing Service



Consumer API

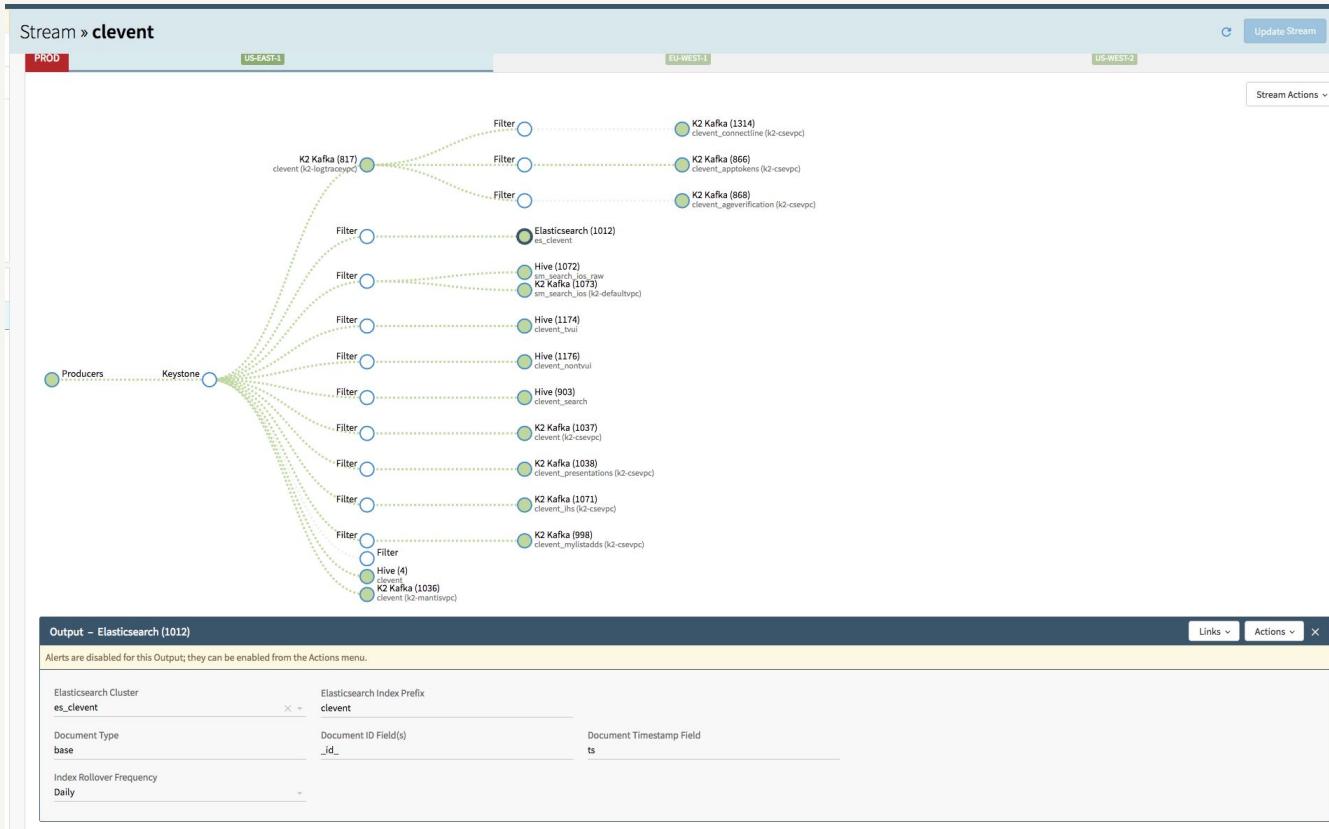
# Keystone is ...

... a single self-contained logical PaaS



# Keystone is ...

## ... a multi-tenants, self-serving tool

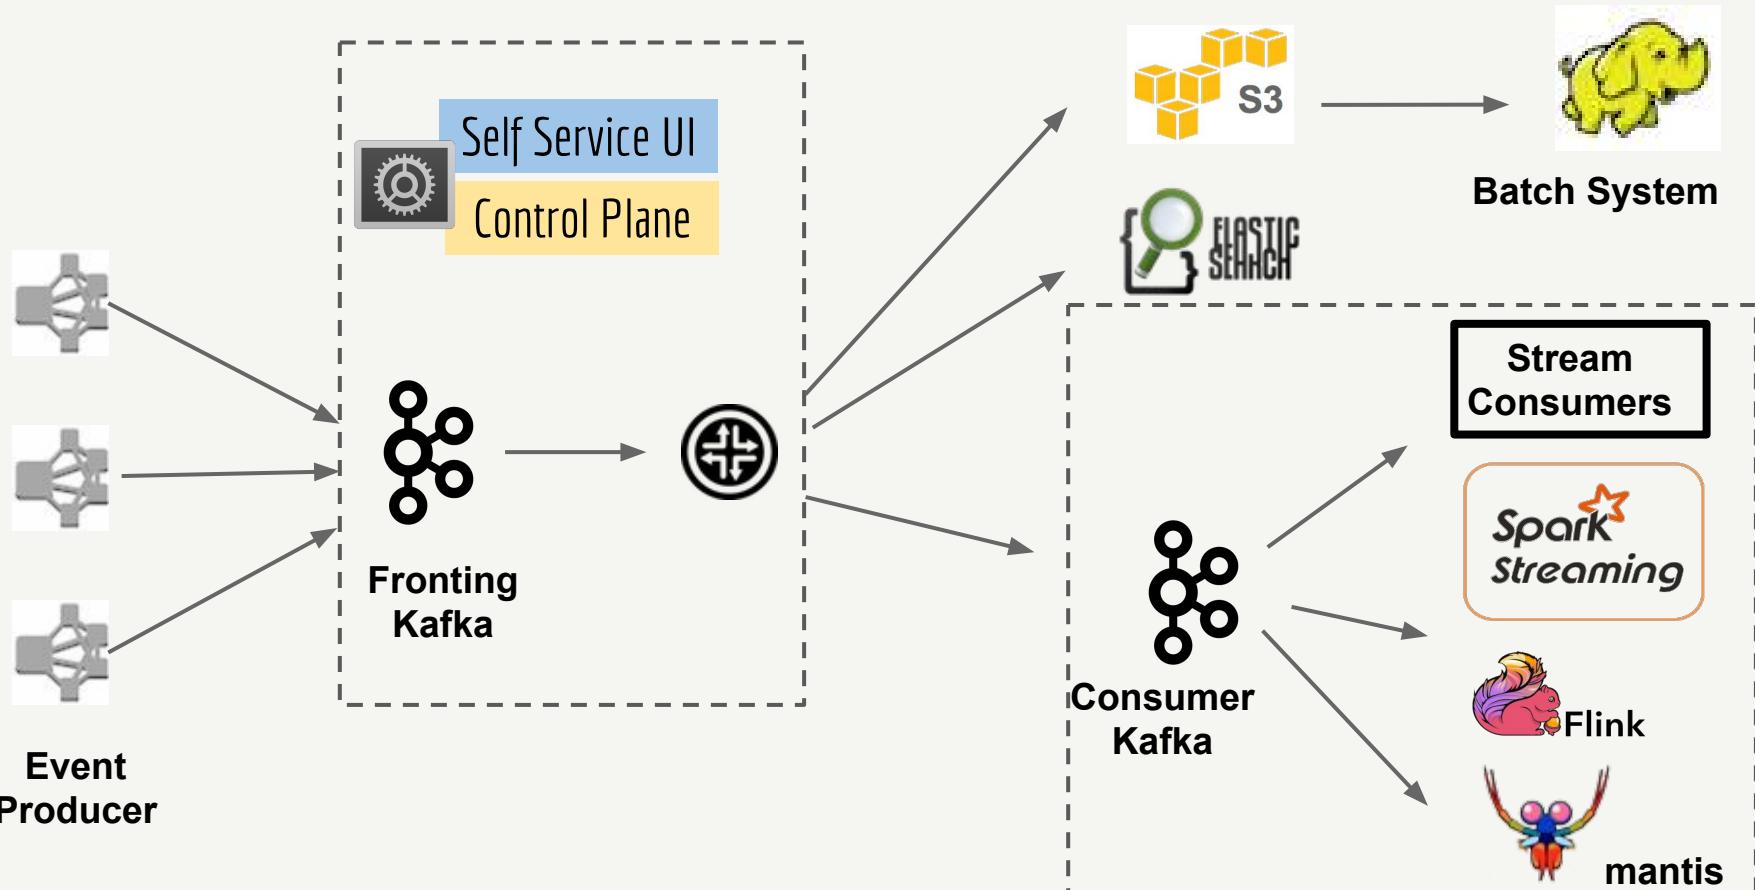


# Keystone is ...

... a self healing, cloud failure tolerant service,  
guarantees at-least-once delivery semantics

... adaptable to changing environments

# Putting together ...



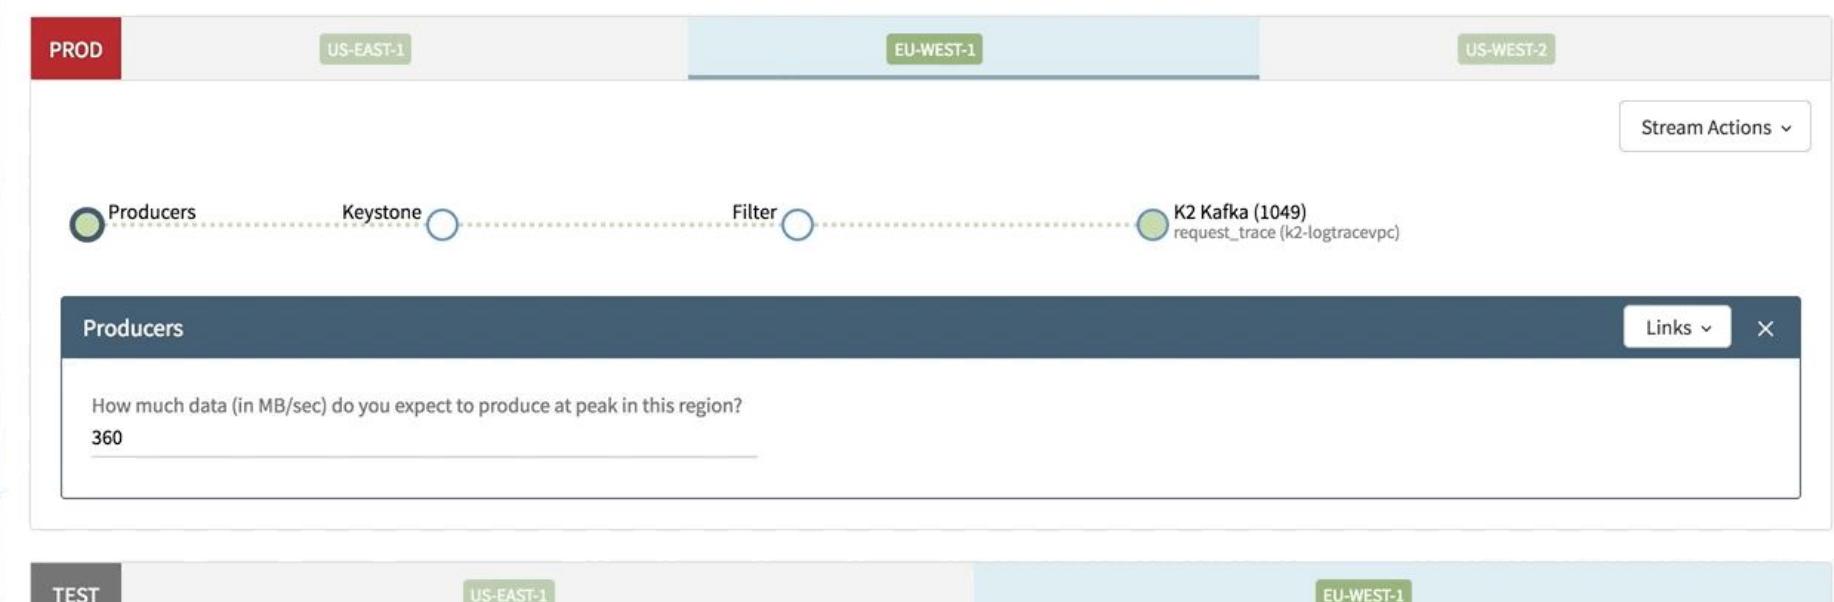
# Challenges Solutions & Principles

| ... the decisions we made to build scalable, reliable and  
  maintainable systems.

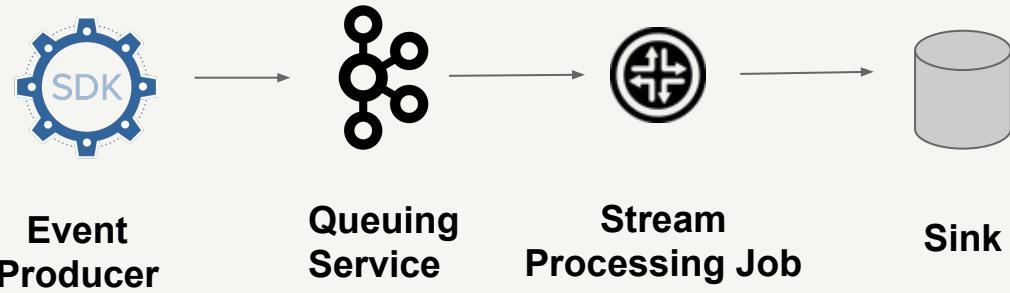
# **Challenge 1:**

## **Scale.**

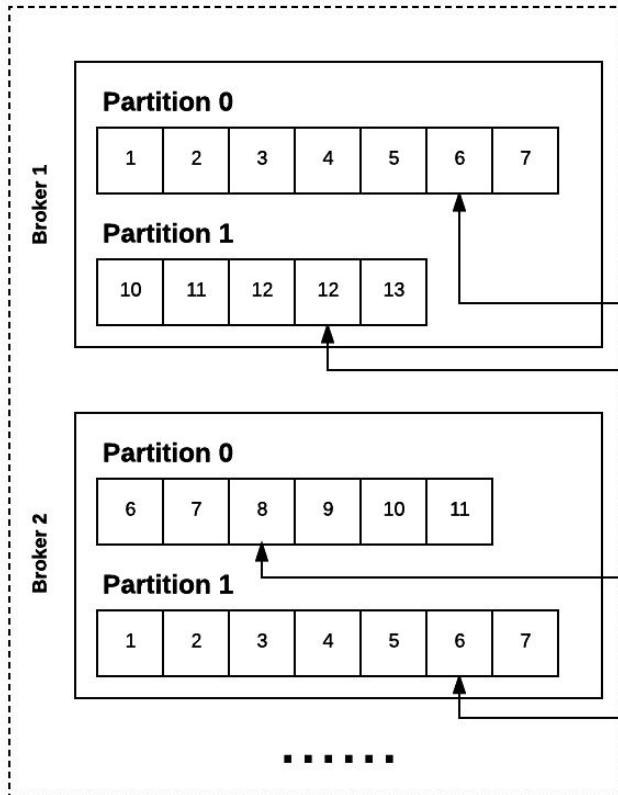
# A Single Stream



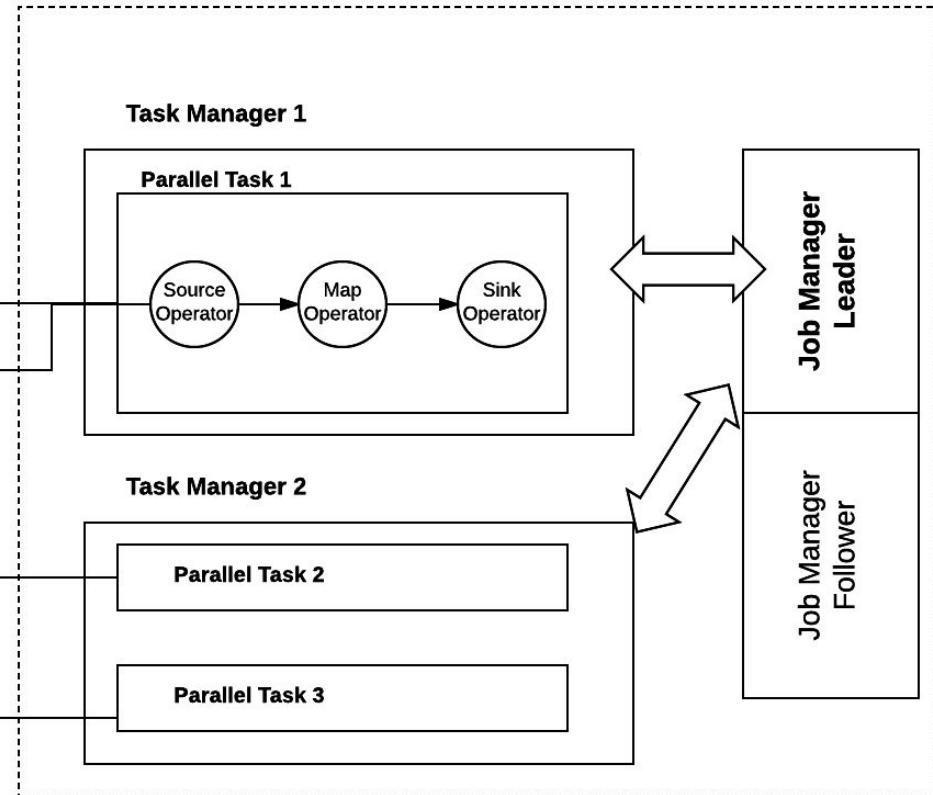
## Anatomy of a Single Stream



## Kafka Cluster



## Flink Job Cluster (Parallel Style)

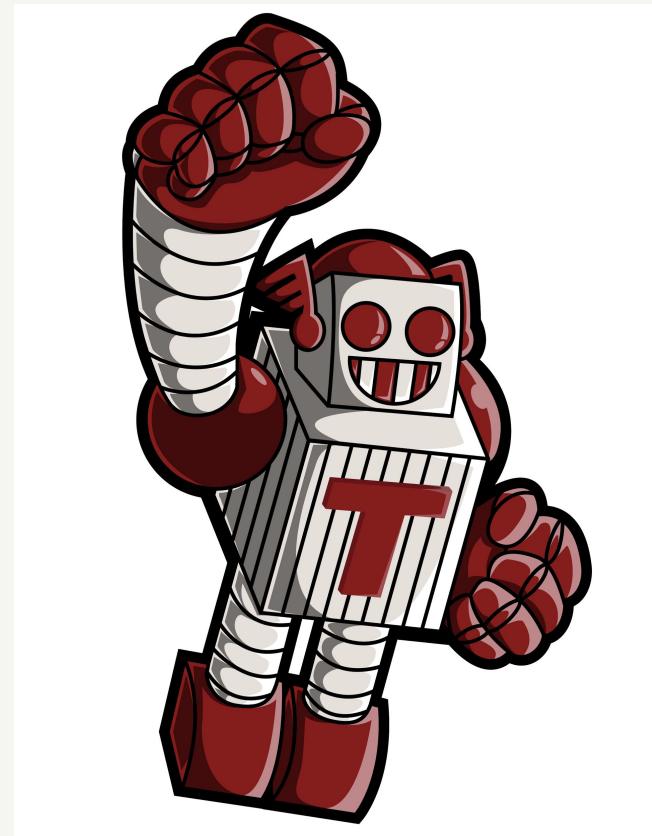


## Separation of Concerns

- Separation of Messaging and Stream Processing services.
- Each service scales individually.
- Each service manages its own states.
- Independently manage service dependencies.
  - Kafka brokers on EC2
  - Streaming Service job on Titus Container Runtime

## Titus Container Runtime

- Resource Provisioning
- Scheduling and bin-packing
- Capacity Guarantees
- Resource isolation
- Per container IP address (underlay via VPN)

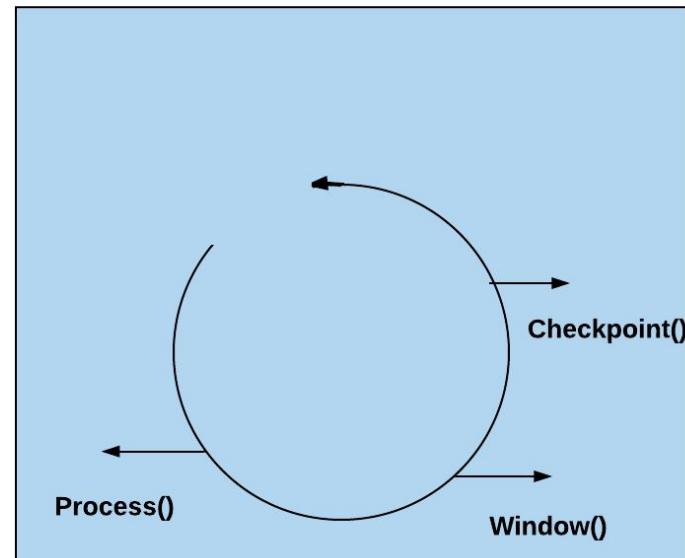


## **Delivery/Processing Semantics**

- At-most once
- At-least once
- Exactly once

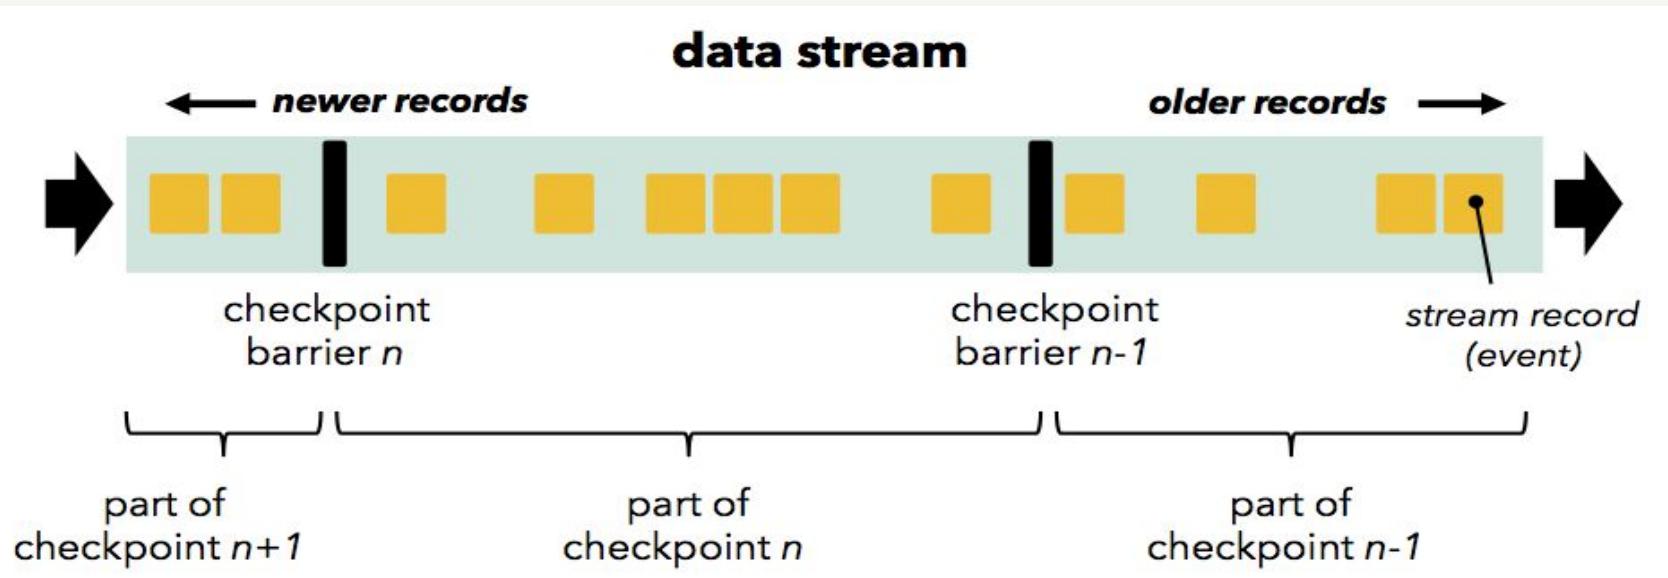
## At-least Once Processing Checkpointing

- Synchronous checkpointing through event loop

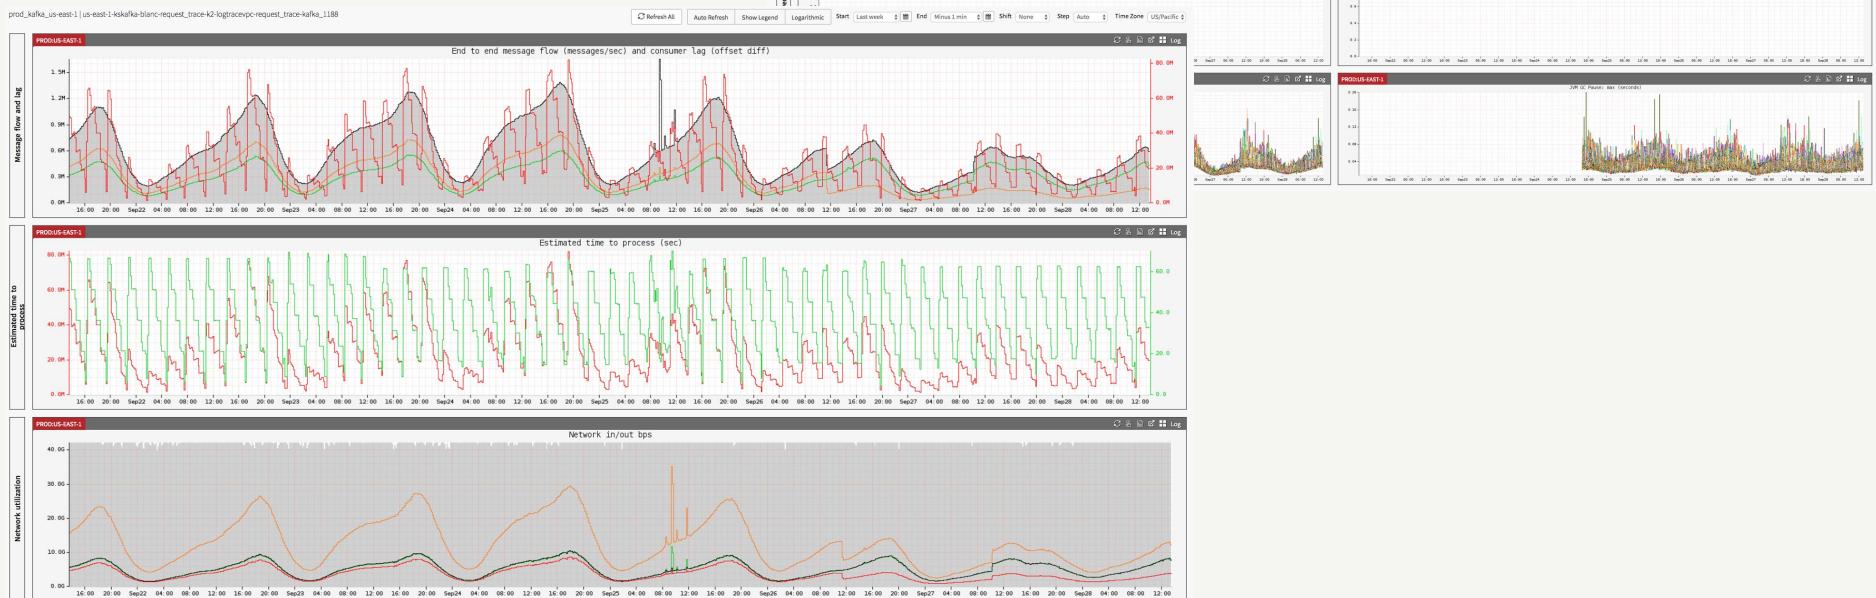


## Exactly Once\* Processing Semantics

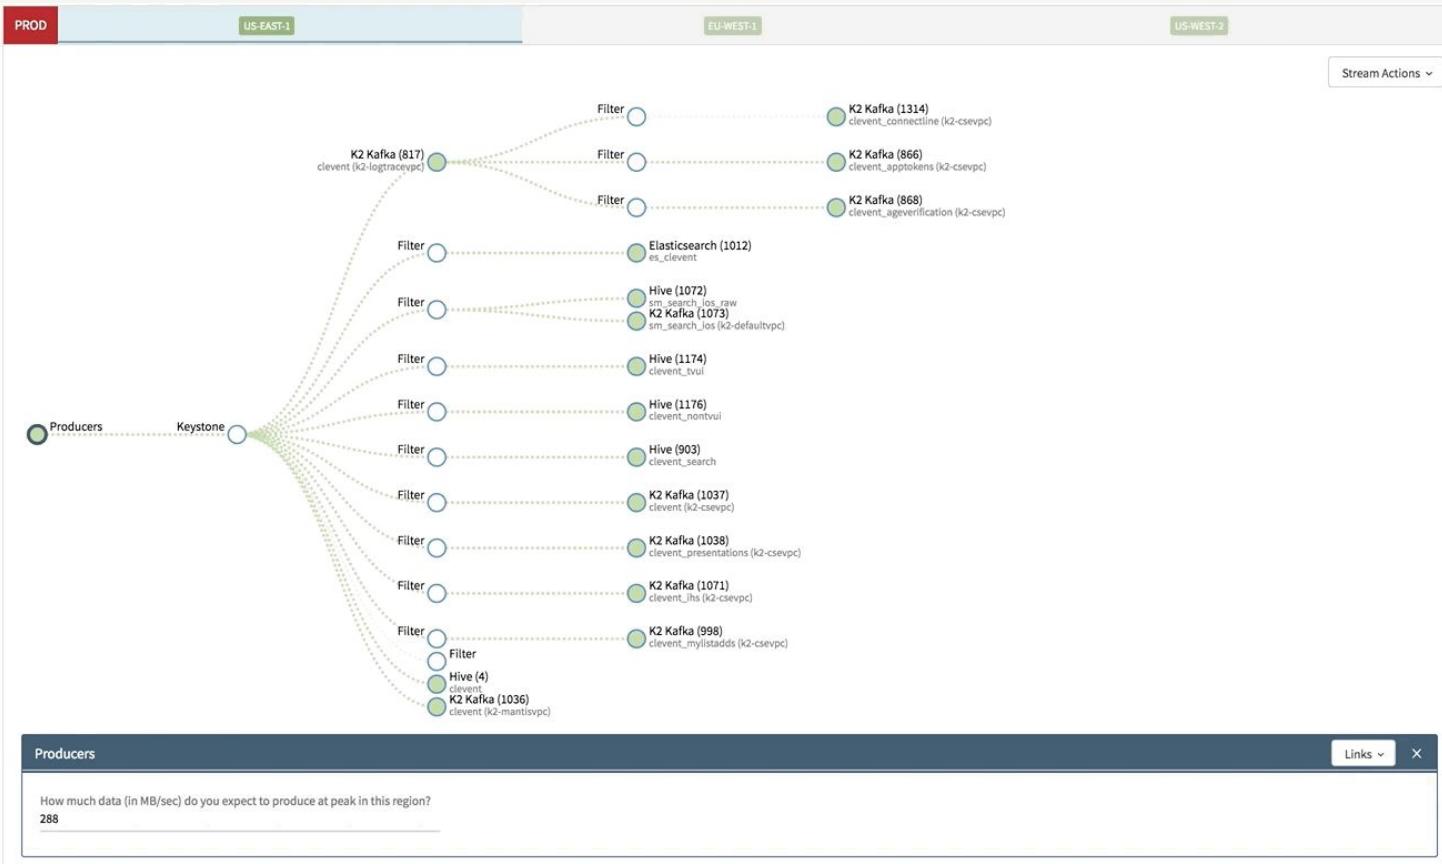
- Lightweight Asynchronous Snapshot (Async Barrier Checkpointing)



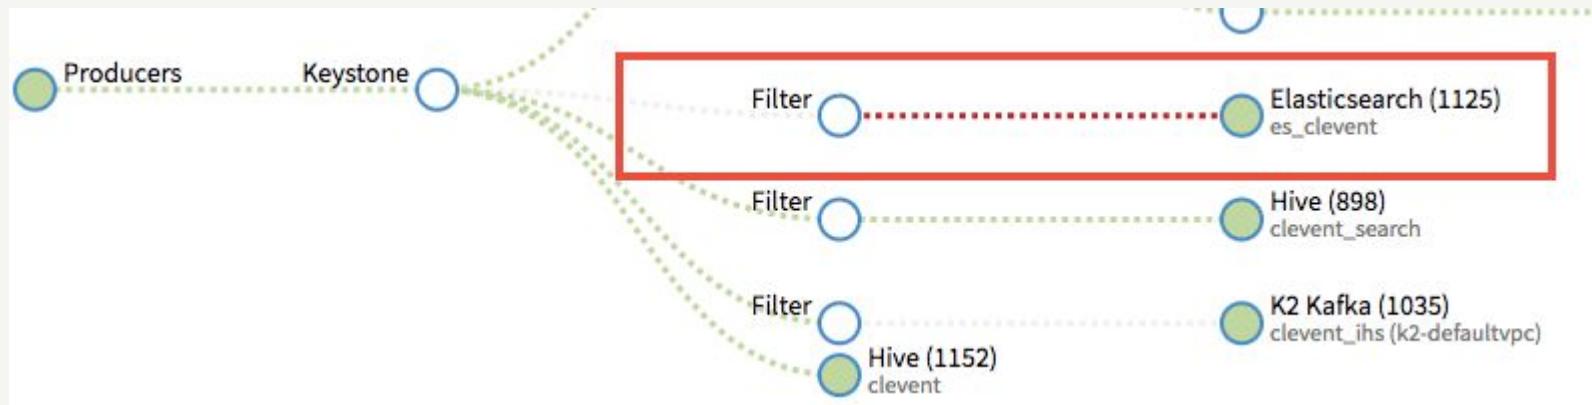
# Per Stream Monitoring & Alerting



# Streams with Fanout

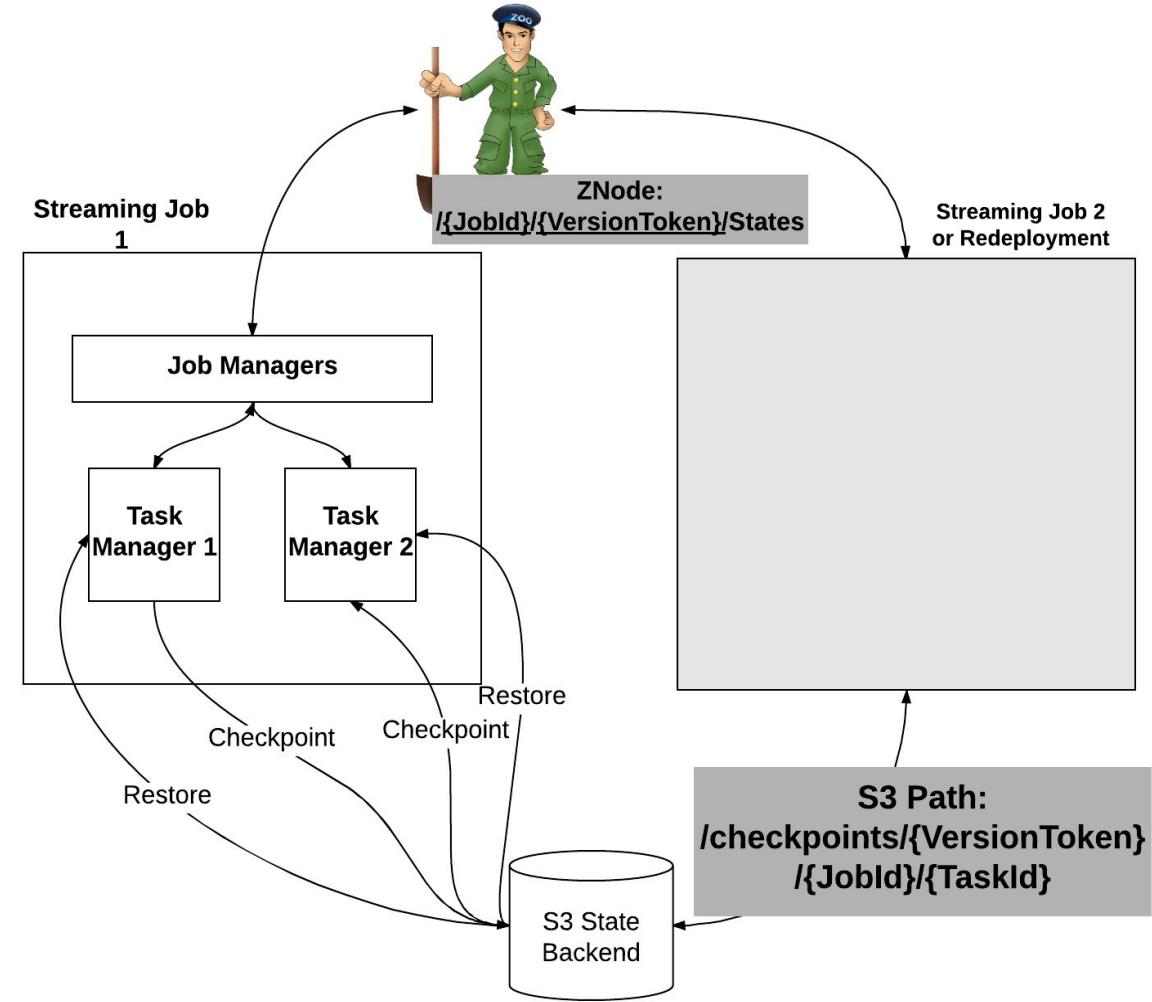


# Logical Isolation

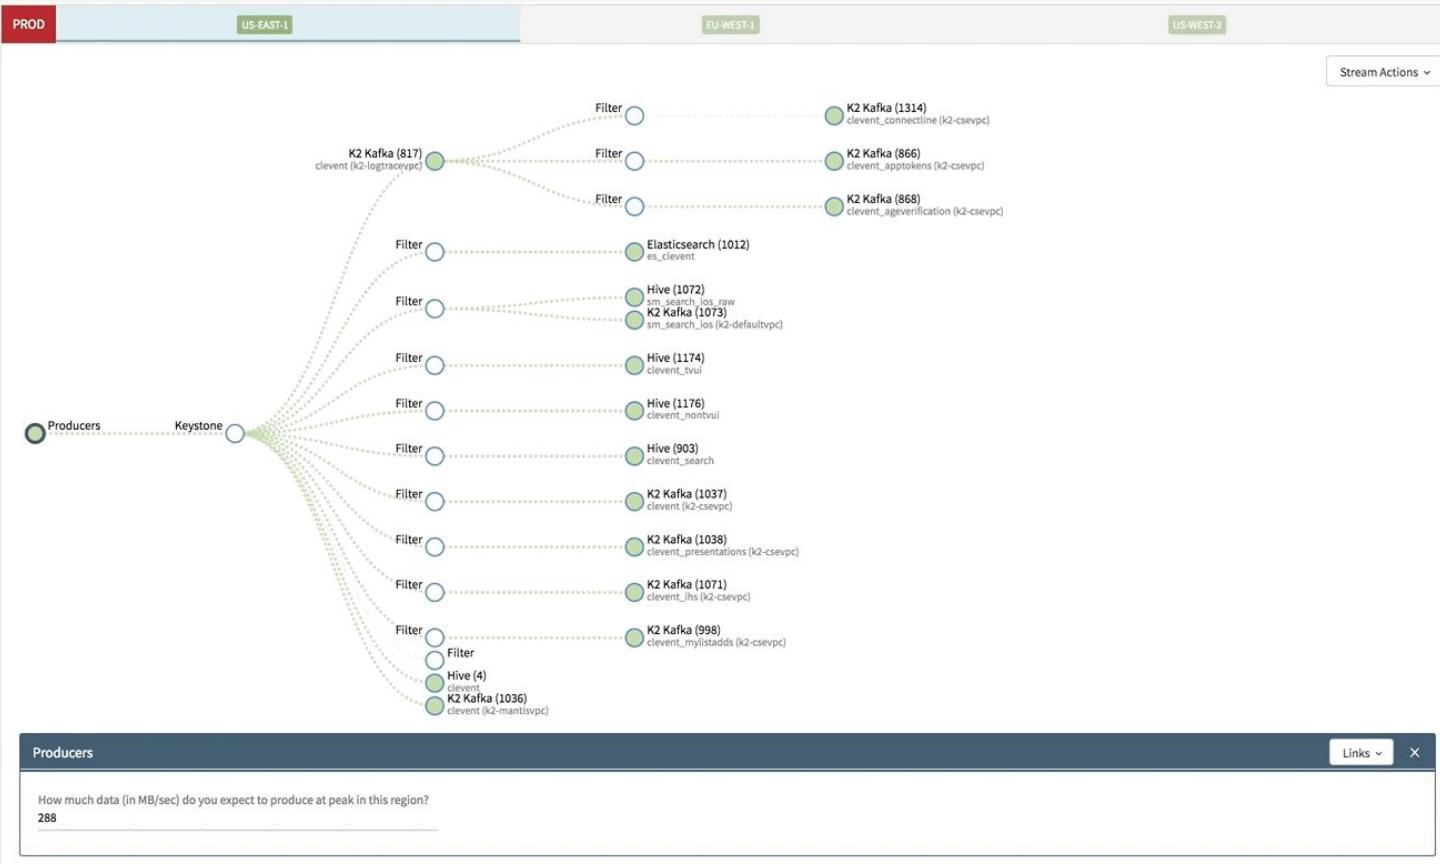


# Logical Isolation

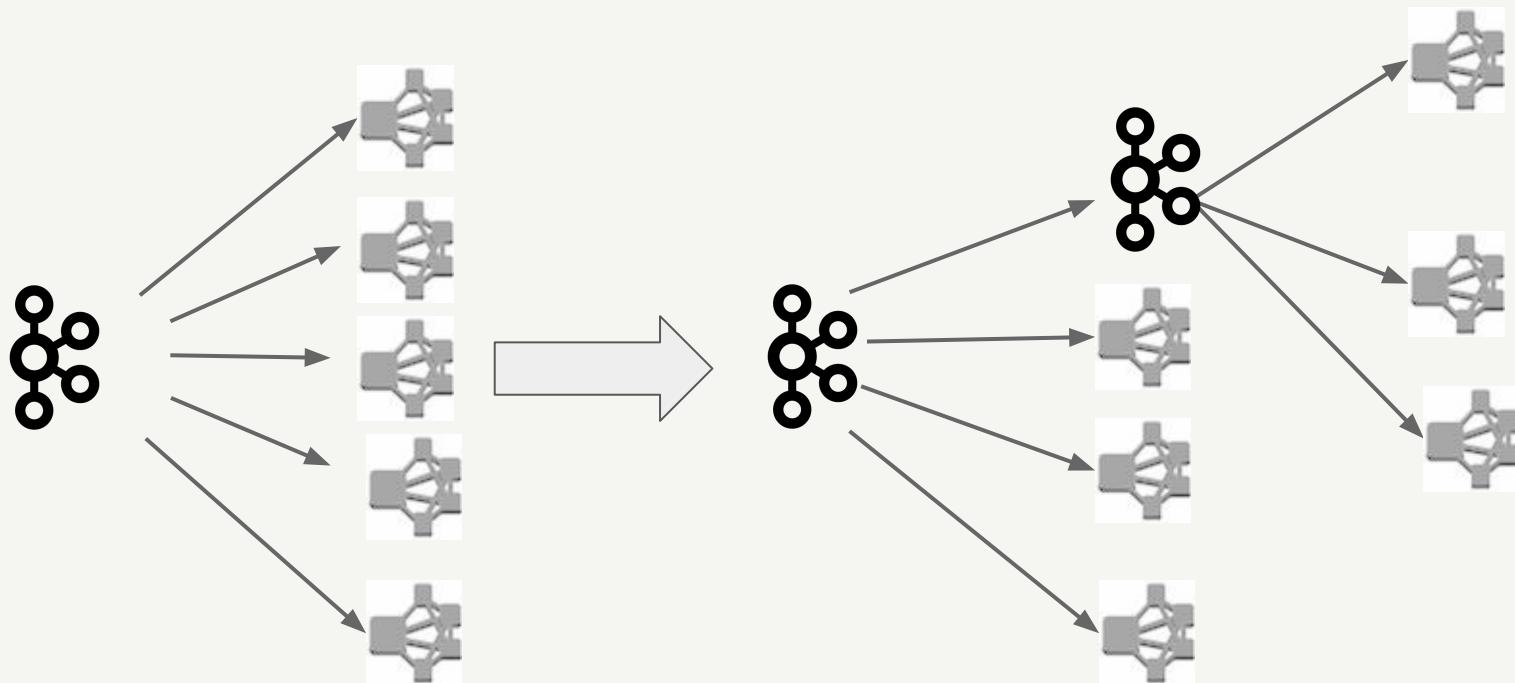
- Streams Level
- Deployments Level
- What about regional  
Island Isolation?



Total Bytes Out = (num Of Consumers + replication factor - 1) \* Bytes In



## Solve Consumer Fanout with Hierarchies



## Total Infrastructure Scale

- 500+ Billion events generated per day
- 1+ Trillion events processed per day
- ~800 Topics
- ~1,800 Streams
- 4000+ Kafka Instances
- ~9,000 Stream Processing Containers

# Kafka Clusters



Search for a Cluster...



Filter:

prod

test

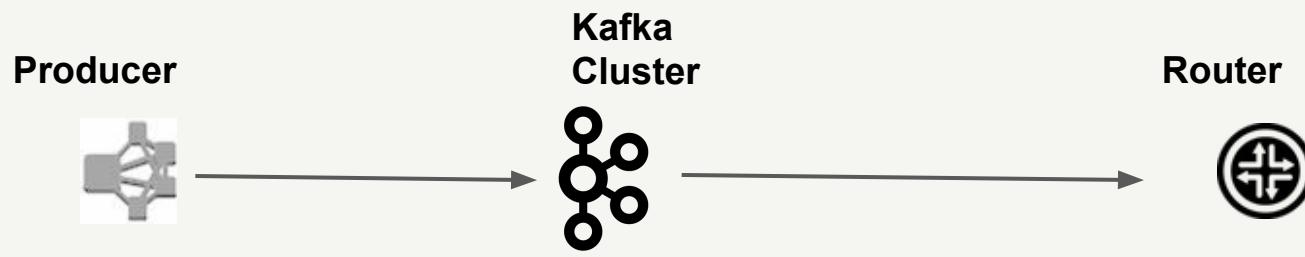
us-east-1

eu-west-1

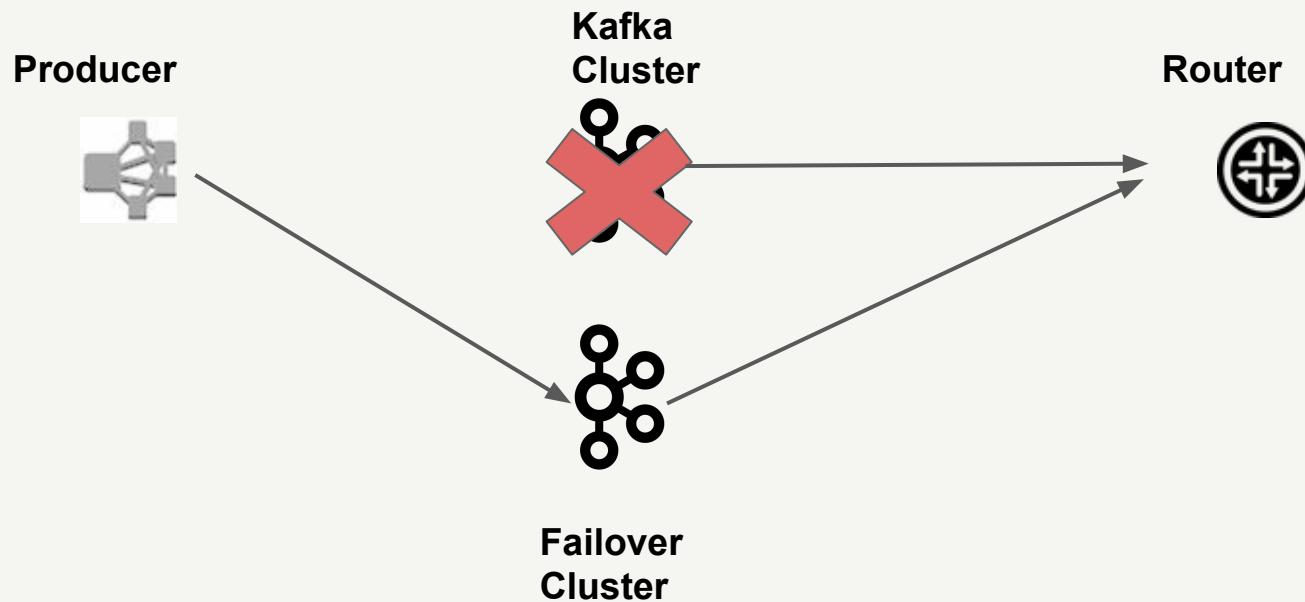
us-west-2

Cluster	Env	Region	State
whitney	prod	us-east-1	NORMAL
rocky	prod	us-east-1	NORMAL
robson	prod	us-east-1	NORMAL
himalayas	prod	us-east-1	NORMAL
everest	prod	us-east-1	NORMAL
elbert	prod	us-east-1	NORMAL
denali	prod	us-east-1	NORMAL
blanc	prod	us-east-1	NORMAL

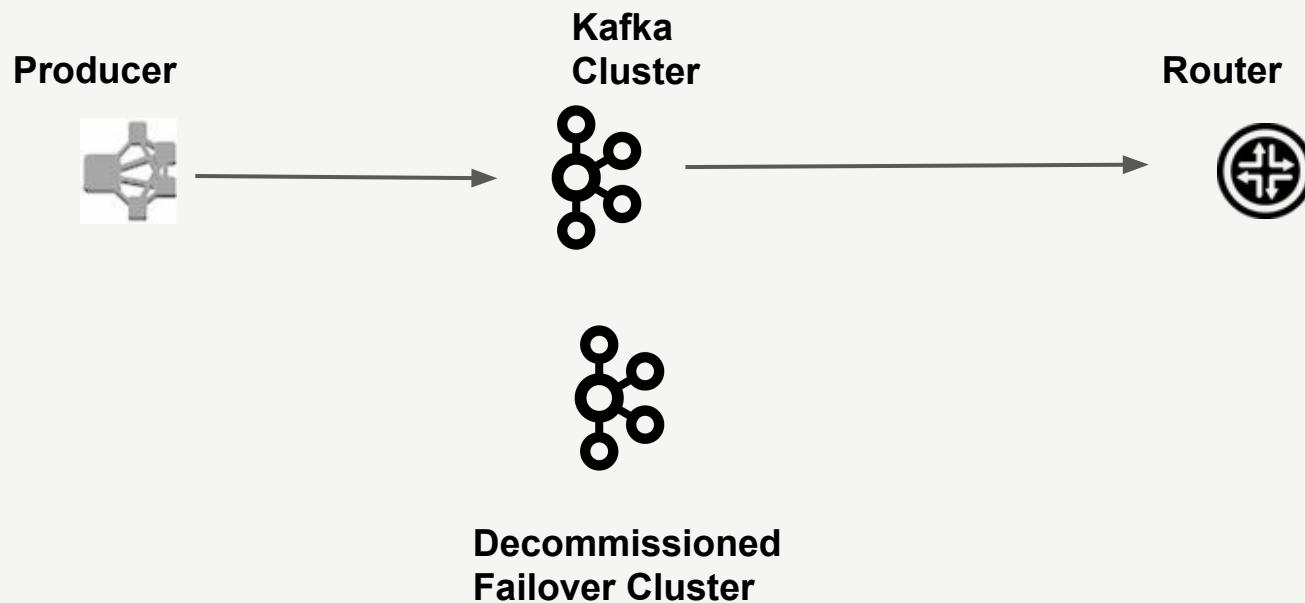
# Kafka Cluster Failover



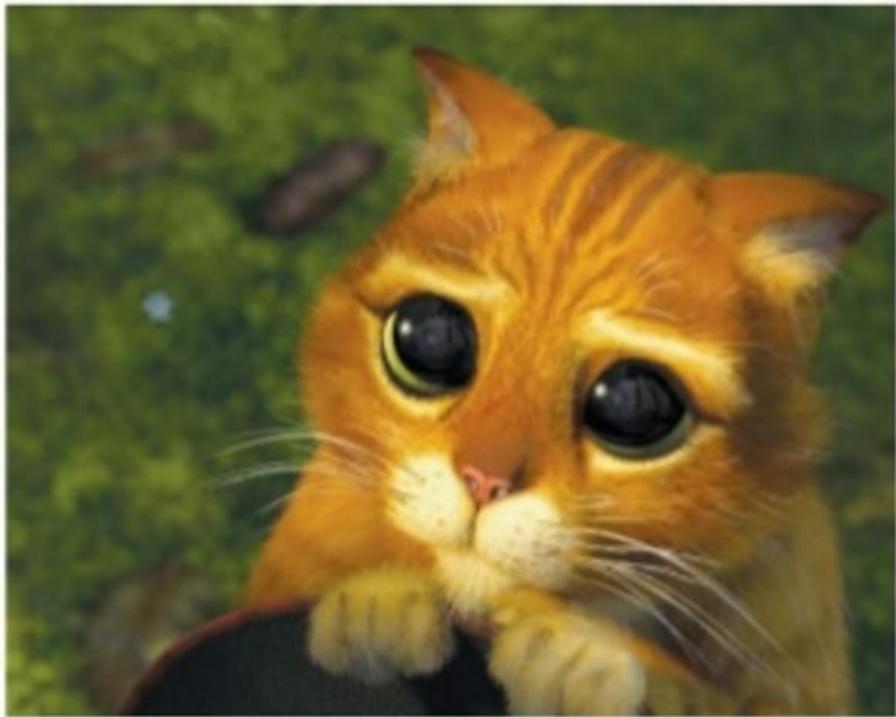
## **Failover**



## **Fallback**



## Pet vs Cattle





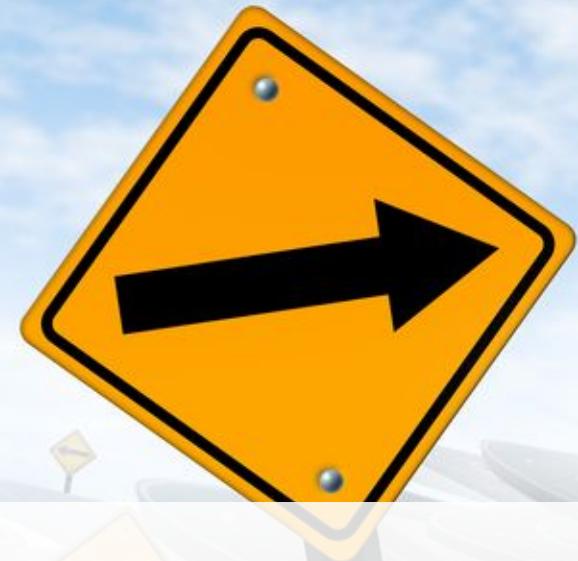
# Immutable Kafka Clusters ?

**Principles:**  
**Failure as a First Class Citizen**  
**Separation of Concerns**  
**Embrace Immutability**

## **Challenge 2:** **Self-serv & Multi-tenants**

## **Diverse Customer Requirements**

- Diverse combination of features
- Diverse platform tradeoffs



**“Change is the Only Constant”**



## **Change Is the Only Constant**

- New streams deployed in a few mins
- Customer changing needs
- Scaling activity
- Infrastructure Upgrades

## **Failure Modes**

- Infrastructure Disaster
- Any component can become temporarily unavailable



Provide Building  
Blocks

# Declarative Reconciliation



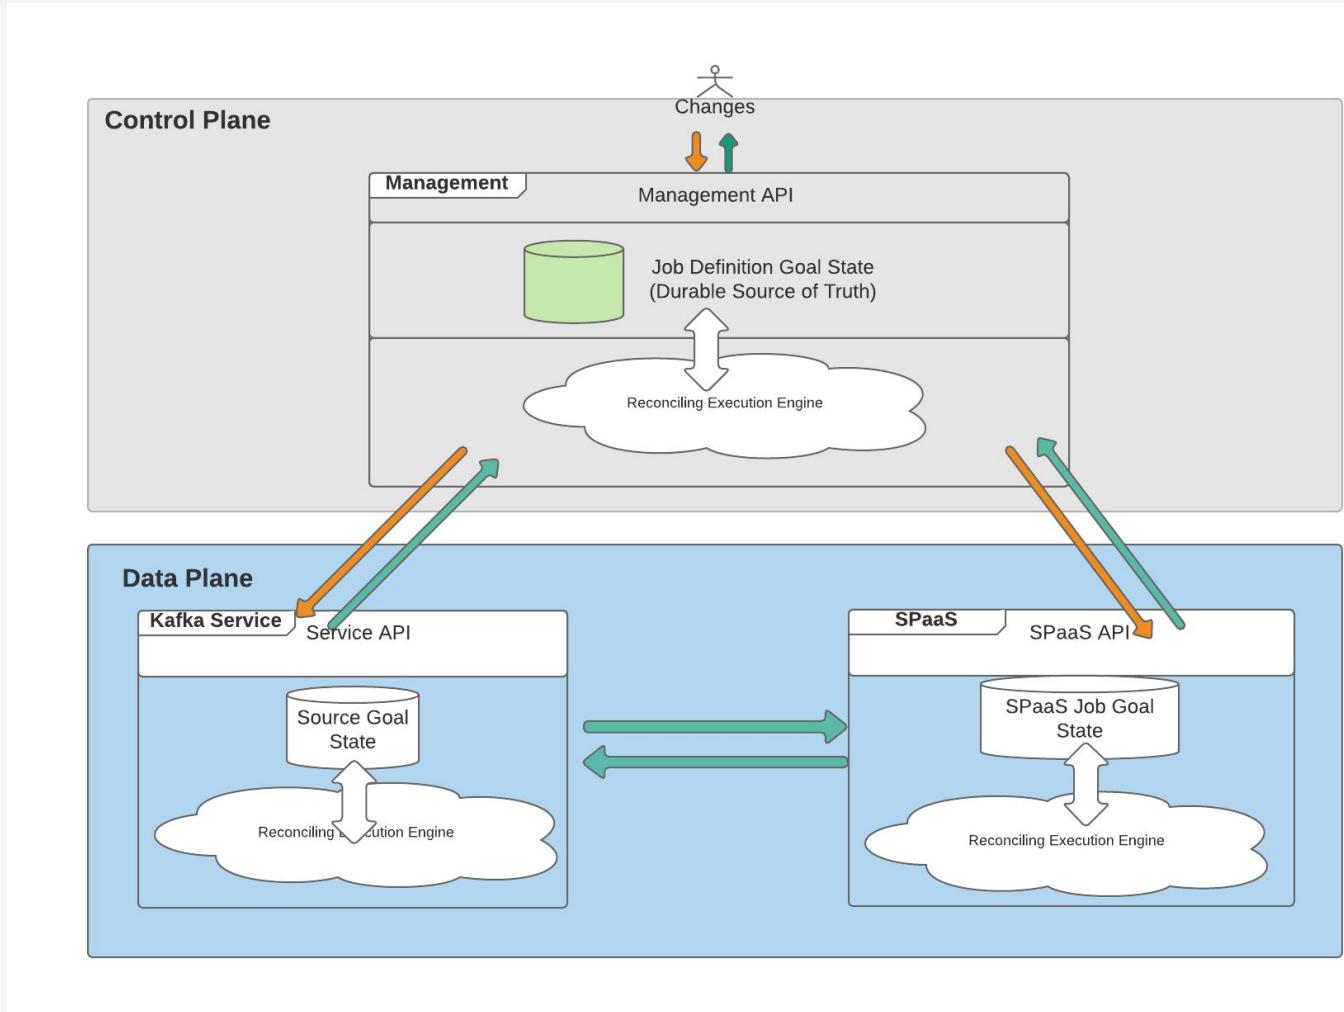
## **Declarative Reconciliation**

- “Declarative” is a communication pattern
- “Reconciliation” to drive the entire system towards goal

## **Declarative Reconciliation**

- Goal States
- Current States

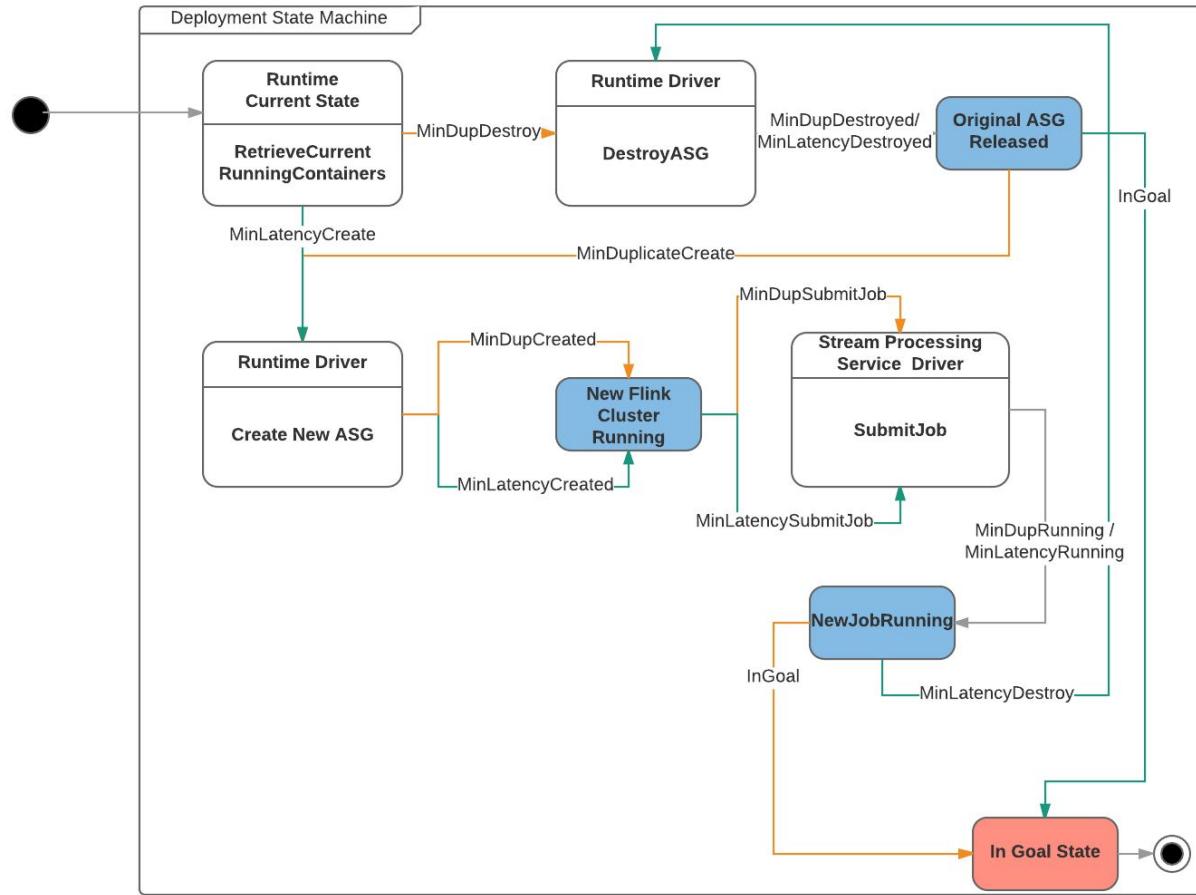
# Layered Reconciliation

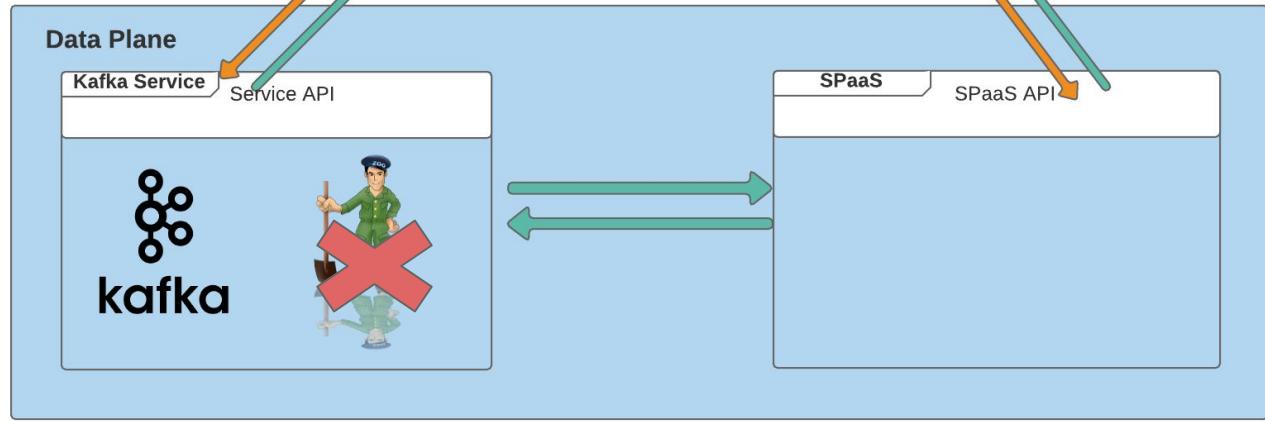
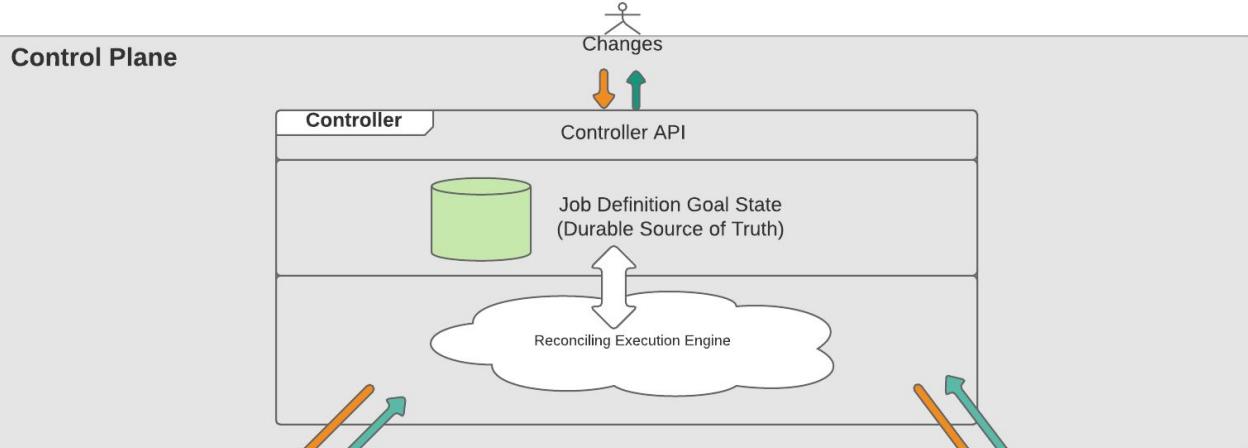


## **Declarative Reconciliation**

- Goal States
- Current States
- State Machine Driven Reconciliation

# State Machine Goal State Driver

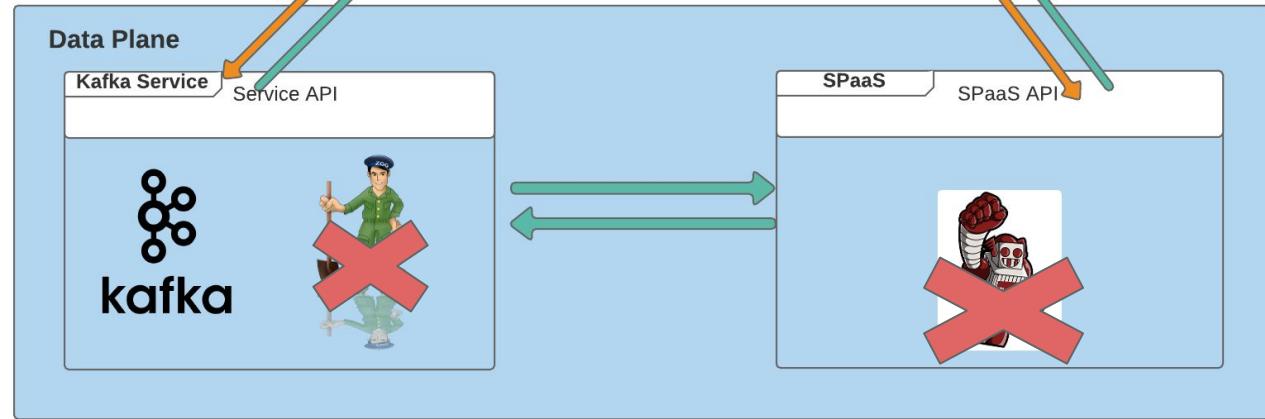
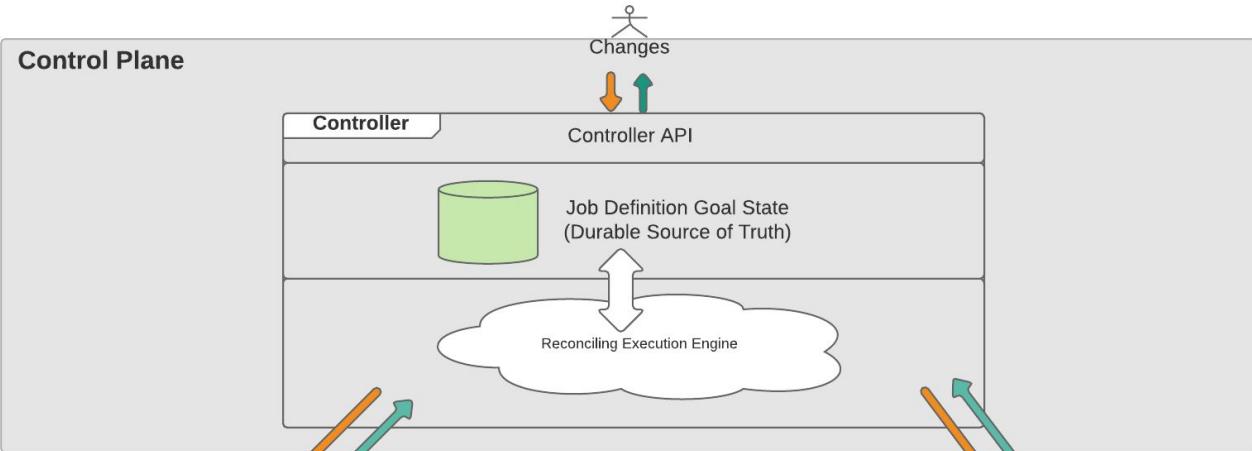




# Single Source of Truth

Allows Eventual Consistency

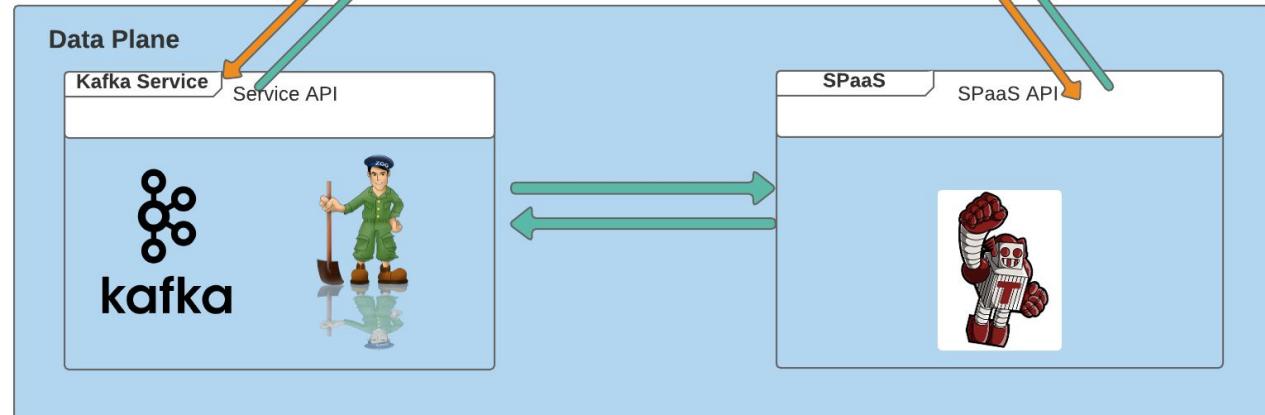
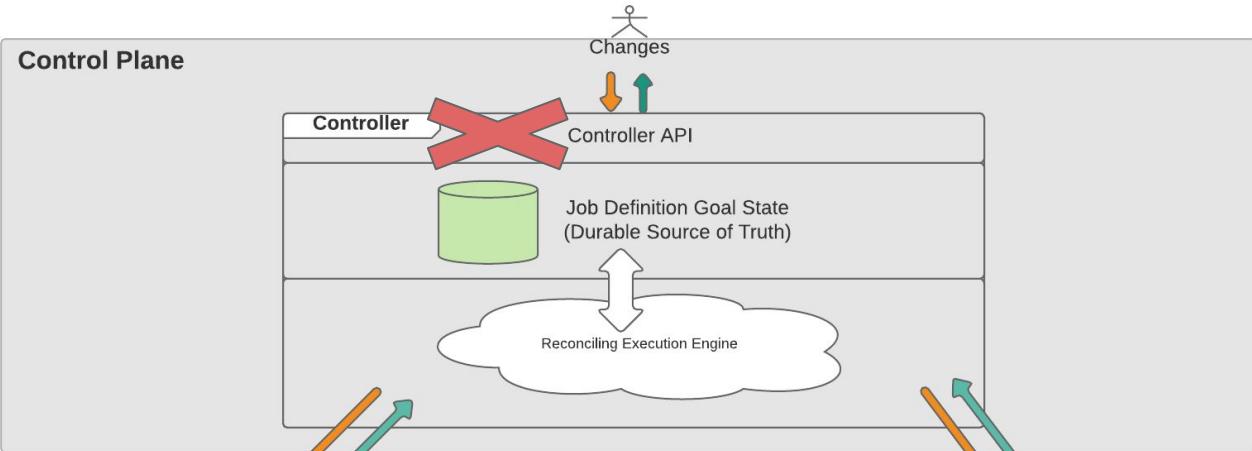
Convergence on Goal State



# Single Source of Truth

Allows Eventual Consistency

Convergence on Goal State



**Principle:**  
**Leverage Reusable Building Blocks**  
**Declarative Reconciliation**  
**Single Source of Truth**

**Thought Experiment:**

**Kitchen Management vs Distributed  
Architecture?**



# Thank you.

## References:

<https://medium.com/netflix-techblog>

<https://www.confluent.io/kafka-summit-sf17/multitenant-multicloud-and-hierarchical-kafka-messaging-service>

We're hiring - <http://bit.ly/NetflixSPaaS>



@ZhenzhongXu (tweet me questions!)

